

# Caching principles and paging performance

Professor Hugh C. Lauer  
CS-3013 — Operating Systems

(Slides include copyright materials from *Operating Systems: Three Easy Steps*, by Remzi and Andrea Arpaci-Dusseau, from *Modern Operating Systems*, by Andrew S. Tanenbaum, 3<sup>rd</sup> edition, and from other sources)

# Fundamental observation

- Paging allows us to treat physical memory as a *cache* of virtual memory
- Therefore, all of the principles and issues of caches apply to paging
- ... especially paging performance

# Definition — *cache*

- A small subset of active items held in small, *fast* storage while most of the items remain in much larger, *slower* storage
- Includes mechanisms for
  - Recognizing what items are in the cache and for accessing them quickly
  - Bringing things into cache and throwing them out again

# Note on caches and caching

- This topic is not adequately covered in Tanenbaum or most other Operating System textbooks
- Silberschatz, Galvin, & Gagne discuss paging performance somewhat
- It is treated extensively in Computer Architecture textbooks

*Caching* is a fundamental, cross-cutting concept with broad application to all areas of computing science and to many areas of system design, embedded systems, etc.

# Paging — *two* examples of caches

## ■ Paged Virtual Memory

- Very large, mostly stored on (slow) disk
- Small *working set* in (fast) RAM during execution

## ■ Page tables

- Very large, mostly stored in (slow) RAM
- Small *working set* stored in (fast) TLB registers

# *Caching* is ubiquitous in computing

## ■ Transaction processing

- Keep records of today's departures in RAM or local storage while records of future flights are on remote database

## ■ Program execution

- Keep the bytes near the current program counter in on-chip memory while rest of program is in RAM

## ■ File management

- Keep disk maps of open files in RAM while retaining maps of all files on disk

## ■ Game design

- Keep nearby environment in cache of each character

■ ...

# Caching issues

*This is a very important list!*

- When to put something in the cache
- What to throw out to create cache space for new items
- How to keep cached item and stored item in sync after one or the other is updated
- How to keep multiple caches in sync across processors or machines
- Size of cache needed to be effective
- Size of cache items for efficiency
- ...

# General observation on *caching*

## ■ We create caches because

- There is not enough fast memory to hold everything we need
- Memory that *is* large enough is too slow

## ■ Performance metric for all caches is *EAT*

- *Effective Access Time*

## ■ Goal is to make overall performance close to cache memory performance

- By taking advantage of *locality* — temporal and spatial
- By burying a small number of accesses to slow memory under many, many accesses to fast memory



# Definition – Effective Access Time (EAT)

- The *average access time* to memory items, where some items are cached in fast storage and other items are not cached...

- ...weighted by  $p$ , the *fault rate*

- $0 \leq p < 1$

Computer architecture textbooks use *Average Access Time* =  $(\text{cache access time}) + p * (\text{miss penalty})$

- $EAT = (1-p) * (\text{cache access time}) + p * (\text{non-cache access time})$

# Goal of Caching

- To take advantage of locality to achieve nearly the same performance of the *fast* memory when most data is in *slow* memory

- I.e., solve *EAT* equation for  $p$

$$(1-p) * (\text{cache\_time}) + p * (\text{non\_cache\_time}) < \text{EAT} \quad \leftarrow$$

$$(1+x) * (\text{cache\_time})$$

$x$  is "acceptable" performance penalty

$$p < x * \frac{\text{cache\_time}}{(\text{non\_cache\_time} - \text{cache\_time})}$$

# Goal of Caching (continued)

- Select *size of cache* and *size of cache items* so that  $p$  is low enough to meet acceptable performance goal
- Usually requires simulation of a suite of benchmarks

# Application to Demand Paging

## ■ ***Page Fault Rate ( $p$ )***

$0 \leq p < 1.0$  (measured in average number of faults / reference)

## ■ ***Page Fault Overhead***

*= fault service time + read page time + restart process time*

- *Fault service time*  $\sim 0.1\text{--}10\ \mu\text{sec}$
- *Restart process time*  $\sim 0.1\text{--}10\text{--}100\ \mu\text{sec}$
- *Read page time*  $\sim 8\text{--}20\ \text{milliseconds!}$

## ■ **Dominated by time to read page in from disk!**

# Demand Paging Performance (continued)

- **Effective Access Time (*EAT*)**

$$= (1-p) * (\text{memory access time}) + \\ p * (\text{page fault overhead})$$

- **Want *EAT* to degrade no more than, say, 10% from true *memory access time***

- i.e.,  $EAT < (1 + 10\%) * \text{memory access time}$

# Performance Example

- *Memory access time* = 100 nanosec =  $10^{-7}$
- *Page fault overhead* = 25 millisec = 0.025
- *Page fault rate* =  $1/1000 = 10^{-3}$
  
- $EAT = (1-p) * 10^{-7} + p * (0.025)$   
 $= (0.999) * 10^{-7} + 10^{-3} * 0.025$   
 $\cong 25 \text{ microseconds per reference!}$
  
- *I.e.,*  
 $= 250 * \text{memory access time!}$

# Performance Goal

- To achieve less than 10% degradation

$$(1-p) * 10^{-7} + p * (0.025) < 1.1 * 10^{-7}$$

*i.e.,*

$$p < (0.1 * 10^{-7}) / (0.025 - 10^{-7})$$

$$\cong 0.0000004$$

- *i.e.,*

1 fault in 2,500,000 accesses!



# Working Set Size

- Assume average swap time of 25 millisec.
- For *memory access time* = 100 nanoseconds
  - Require < 1 page fault per 2,500,000 accesses
- For *memory access time* = 1 microsecond
  - Require < 1 page fault per 250,000 accesses
- For *memory access time* = 10 microseconds
  - Require < 1 page fault per 25,000 accesses



# Object Lesson

- Working sets *must* get larger in proportion to memory speed!
  - Disk speed  $\sim$  constant (nearly)
- I.e., faster computers *need* larger physical memories to exploit the speed!

# Class Discussion

- 1. What is first thing to do when the PC you bought last year becomes too slow?**
- 2. What else might help?**
- 3. Can we do the same analysis on TLB performance?**



# TLB fault performance

## ■ Assumptions

- $m = \text{memory access time} = 100 \text{ nsec}$
- $t = \text{TLB load time from memory} = 300 \text{ nsec}$   
 $= 3 * m$

## ■ Goal is < 5% penalty for TLB misses

- i.e.,  $EAT < 1.05 * m$

## ■ How low does TLB fault rate need to be?

# TLB fault performance

## ■ Assumptions

- $m = \text{memory access time} = 100 \text{ nsec}$
- $t = \text{TLB load time from memory} = 300 \text{ nsec}$   
 $= 3 * m$

## ■ Goal is < 5% penalty for TLB misses

- i.e.,  $EAT < 1.05 * m$

## ■ $EAT = (1-p) * m + p * t < 1.05 * m \Rightarrow$

$$p < (0.05 * m) / (t - m)$$

$$= 0.05 * m / (2 * m)$$

$$= 0.025$$

## ■ i.e., TLB fault rate should be < 1 per 40 accesses!

# TLB fault performance (continued)

■ **Q: How large should TLB be so that TLB faults are not onerous, in these circumstances?**

■ **A: Somewhat less than 40 entries**

- Assuming a reasonable degree of locality!

**Note: pathological cases may prevent goal altogether!**

# What if Software Loaded TLB?

- E.g., with hashed or inverted page tables?
- Assume TLB load time is  $100 * m$
- *Work out on white board*

$$p < x * \frac{cache\_time}{(non\_cache\_time - cache\_time)}$$

# Summary of this Topic

- A quantitative way of estimating how large the *cache* needs to be to avoid excessive thrashing, where
  - *Cache* = Working set in physical memory
  - *Cache* = TLB size in hardware
- Applicable to all forms of *caching*

# General Observation on *Caching*

## ■ We create caches because

- There is not enough fast memory to hold everything we need
- Memory that *is* large enough is too slow

## ■ Performance metric for all caches is EAT

- *Effective Access Time*

## ■ Goal is to make overall performance close to cache memory performance

- By taking advantage of *locality* — temporal and spatial
- By burying a small number of accesses to slow memory under many, many accesses to fast memory



# Cache Applications

- ***Physical memory: cache of virtual memory***
  - I.e., RAM over disk
- ***TLB: cache of page table entries***
  - I.e., Registers over RAM
- ***Processor L2 cache: over RAM***
  - I.e., nanosecond memory over 10's of nanoseconds
- ***Processor L1 cache: over L2 cache***
  - I.e., picosecond registers over nanosecond memory
- ...

# Cache Applications (continued)

- **Recently accessed blocks of a file**

- I.e., RAM over disk blocks

- **Today's airline flights**

- I.e., local disk over remote disk

# Questions?