

Reminders:—

- **Project 4 due *Friday, February 24***
 - Late penalties described in syllabus on InstructAssist
- **Any project submitted after 4 PM on Tuesday, Feb 27 *will not be graded!***
- **Any project not demonstrated by 6 PM on Thursday, March 1, *will not be graded***
- **Schedule your demos ASAP!**
 - If not enough slots, be sure to ask for more *early!*

Input and Output

Professor Hugh C. Lauer

CS-3013 — Operating Systems

(Slides include copyright materials from *Operating Systems: Three Easy Steps*, by Remzi and Andrea Arpaci-Dusseau, from *Modern Operating Systems*, by Andrew S. Tanenbaum, 3rd edition, and from other sources)

While a typical Linux or Windows distribution includes support for > 50 file systems ...

... these systems are called upon to support 1000's of I/O devices and subsystems

The I/O subsystem

- The largest, most complex subsystem in an OS
- Most lines of code
- Highest rate of code changes
- Where OS engineers most likely to work
- Difficult to test thoroughly
- **Make-or-break issue for any system**
 - Big impact on performance and perception
 - Bigger impact on acceptability in market


Overview

- What is I/O?
- Principles of I/O hardware
- Principles of I/O software
- Methods of implementing input-output activities
- Organization of device drivers
- Specific kinds of devices

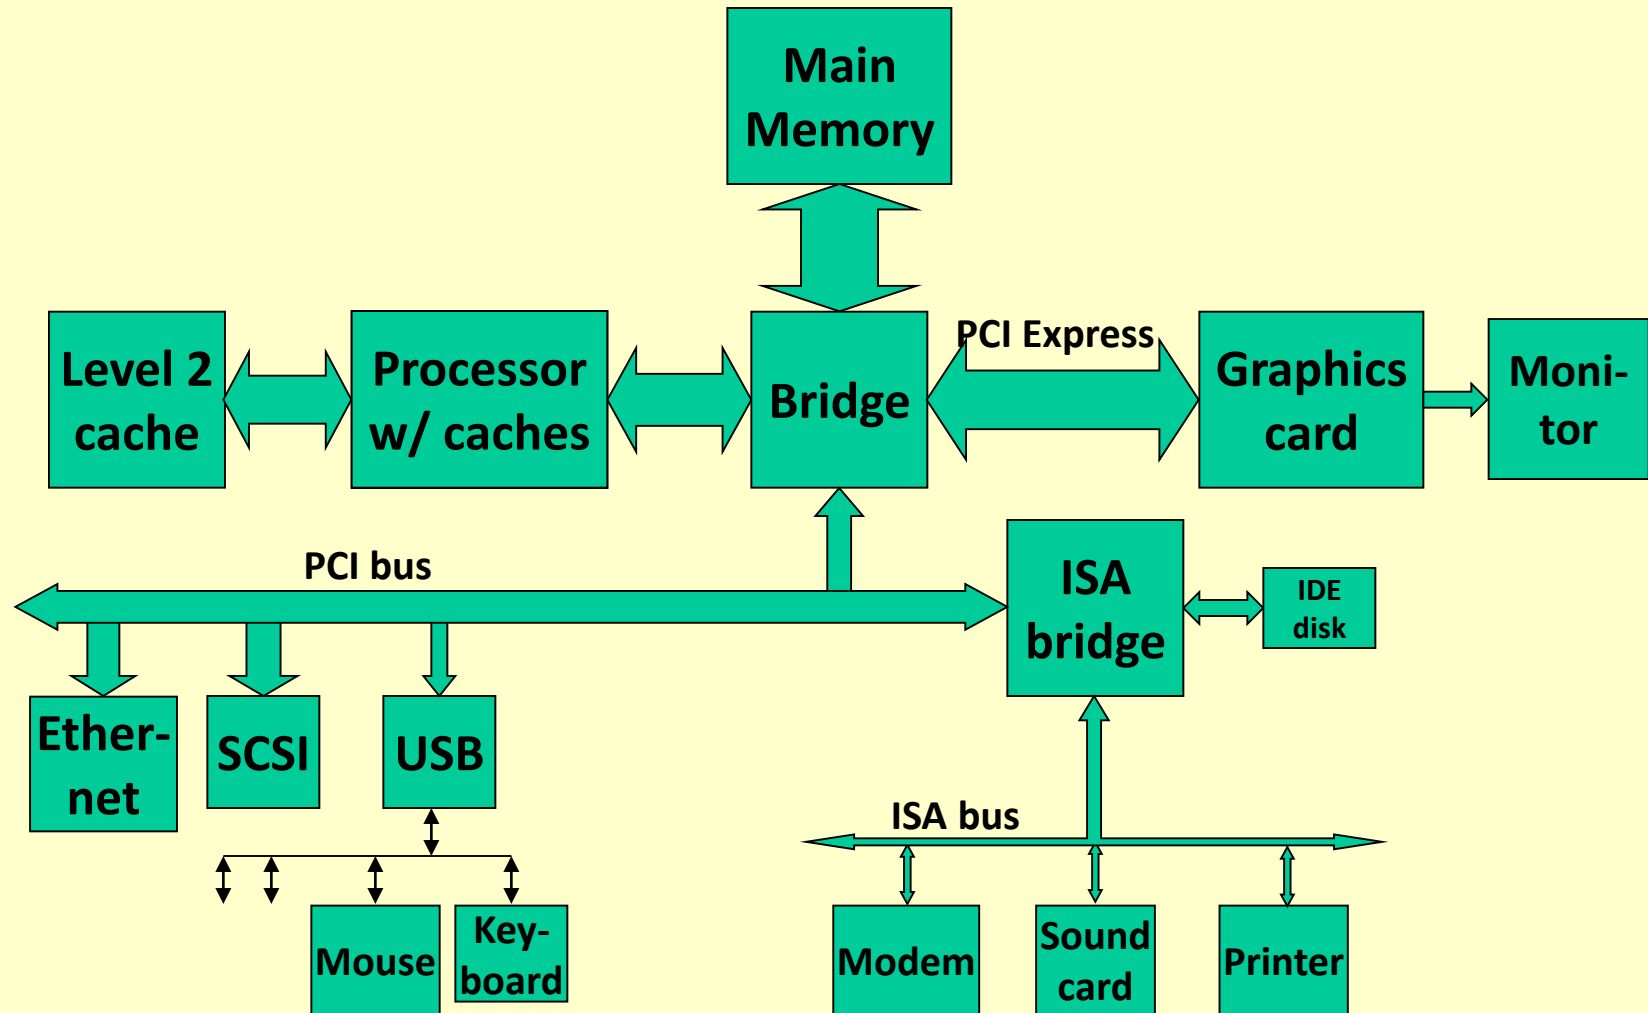
(OSTEP, §36-38)

- Treats disks as I/O devices rather than repositories for files


Hardware organization (simple, naïve)



Hardware organization (PC not long ago)



Hardware organization (more recent PC)



Kinds of I/O devices

■ Character (and sub-character) devices

- Mouse, character terminal, joy stick, some keyboards

■ Block transfer

- Disk, tape, CD, DVD
- Network

■ Clocks

- Internal, external

■ Graphics

- GUI, games

■ Multimedia

- Audio, video

■ Other

- Sensors, cameras, controllers, touch screens, etc.



Controlling an I/O device

■ A function of host Processor architecture

- Two approaches:– Special instructions vs. memory-mapped

■ Special I/O instructions

- Opcode to stop, start, query, etc.
- Separate I/O address space
- Kernel mode only

■ Memory-mapped I/O control registers

- Each register has a physical memory address
- Writing to data register is output
- Reading from data register is input
- Writing to control register causes action
- Can be mapped to kernel or user-level virtual memory



Character device (example)

■ ***Data register:***

- Register or address where data is read from or written to
- Very limited capacity (at most a few bytes)

■ ***Action register:***

- When writing to register, causes a physical action
- Reading from register yields zero

■ ***Status register:***

- Reading from register provides information
- Writing to register is no-op

Block transfer device (example)

■ ***Buffer address register:***

- Points to area in *physical* memory to read or write data
- or*

■ **Addressable *buffer* for data**

- E.g., network cards, modern hard drives

Keyboard/mouse

■ ***Action register:***

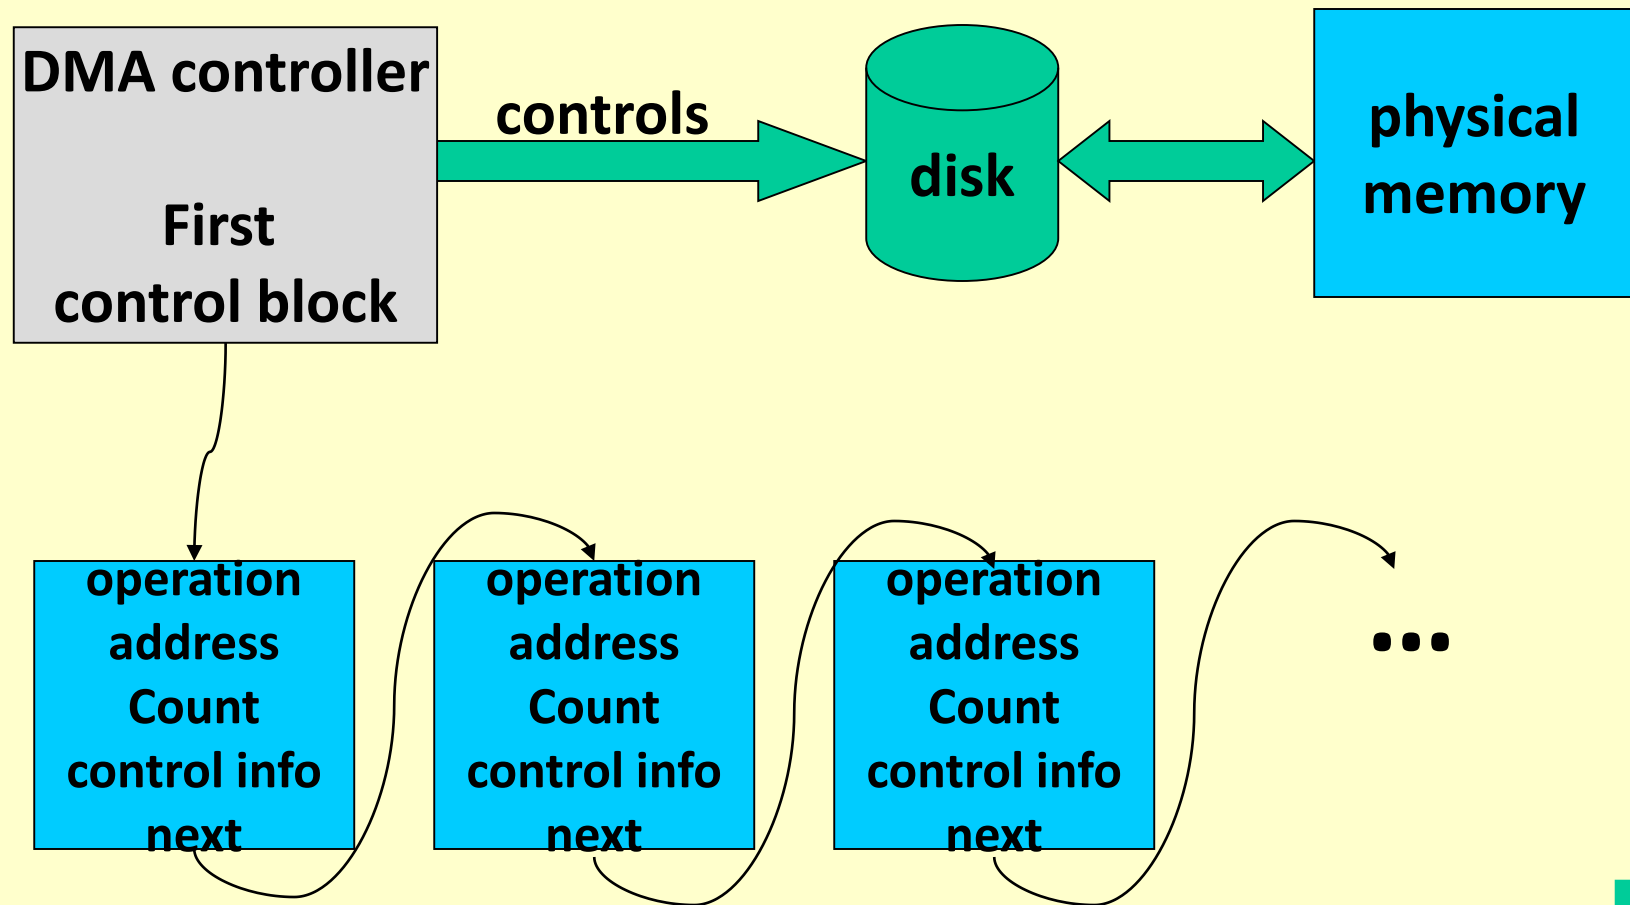
- When writing to register, initiates a physical action or data transfer
- Reading from register yields zero

■ ***Status register:***

- Reading from register provides information
- Writing to register is no-op

DMA (Direct Memory Access)

- **Block devices to autonomously read from and/or write to main memory**
 - (Usually) physical addresses
 - Compete with processor(s) for access to bus, memory
- **Transfer address**
 - Points to location in physical memory
- **Action register:**
 - Initiates a reading of control block chain to start actions
- **Status register:**
 - Reading from register provides information



Questions?

Overview

- What is I/O?
- Principles of I/O hardware
- **Principles of I/O software**
- Methods of implementing input-output activities
- Organization of device drivers
- Specific kinds of devices

(OSTEP, §36-38)

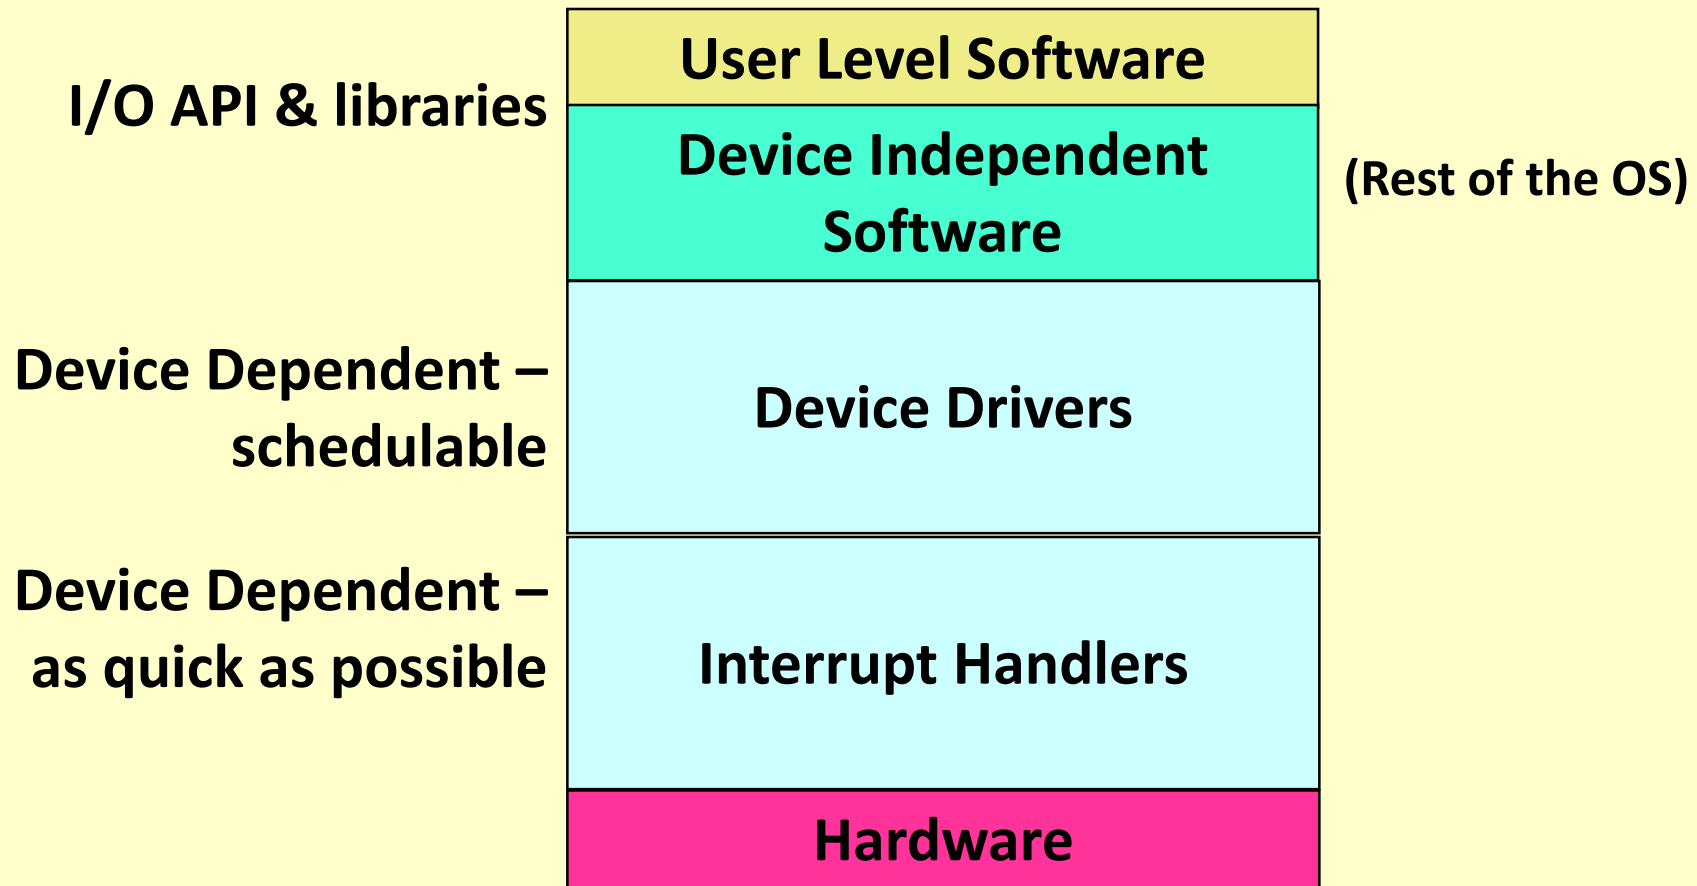
Requirements of I/O software

- ***Efficiency*** – Do not allow I/O operations to become system bottleneck
 - Especially slow devices
- ***Device independence*** – isolate OS and application programs from device specific details and peculiarities
- ***Uniform naming*** – support a way of naming devices that is scalable and consistent
- ***Error handling*** – isolate the impact of device errors, retry where possible, provide uniform error codes
 - Errors are abundant in I/O
- ...

Principles of I/O software (continued)

- ...
- ***Buffering*** – provide uniform methods for storing and copying data between physical memory and the devices
- ***Uniform data transfer modes*** – synchronous and asynchronous, read, write, ..
- ***Controlled device access*** – sharing and transfer modes
- ***Uniform driver support*** – specify interfaces and protocols that drivers must adhere to

I/O software “stack”



Three common ways of I/O

- Programmed I/O
- Interrupt-Driven I/O
- I/O using DMA

Programmed I/O (polling)

- **Used when device and controller are relatively quick to process an I/O operation**
- **Device driver**
 - Gains access to device
 - Initiates I/O operation
 - Loops testing for completion of I/O operation (busy wait)
 - If there are more I/O operations, repeat
- **Used in following kinds of cases**
 - Service interrupt time > Device response time
 - Device has no interrupt capability
 - Embedded systems where Processor has nothing else to do

Programmed I/O example — bitmapped keyboard & mouse

- **Keyboard & mouse buttons implemented as 128-bit read-only register**
 - One bit for each key and mouse button
 - $0 = \text{“up”}$; $1 = \text{“down”}$
- **Mouse position implemented as pair of counters**
 - One increment per unit of motion in each of x and y directions
- **Periodic clock interrupt (e.g., every 10 msec)**
 - Reads keyboard register, compares to previous copy
 - Determines key & button transitions up or down
 - Decodes transition stream to form character and button sequence
 - Reads and compares mouse counters to form motion sequence

**Advantage:— key positions & mouse behavior
are programmed entirely in software**

Other programmed I/O examples

- **Check status of device**
- **Read from disk or boot device at boot time**
 - No OS present, hence no interrupt handlers
 - Needed for bootstrap loading of the inner portions of kernel
- **External sensors or controllers**
 - Real-time control systems

Questions about programmed I/O?

Interrupt-driven I/O

- **Interrupts occur on I/O events**
 - operation completion
 - Error or change of status
- **Data transferred by interrupt handler**
 - Reading or writing device registers

Processor participates in *every* byte or word transferred!

Interrupt request lines (IRQs)

■ Every device is assigned an *IRQ*

- Used when raising an interrupt
- Interrupt handler can identify the interrupting device

■ Assigning IRQs

- In older and simpler hardware, physical wires and contacts on device or bus
- In most modern PCs, etc., assigned dynamically at boot time

Handling interrupts (Linux style)

■ Terminology

- *Interrupt context* – kernel operating not on behalf of any process
- *Process context* – kernel operating on behalf of a particular process
- *User context* – process executing in user-space virtual memory

■ ***Interrupt Service Routine (ISR), also called Interrupt Handler***

- The function that is invoked when an interrupt is raised
- Identified by IRQ
- Operates on *Interrupt stack* (as of Linux kernel 2.6)
 - *One* interrupt stack per processor or core
 - Approx 4-8 kbytes
 - *All handlers share same stack on processor!*

■ ...

Handling interrupts (Linux style – continued)

Definitions!

- ...
- ***Top half*** – does minimal, time-critical work necessary
 - Acknowledge interrupt, reset device, copy buffer or registers, etc.
 - Interrupts (usually) disabled on current processor
- ***Bottom half*** – the part of the ISR that can be deferred to more convenient time
 - Completes I/O processing; does most of the work
 - Interrupts enabled (usually)
 - Communicates with processes
 - Possibly in a kernel thread (or even a user thread!)

Interrupt-driven I/O example — software time-of-day clock

■ Interrupt occurs at fixed intervals

- 50 or 60 Hz
- 1 kHz in more modern processors

■ Service routine (*top half*):—

- Adds one *tick* to clock counter

■ Service routine (*bottom half*):—

- Checks list of *soft timers*
- Notifies tasks of any expired timers

Note that this looks a lot like programmed I/O

- Except for bottom-half processing

Other interrupt-driven I/O examples

■ Very “slow” character-at-a-time terminals

- Mechanical printers (15 characters/second)
- Some keyboards (one character/keystroke)
 - Command-line completion in many Unix systems
- Game consoles
- Serial modems
- Many I/O devices in “old” computers
 - Paper tape, punched cards, etc.

Programmed I/O vs. interrupt-driven I/O

■ Programmed I/O

- A process or thread pro-actively works the device, gets or puts the data, does everything else.

■ Interrupt-driven I/O

- Device operates autonomously, let's processor know when it is done or ready

■ Both

- *Processor participates in transfer of every byte or word.*

DMA interrupt handler

Device transfers
data itself

- **Service Routine – *top half* (interrupts disabled)**
 - Does as little work as possible and returns
 - (Mostly) notices completion of one transfer, starts another
 - (Occasionally) checks for status
 - Setup for more processing in bottom half

- **Service Routine – *bottom half* (interrupts enabled)**
 - Compiles control blocks from I/O requests
 - Manages & pins buffers, translates virtual to physical addresses
 - Posts completion of transfers to requesting applications
 - Unpin and/or release buffers
 - Possibly in a kernel thread

DMA examples

- Disks, CD-ROM readers, DVD readers
- Ethernet & wireless “modems”
- Tape and bulk storage devices
- Common themes:–

Not GUIs

- Device controller has space to buffer a (big) block of data
- Controller has intelligence to update physical addresses and transfer data
- Controller (often) has intelligence to interpret a sequence of control blocks without processor help

- ***Processor does not touch data during transfer!***

Buffering

■ DMA devices need memory to read from, write to

- Must be contiguous pages
- (Usually) physical addresses

■ *Double buffering*

- One being filled (or emptied) by device
- Other being emptied (or filled) by application
- Special case of producer-consumer with $n = 2$

Producer-Consumer:– see OSTEP §30.2

Digression:– error detection and correction

- **Most data storage and network devices have hardware error detection and correction**
- **Redundancy code added during writing**
 - Parity: detects 1-bit errors, not 2-bit errors
 - Hamming codes
 - Corrects 1-bit errors, detects 2-bit errors
 - Cyclic redundancy check (CRC)
 - Detects errors in string of 16- or 32-bits
 - Reduces probability of undetected errors to very, very low
- **Check during reading**
 - Report error to device driver
- **Error recovery:– one of principal responsibilities of a device driver!**

Questions?

Overview

- What is I/O?
- Principles of I/O hardware
- Principles of I/O software
- Methods of implementing input-output activities
- **Organization of device drivers**
- Specific kinds of devices

(OSTEP, §36-38)

Device Drivers

- **Organization**
- **Uniform interfaces to OS**
- **Uniform buffering strategies**
- **Hide device idiosyncrasies**

Device Drivers

- **Device Drivers are dependent on both the OS & device**
- ***OS dependence***
 - Meet the interface specs of the device independent layer
 - Use the facilities supplied by the OS – buffers, error codes, etc.
 - Accept and execute OS commands – e.g. *read*, *open*, etc.
- ***Device Dependent***
 - Actions during Interrupt Service routine
 - Translate OS commands into device operations
 - E.g read block *n* becomes a series of setting and clearing and interpreting device registers or interfaces
 - Note that some device drivers have layers
 - Strategy or policy part to optimize arm movement or do retries; plus a mechanism part the executes the operations

OS Responsibility to Device Driver

■ Uniform API

- *Open, Close, Read, Write, Seek* functions
- `ioctl` function as escape mechanism

■ Buffering

- Kernel functions for allocating, freeing, mapping, pinning buffers

■ Uniform naming

- `/dev/(type)(unit)`
 - type defines driver; unit says which device

■ Other

- Assign interrupt level (IRQ)
- Protection (accessibility by application, user-space routines)
- Error reporting mechanism

Abstract Overview

■ Think of I/O subsystem as a C++ or Java-style abstract class

- Uniform interface in form of specific operations (methods) and services
- Uniform state information


■ Each I/O driver implements a subclass

- Own methods for uniform interface
- Additional state info & methods for specific needs

■ However, no Java or C++ support in kernel

- Must implement everything in long-hand in C
- Just like with file systems!

Operating System Organization



**Dynamically loaded
at run time!**

**Dynamically loaded
at run time!**

Installing Device Drivers

■ Classic Unix

- Create and compile driver to .o file
- Edit and re-compile device table to add new device
- Re-link with .o files for OS kernel \Rightarrow new boot file

■ Classic Macintosh

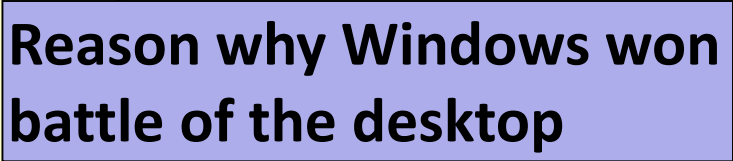
- Submit to Apple for verification, approval, and inclusion

■ MS-DOS and Windows

- Dynamic driver loading and installation
- Special driver-level debuggers available; *open device environment*
- Certification program for trademarking

■ Linux

- Dynamic driver loading and installation
- *Open device environment*



**Reason why Windows won
battle of the desktop**

Dynamic Device Configuration

■ At boot time:–

1. Probe hardware for inventory of devices & addresses
2. Map devices to drivers (using table previously created)
3. Load necessary drivers into kernel space, register in interrupt vector (.sys files in Windows)

■ Run time:–

1. Detect interrupt from newly added device
2. Search for driver, or ask user; add to table
3. Load into kernel space, register in interrupt vector

Allocating and Releasing Devices

- **Some devices can only be used by one application at a time**
 - CD-ROM recorders
 - GUI interface
- **Allocated at *Open()* time**
- **Freed at *Close()* time**

Overview


- What is I/O?
- Principles of I/O hardware
- Principles of I/O software
- Methods of implementing input-output activities
- Organization of device drivers
- **Specific kinds of devices**

(OSTEP, §36-38)

A special kind of device —graphical user interface

- ***aka, the bitmapped display***
 - In IBM language:– “all points addressable”
- **300K *pixels* to 2M *pixels***
- **Each pixel may be separated written**
- **Collectively, they create**
 - Windows
 - Graphics
 - Images
 - Videos
 - Games


GUI Device — early days



- **Bitmap in main memory**
- **All output via library routines to bitmap**
 - Entirely (or mostly) in user space
- **Controller, an automaton to do:—**
 - D-A conversion (digital to analog video)
 - 60+ Hz refresh rate
 - “clock” interrupt at top of each frame

GUI Device — Displaying Text

- **Font:** an array of bitmaps, one per character
 - Designed to be pleasing to eye
- ***bitblt*:** (*Bit-oriented Block Transfer*)
 - An operation to copy a rectangular array of pixels from one bitmap to another



GUI Device – Color

- ***Monochrome: one bit per pixel***
 - *foreground vs. background*
- ***Color: 2-32 bits per pixel***
- ***Direct vs. Color palette***
 - *Direct: (usually) 8 bits each per Red, Green, Blue*
 - *Palette: a table of length 2^p , for p -bit pixels*
Each entry (usually) 8 bits each for RGB

GUI Device – Cursor

- **A small bitmap to overlay main bitmap**
- **Hardware support**
 - Substitute cursor bits during each frame
- **Software implementation**
 - *Bitblt* area under cursor to temporary bitmap
 - *Bitblt* cursor bitmap to main bitmap
 - Restore area under cursor from temporary bitmap
- **Very, very tricky!**
 - Timing is critical for smooth appearance
 - Best with double-buffered main bitmap

GUI Device – Window

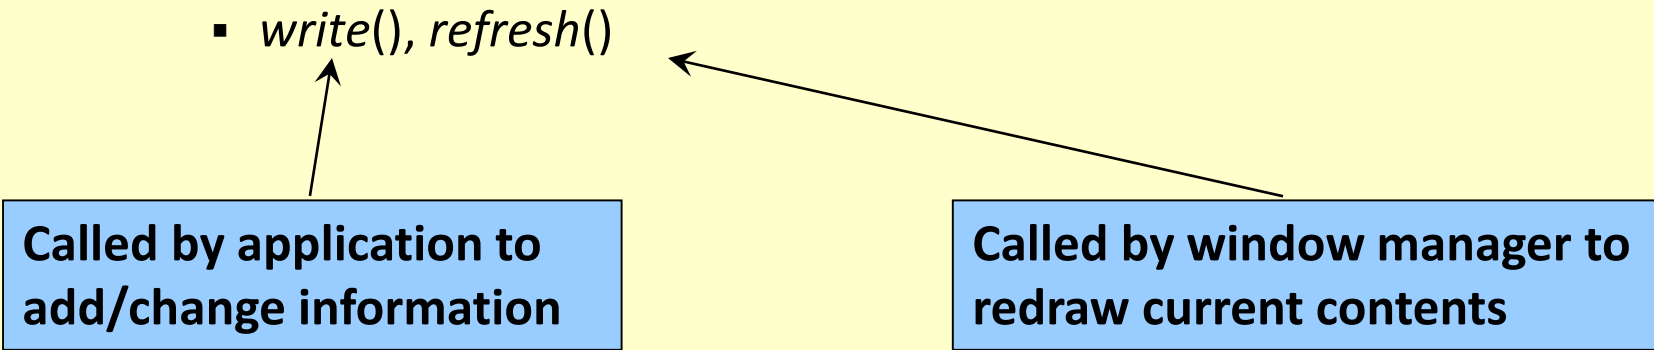
■ A virtual bitmap

- *size, position, clipping boundaries*
- *font, foreground and background colors*
- A list of operations needed to redraw contents

■ Operations to window itself:–

- *write(), refresh()*

Called by application to
add/change information



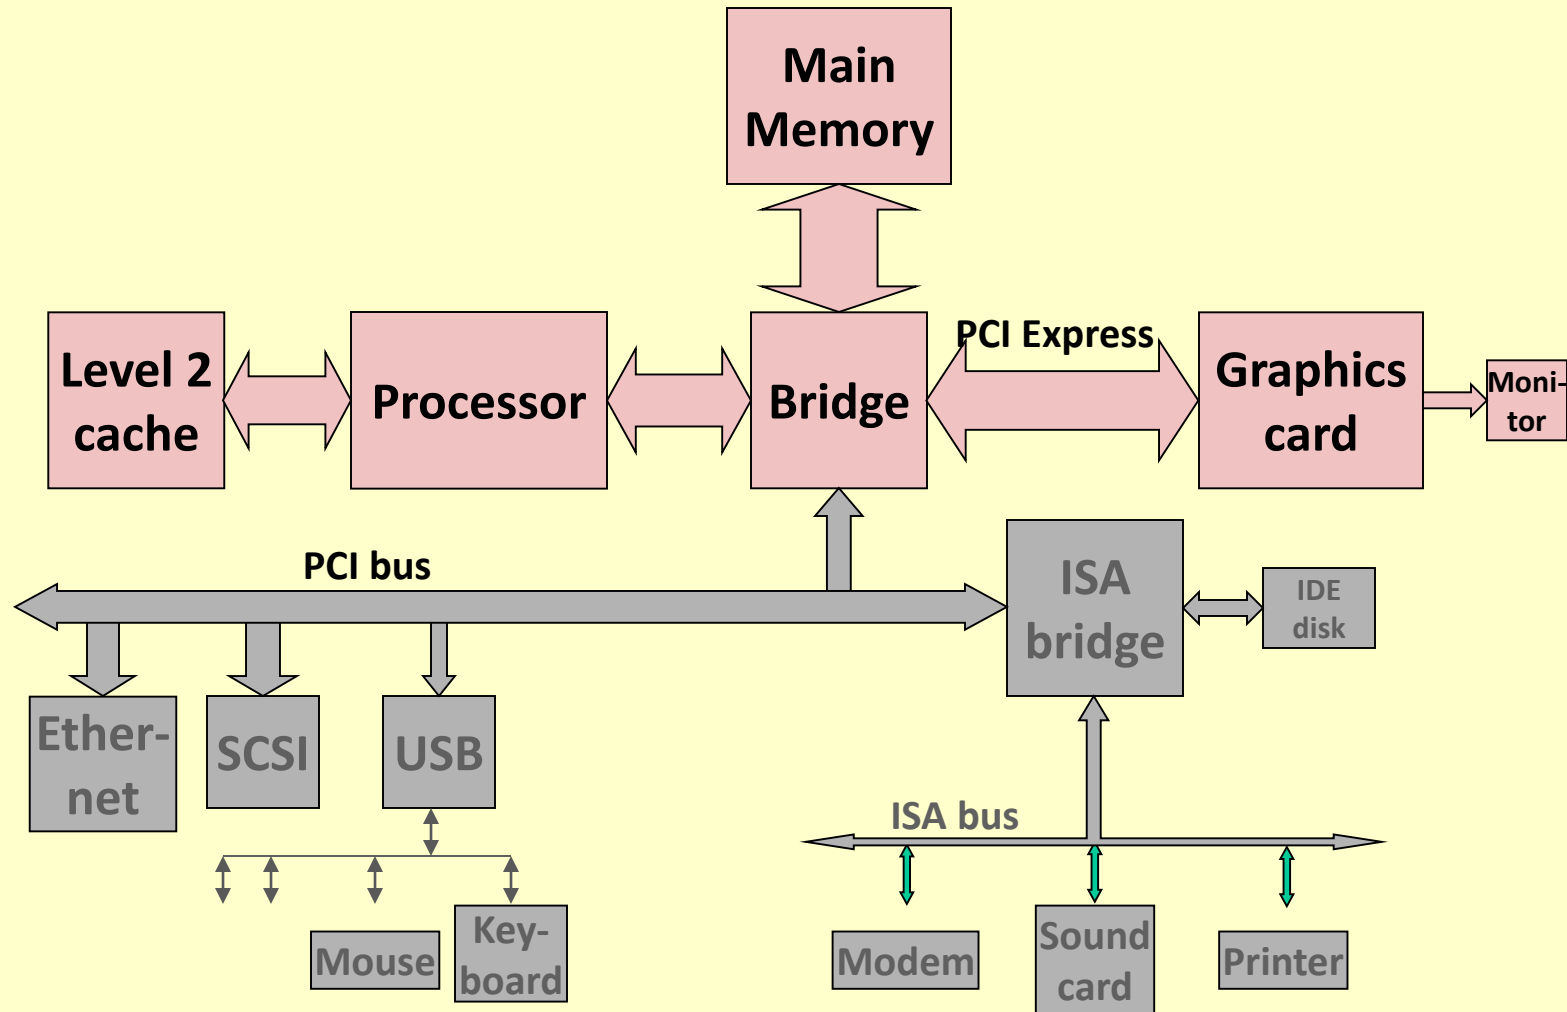
```
graph TD; A[Called by application to add/change information] --> B[write()]; C[Called by window manager to redraw current contents] --> D[refresh()];
```

Called by window manager to
redraw current contents


GUI Device — Text Window

- **Character terminal emulated in a window**
 - RS-232 character set and controls
 - */dev/tty*
- **Operates like a character terminal with visible, partially obscured, or completely covered**

Modern GUI devices



Modern GUI Devices (continued)



- **Double-buffered bitmap in Graphics card**
 - Graphics and information written/drawn in *back* buffer
 - Monitor refreshes from main buffer (60+ Hz)
- ***Refresh* interrupt at start of every frame**
 - *Bitblt* to substitute cursor
- **Processor writes text, etc.**
- **Graphics processor (GPU) draws images, vectors, polygons**
- **Window manager orders redraw when necessary**

Questions?

Reading Assignment

- OSTEP §37-39