

# Introduction to Concurrency (Processes, Threads, Interrupts, etc.)

Professor Hugh C. Lauer  
CS-3013, Operating Systems

(Slides include copyright materials from *Operating Systems: Three Easy Steps*, by Remzi and Andrea Arpaci-Dusseau, from *Modern Operating Systems*, by Andrew S. Tanenbaum, 3<sup>rd</sup> edition, and from other sources)

# Concurrency — things happening at the same time

- Since the beginning of computing, management of concurrent activity has been a *central* issue
- Concurrency between computation and input or output
- Concurrency between computation and user
- Concurrency between independent activities that take place at same time
- Concurrency between parts of large computations that are divided up to improve performance
- ...

# Early 1960s

- Programmers tried to write programs that would read from input devices and write to output devices in parallel with computing

- Card readers, paper tape, line printers, etc.

- Challenges

- Keeping the *buffers* organized
- Synchronizing between I/O and computation
- Computation getting ahead of I/O
- I/O activity getting ahead of computation

**Definition:** *buffer* – a region of memory from which an I/O device gets data or into which an I/O device puts data

# Late 1960s — Shared Computing Services

- **Multiple simultaneous, independent users of large computing facilities**
  - E.g., *Time Sharing* systems of university computing centers
- **Data centers of large enterprises**
  - Multiple accounting, administrative, and data processing activities over common databases
- ...

# Modern Workstations and PCs

- **Multiple windows in personal computer doing completely independent things**
  - *Word, Excel, Photoshop, E-mail, music, etc.*
- **Multiple activities within one application**
  - E.g., in Microsoft *Word*
    - Reading and interpreting keystrokes
    - Displaying what you typed
    - Formatting line and page breaks
    - Spell checking
    - Hyphenation
    - ....

Impossible to do all of these things in one single-threaded program and still get the performance and responsiveness expected by users.

# Modern Game Implementations

- **Multiple characters in game**
  - Concurrently & independently active
- **Multiple constraints, challenges, and interactions among characters**
- **Multiple players**

# Traditional Challenge for OS

- **Useful set of abstractions that help to**
  - Manage concurrency
  - Manage synchronization among concurrent activities
  - Communicate information in useful way among concurrent activities
  - Do it all efficiently

# Technological Pressure

- From early 1950s to early 2000s, single processor computers increased in speed by  $2\times$  every 18 months or so
  - Moore's Law
- Multiprocessing was somewhat of a *niche* problem
  - I.e., computing systems with more than one CPU
  - Specialized computing centers, techniques



# Technological Pressure (continued)

- No longer!
- Modern microprocessor clock speeds are no longer increasing as predicted by Moore's Law
- Microprocessor density on chips still is!

*Sort of!*

⇒ multi-threaded and multi-core processors are now *de facto* standard

- Even on low-end PCs!

# Modern Challenge

- **Methods and abstractions to help software engineers and application designers ...**
  - Take advantage of inherent concurrency in modern application systems
  - Exploit multi-processor and multi-core architectures that are becoming ubiquitous
  - Do so with relative ease

# Fundamental Abstraction

- **Process**

Ostep, Chapter 4

- ... *aka* **Task**

- ... *aka* **Thread**

- ... *aka* **Job**

- ... *aka* [other terms]

# Definition

- ***Process (generic):— A particular execution of a particular program.***
  - Requires time, space, and (perhaps) other resources
- **Separate from all other executions of the same program**
  - Even those at the same time!
- **Separate from executions of other programs**

# ***Process*** (continued)

## ■ Can be

- Interrupted
- Suspended
- Blocked
- Unblocked
- Started or continued

## ■ Fundamental *abstraction* of all modern operating systems

## ■ Note: “Process” in Unix, Linux, and Windows is a heavyweight concept with more implications than this simple definition

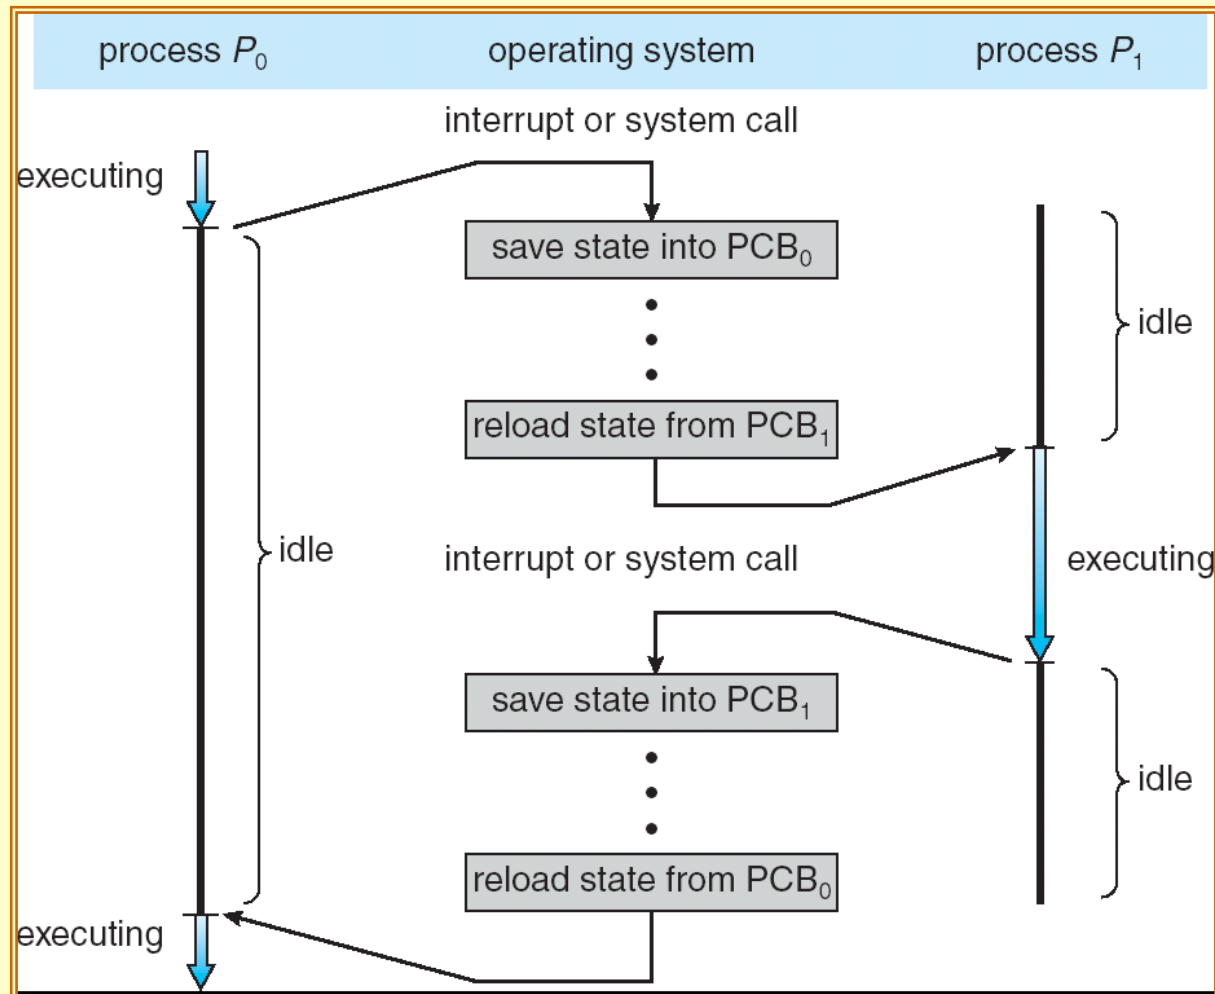
# Process (a generic term – continued)

- Concept emerged and evolved in 1960s
- Intended to make sense out of mish-mash of concurrent programming techniques that bedeviled software engineers
- Analogous to police or taxi dispatcher!

# Background – *Interrupts*

- A mechanism in (nearly) all computers by which a running program can be suspended in order to cause processor to do something else
- Two kinds:–
  - *Traps* – synchronous, caused by running program
    - Deliberate: e.g., system call
    - Error: divide by zero
  - *Interrupts* – asynchronous, spawned by some other concurrent activity or device.
- Essential to the usefulness of computing systems

# Switching from process to process





# Hardware Interrupt Mechanism

## ■ Upon receipt of electronic signal, the processor

- Saves current PSW to a fixed location
- Loads new PSW from another fixed location

## ■ Definition: PSW — *Program Status Word*

- *Program counter*
- Condition code bits (comparison results)
- Interrupt enable/disable bits
- Other control and mode information
  - E.g., privilege level, access to special instructions, etc.

## ■ Occurs *between* machine instructions

- An abstraction in modern processors (see OSTEP, §34.6)

# Interrupt handler

```
/* Enter with interrupts disabled */
```

```
Save registers & state of interrupted computation
```

```
Load registers & state needed by handler
```

```
Examine cause of interrupt
```

```
Take appropriate action (brief)
```

```
Reload registers & state of interrupted computation
```

```
Reload interrupted PSW and re-enable interrupts
```

or

```
Load registers & state of another computation
```

```
Load its PSW and re-enable interrupts
```

# Requirements of interrupt handlers

- Fast
- Avoid possibilities of interminable waits
- Must not count on correctness of interrupted computation
- Must not get confused by multiple interrupts in close succession
- ...
- More challenging on multiprocessor systems

# Result

- **Interrupts make it possible to support concurrent activities**
  - Even on machines with only one processor
- **Don't help in establishing some kind of orderly way of thinking**
- **Need something more**

# Result (continued)

- **Hence, emergence of generic concept of *process***
  - (or whatever it is called in a particular operating system and environment)
- **Notion of *process* allows us to abstract interrupts and interleaving so that we can concentrate on each executing program separately**

# Information needed to implement processes

## ■ PSW (program status word)

- Program counter
- Condition codes
- Control information – e.g., privilege level, priority, etc

## ■ Registers, stack pointer, etc.

- Whatever hardware resources needed to compute

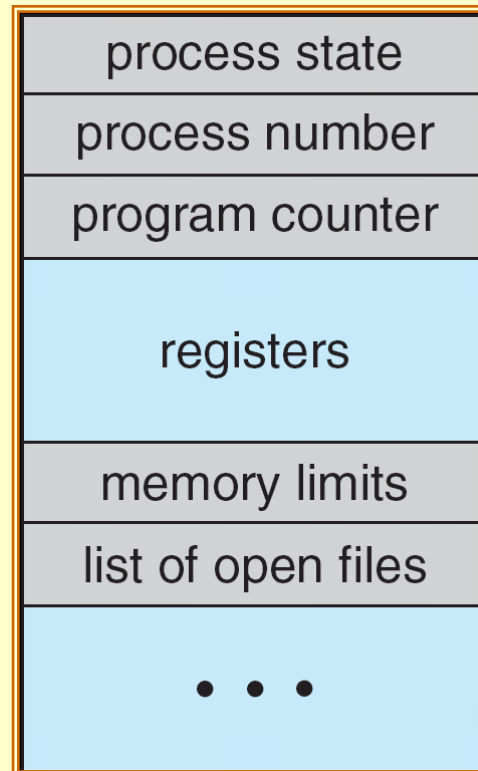
## ■ Administrative information for OS

- Owner, restrictions, resources, etc.

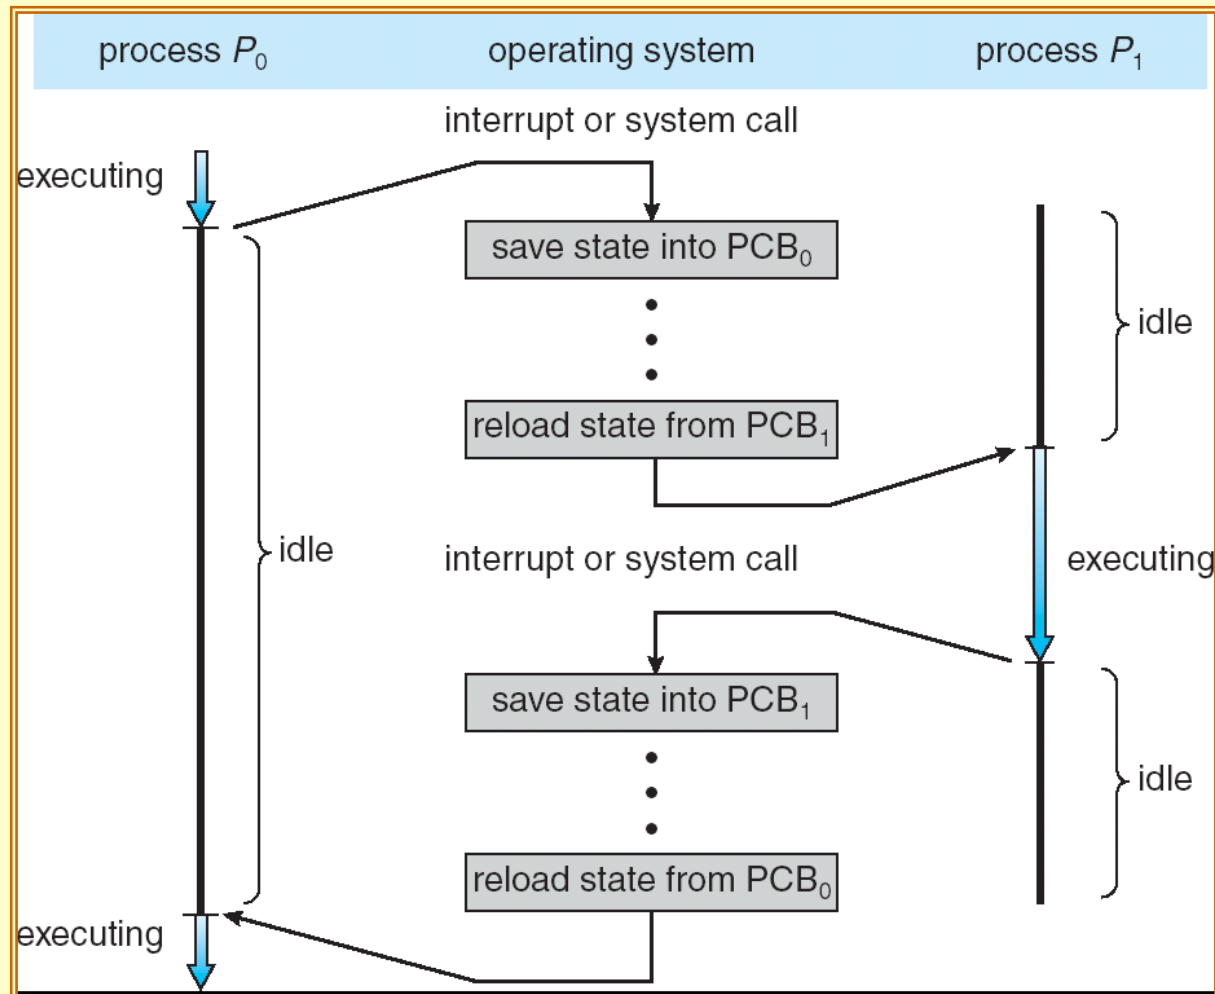
## ■ Other stuff ...

# Process Control Block (PCB)

## (example data structure in an OS)



# Switching from process to process

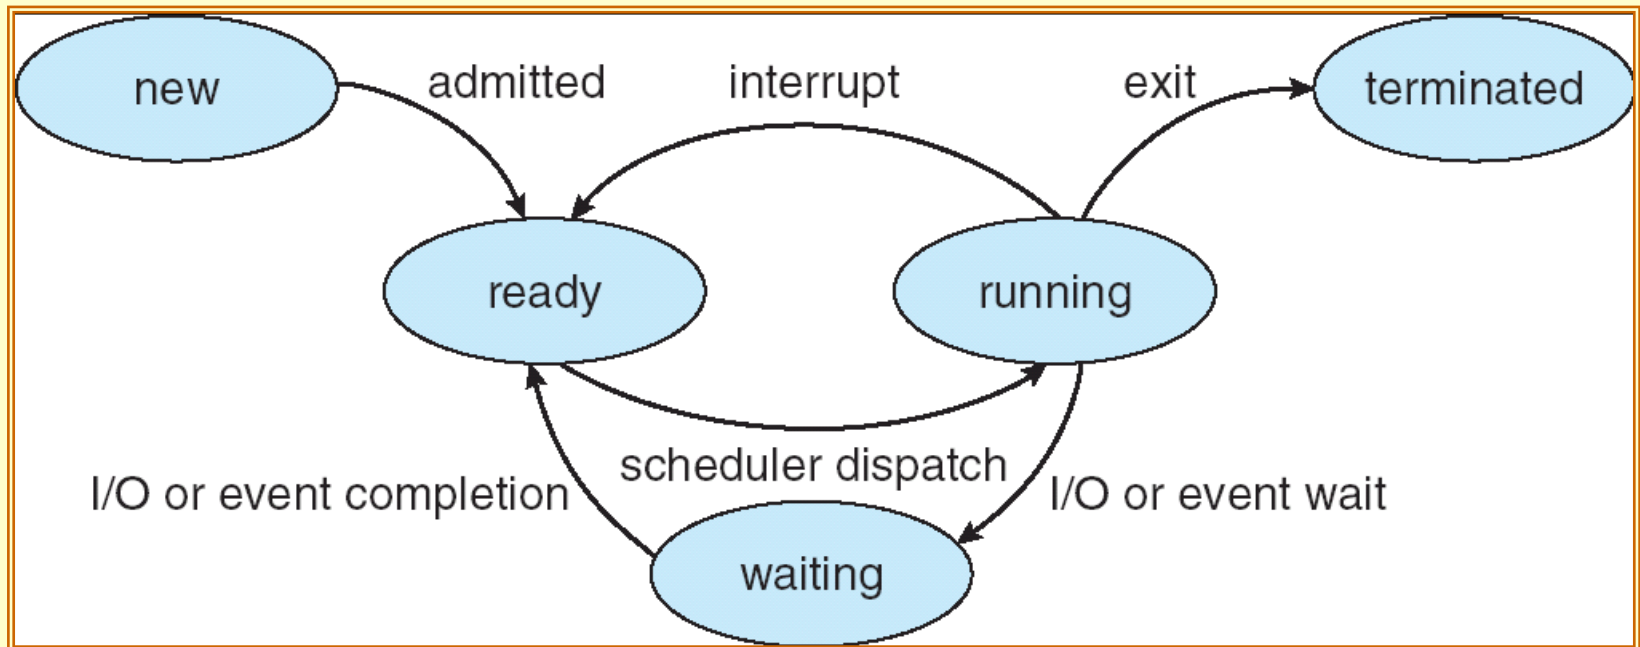




# Result

- A very clean way of thinking about separate computations
- Processes can *appear* be executing in parallel
  - Even on a single processor machine
- Processes really *can* execute in parallel
  - Multi-processor, multi-core, or multi-threaded hardware
- ...

# Process States



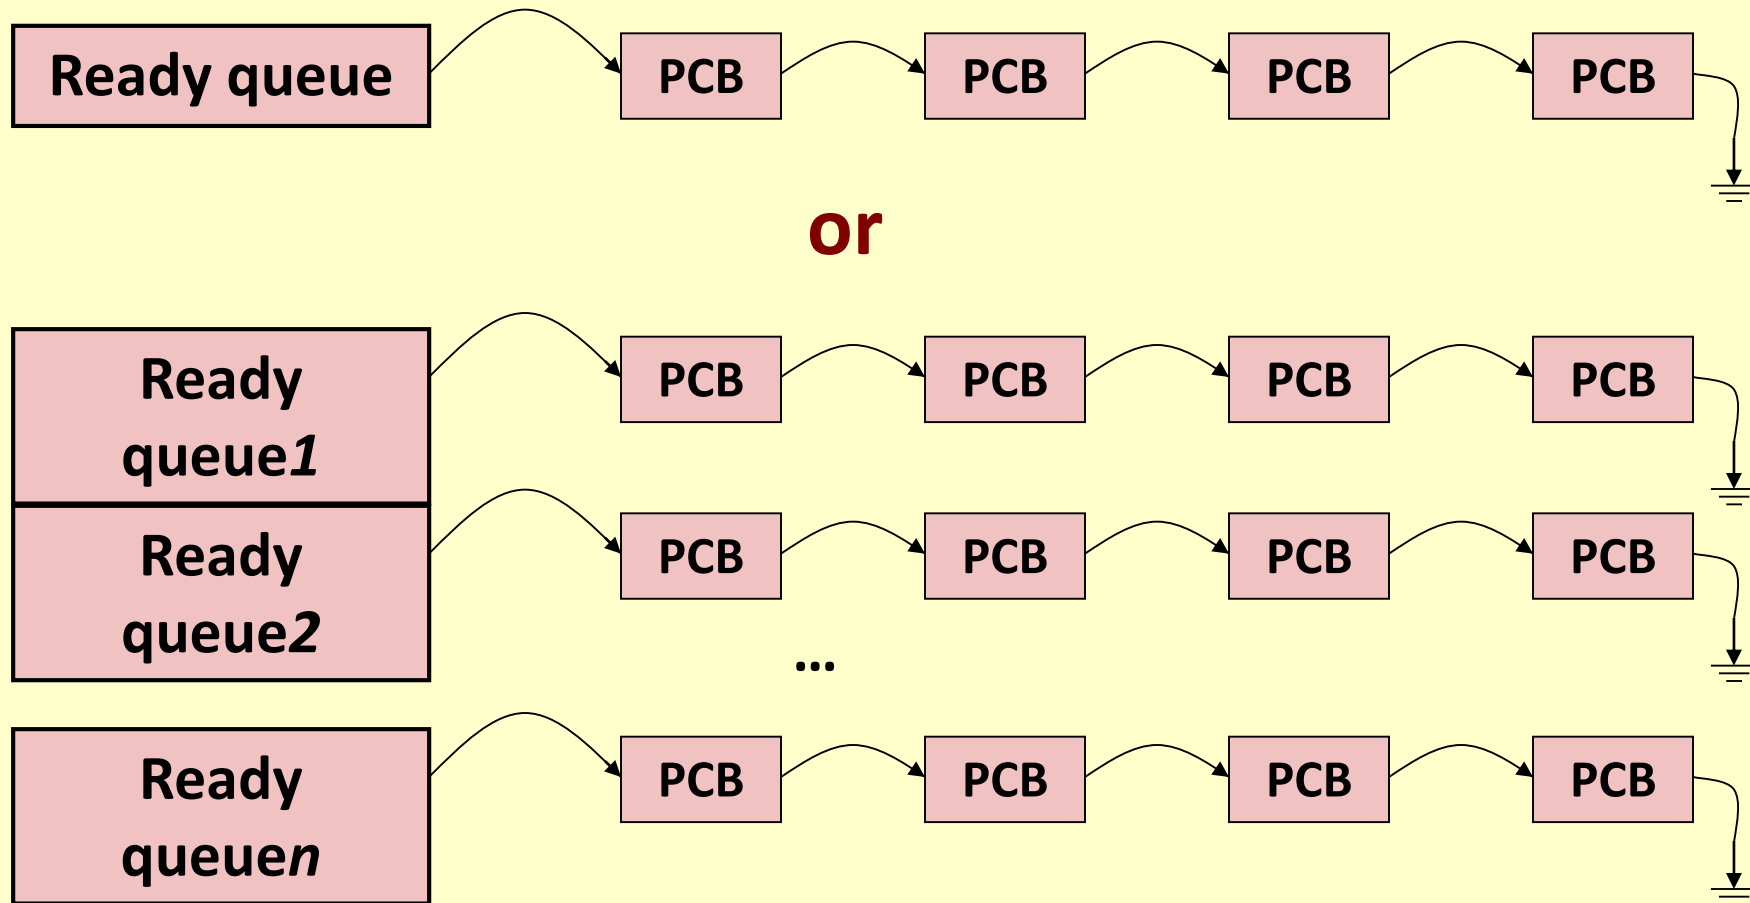
# The Fundamental Abstraction of the OS

- Each process has its own “virtual” processor
- Each process can be thought of as an independent computation
  - Decoupled physical processors from the running of programs!
- On a fast enough physical processor, processes can look like they are really running concurrently
- On multi-processor or multi-core systems, processes really *do* run concurrently

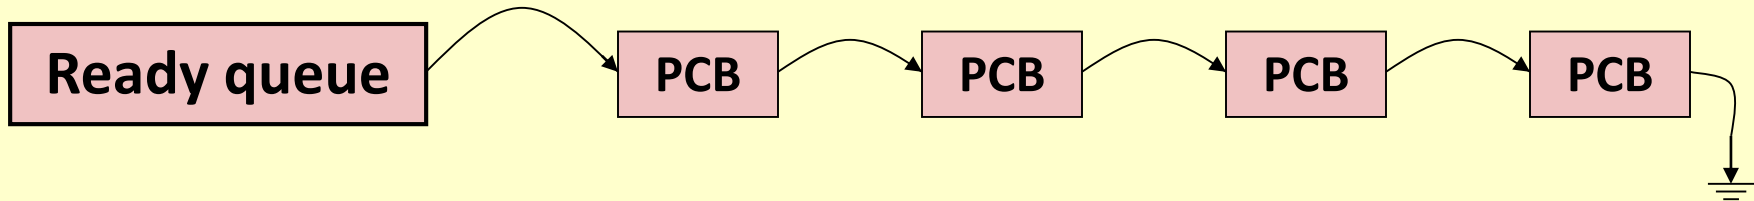
# Questions?



# Implementation of Processes



# Implementation



## ■ Action – *dispatch* a process to CPU

- Remove first PCB from ready queue
- Load hardware registers and PSW
- *Return from interrupt* or trap

## ■ Action – *interrupt* a process

- Save PSW and hardware registers in PCB
- If not blocked, insert PCB back into *ReadyQueue* (in some order); otherwise, link it to some other queue or list
- Take appropriate action
- Dispatch same or another process from *ReadyQueue*

# Timer interrupts

- Can be used to enforce “fair sharing”
- Current process goes to end of *ReadyQueue*
  - *After* other processes of equal or higher priority
- Simulates concurrent execution of multiple processes on same processor

# Processes – Switching

- When a process is *running*, its hardware state is in the processor – PC, processor registers, etc.
- When the OS *suspends* running a process, it saves the hardware state in the PCB
- When the OS *dispatches* a process, it restores the hardware state from the PCB



# Definition – *Context Switch*

- **The act of switching a processor from one process to another**
  - E.g., upon interrupt or some kind of wait for event
- **Not a big deal in simple systems and processors**
- ***Very big deal* in large systems such as**
  - Linux and Windows
  - Pentium, Core i7, etc.

**Many microseconds!**

# Definition — *Scheduling*

- The art and science of deciding *which* process to dispatch next ...
- ... and for how long ...
- ... and on which processor

*Topic for later in this course*

# Questions?

Next Topic – Processes in  
Unix, Linux, and Windows