

Processes in Unix, Linux, and Windows

Professor Hugh C. Lauer
CS-3013, Operating Systems

(Slides include copyright materials from *Operating Systems: Three Easy Steps*, by Remzi and Andrea Arpaci-Dusseau, from *Modern Operating Systems*, by Andrew S. Tanenbaum, 3rd edition, and from other sources)

Processes in Unix, Linux, and Windows

- In previous topic, we used “*process*” in a generic way to represent the abstraction of concurrency
- Unix pre-empted generic term “*process*” to mean something very specific
- Linux, Windows, and other OSes adopted Unix definition

Process in Unix-Linux-Windows comprises

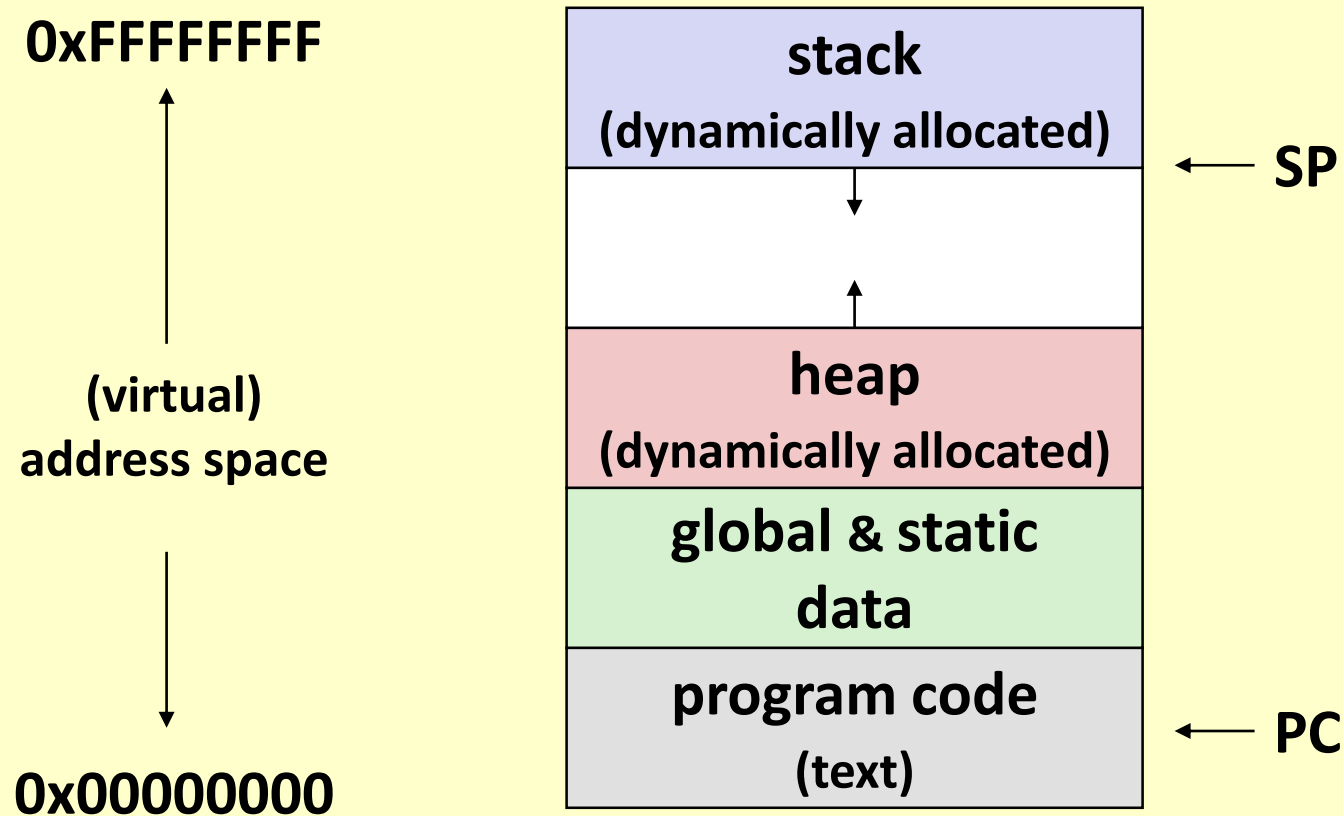
- an *address space* – usually protected and virtual – mapped into memory
- the *code* for the running program
- the *data* for the running program
- an *execution stack* and *stack pointer* (SP); also *heap*
- the *program counter* (PC)
- a set of processor *registers* – general purpose and status
- a set of system *resources*
 - files, network connections, pipes, ...
 - privileges, (human) user association, ...
- ...

Reading Assignment

- **OSTEP**
 - Chapters 4, 5, and 6

- **(optional) R. Love, *Linux Kernel Development***
 - Chapter 3, thru p. 33

Process Address Space (traditional Unix)



Processes in the OS – Representation

- To users (and other processes) a process is identified by its *Process ID* (PID)
- In the OS, processes are represented by entries in a *Process Table* (PT)
 - PID is index to (or pointer to) a PT entry
 - PT entry = *Process Control Block* (PCB)
- PCB is a large data structure that contains (or points to) all info about the process
 - Linux – defined in `task_struct` (over 70 fields)
 - Some of which point to other data structures
 - See `include/linux/sched.h`
 - Windows XP – defined in *EPROCESS* – about 60 fields

Processes in the OS – PCB

■ Typical PCB contains:

- execution state
- PC, SP & processor registers – stored when process is not in *running* state
- memory management info
- privileges and owner info
- scheduling priority
- resource info
- accounting info

Process – starting and ending

■ Processes are created ...

- When the system boots
- By the actions of another process (more later)
- By the actions of a user (via another process)

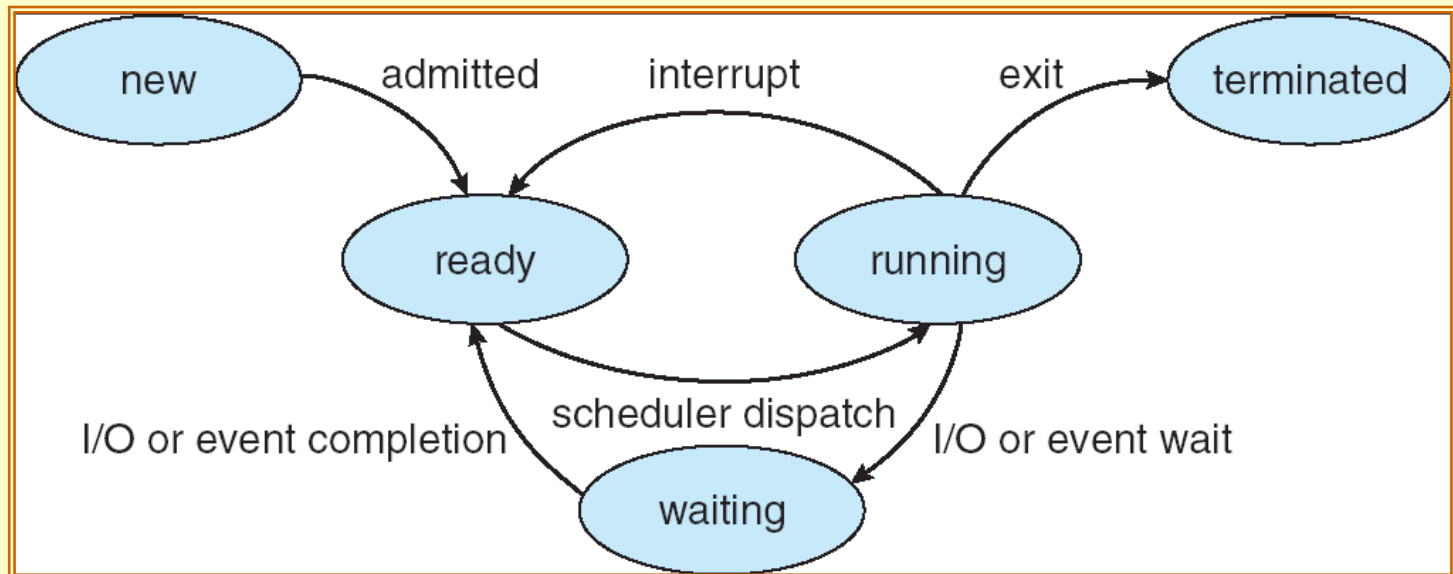
■ Processes terminate ...

- Normally – exit
- Voluntarily on an error
- Involuntarily on an error
- Terminated (killed) by action of
 - a user or
 - another process

Processes – States

■ Process has an execution state

- *ready*: waiting to be assigned to CPU
- *running*: executing on the CPU
- *waiting*: waiting for an event, e.g. I/O



Processes – State Queues

- The OS maintains a collection of *process state* queues
 - typically one queue for each state – e.g., ready, waiting, ...
 - each PCB is put onto a queue according to its current state
 - as a process changes state, its PCB is unlinked from one queue, and linked to another
- Process state and the queues change in response to events – interrupts, traps



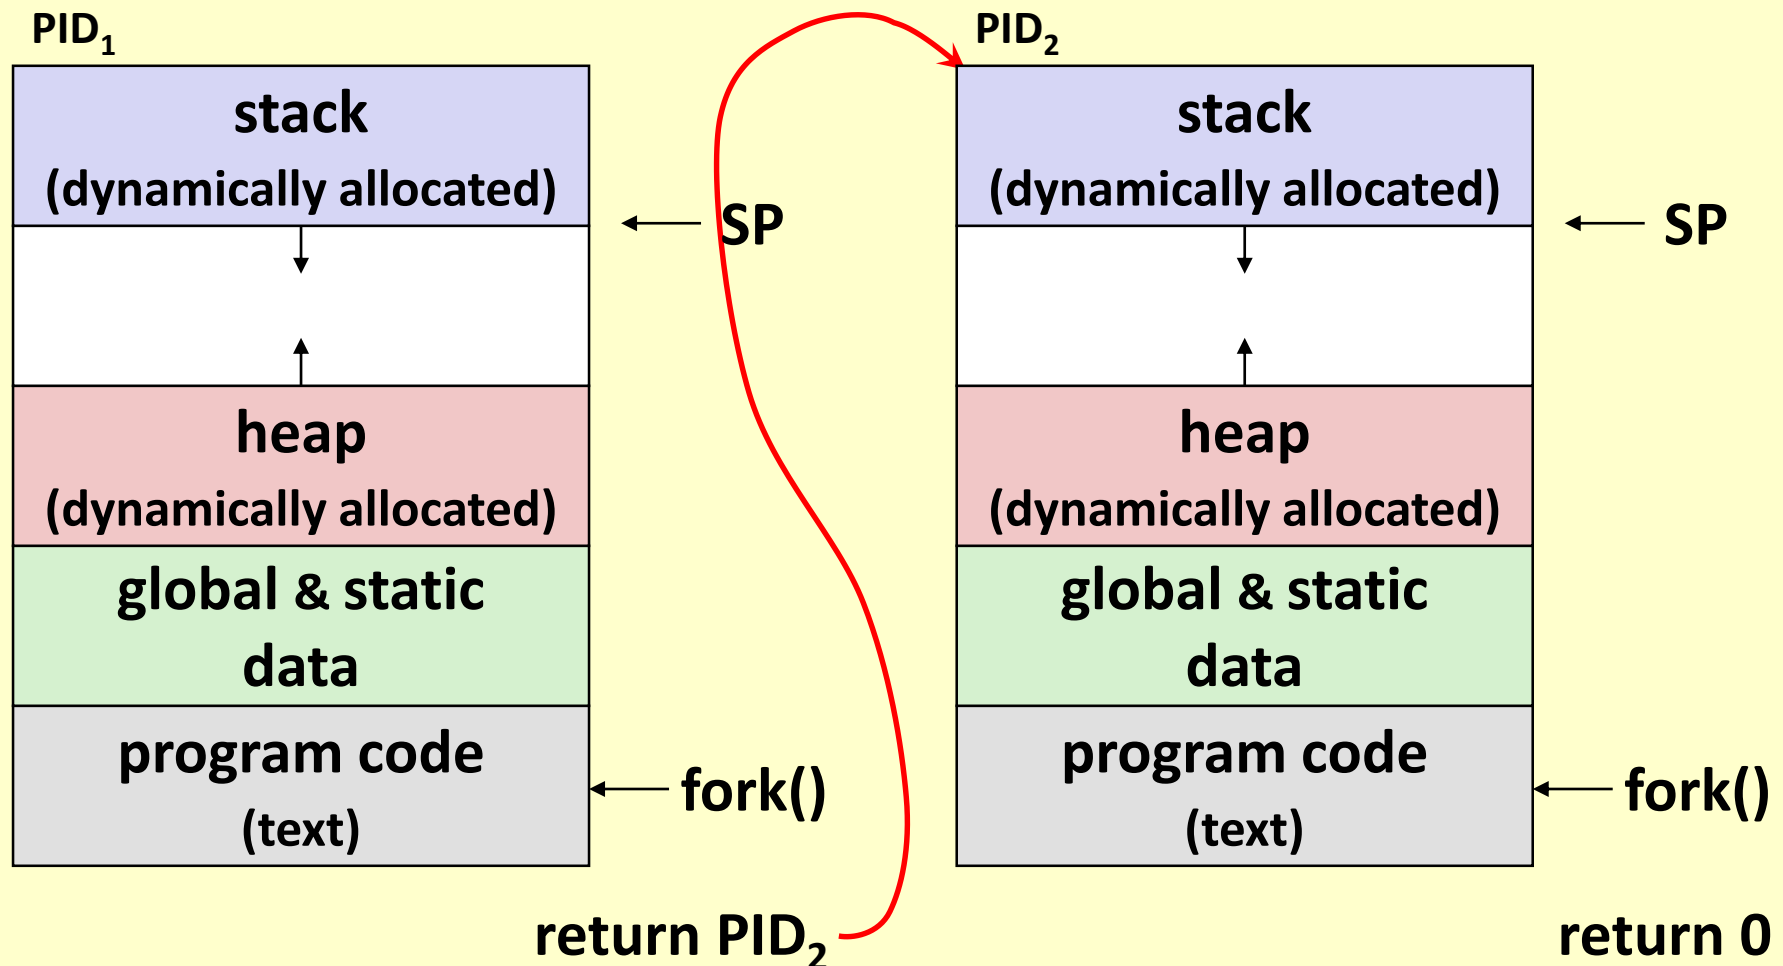
Processes – Privileges

- Users are given privileges by the system administrator
- Privileges determine what *rights* a user has for an *object*.
 - Unix/Linux – Read | Write | eXecute by user, group and “other” (i.e., “world”)
 - Windows NT/XP/7 – Access Control List
- Processes “inherit” privileges from user
 - or from creating process

Process Creation – Unix & Linux

- **Create a new (child) process – `fork()` ;**
 - Allocates new PCB
 - Clones the calling process (almost exactly)
 - Copy of calling process address space
 - Copy of all registers, condition codes, etc.
 - Copies resources in kernel (e.g. pointers to files)
 - Places new PCB on *Ready queue*
 - Return from **`fork()`** call
 - 0 for child
 - child PID for parent

Process Creation – Unix & Linux



Example of *fork()*

```
int main(int argc, char **argv)
{
    char *name = argv[0];

    int child_pid = fork();

    if (child_pid == 0) {
        printf("Child of %s sees PID of %d\n",
              name, child_pid);
        return 0;
    } else {
        printf("I am the parent %s. My child is %d\n",
              name, child_pid);
        return 0;
    }
}
```

```
% ./forktest
Child of forktest sees PID of 0
I am the parent forktest. My child is 486
```

Result – Two identical processes

Parent

```
int main(int argc, char **argv)
{
    char *name = argv[0];
    int child_pid = fork();
    if (child_pid == 0) {
        printf("Child of %s sees PID"
              " %d\n", name, child_pid);
        return 0;
    } else {
        printf("I am the parent %s. "
              "My child is %d\n", name,
              child_pid);
        return 0;
    }
}
```

False

Child

```
int main(int argc, char **argv)
{
    char *name = argv[0];
    int child_pid = fork();
    if (child_pid == 0) {
        printf("Child of %s sees PID"
              " %d\n", name, child_pid);
        return 0;
    } else {
        printf("I am the parent %s. "
              "My child is %d\n", name,
              child_pid);
        return 0;
    }
}
```

True

Only difference

Questions?

Starting New Programs

■ Unix & Linux:—

- `int exec (char *prog, char **argv)`
- Check privileges and file type
- Loads program at path *prog* into address space
 - Replacing previous contents of address space!
 - Execution starts as function call to `main()`
- Initializes context – e.g. passes arguments
 - `*argv`
- Place PCB on *ready queue*
- Preserves, pipes, open files, privileges, etc.

Executing a New Program (Linux-Unix)

- **`fork()` followed by `exec()`**
- **Creates a new process as clone of previous one**
 - I.e., same program, but different execution of it
- **First thing that clone does is to replace itself with new program**

Fork + Exec – shell-like

```
int main(int argc, char **argv)
{ char *argvNew[5];
  int pid;
  if ((pid = fork()) < 0) {
    printf( "Fork error\n");
    exit(1);
  } else if (pid == 0) { /* child process */
    argvNew[0] = "/bin/ls"; /* i.e., the new program */
    argvNew[1] = "-l";
    argvNew[2] = NULL;
    if (execvp(argvNew[0], argvNew) < 0) {
      printf( "Execvp error\n");
      exit(1); /* program should not reach this point */
    }
  } else { /* parent */
    wait(pid); /* wait for the child to finish */
  }
}
```

Waiting for a Process

■ Multiple variations of *wait* function

- Including non-blocking *wait* functions

■ Waits until child process terminates

- Acquires termination code from child
- Child process is destroyed by kernel

■ ***Zombie***:— a process that had never been *waited* for

- Hence, cannot go away!
- See Love, *Linux Kernel Development*, pp 37-38
- See OSTEP, §4.5, p. 32

Processes – Windows

- **Windows NT/XP/10 – combine `fork` & `exec`**
 - `CreateProcess` (10 arguments)
 - No parent child relationship
 - *Note* – privileges are required to create a new process
- Much more about processes, threads, etc., in Windows than we can cover in this course

Traditional Unix

- ***Processes are in separate* address spaces**
 - By default, no shared memory
- ***Processes are unit of scheduling***
 - A process is *ready, waiting, or running*
- ***Processes are unit of resource allocation***
 - Files, I/O, memory, privileges, ...
- ***Processes are used for (almost) everything!***

Windows and Linux

- *Threads* (next topic) are units of scheduling
- Threads are used for everything

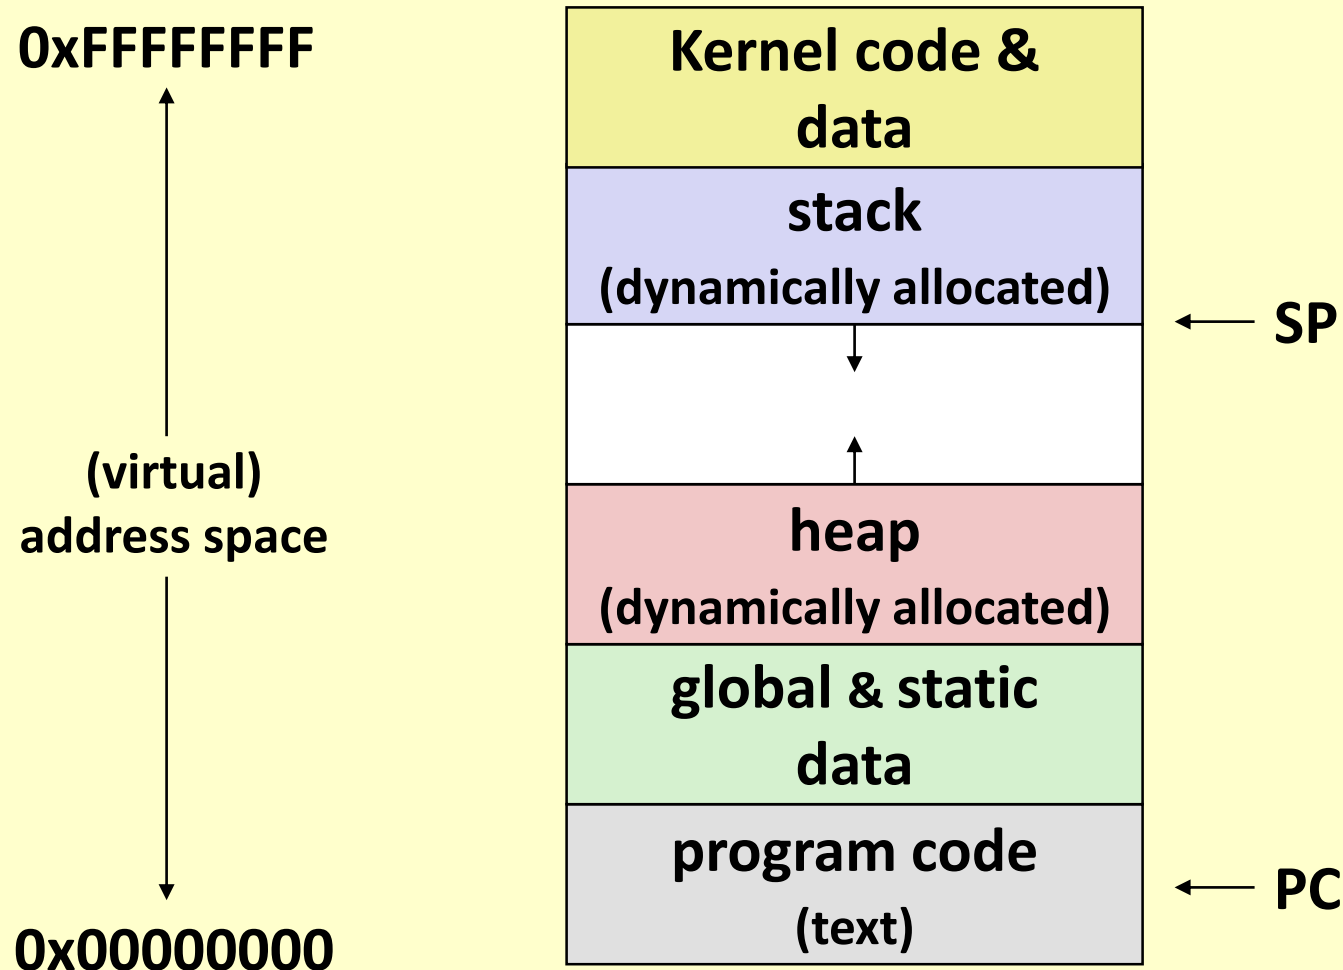
Non-Traditional Unix — e.g., iPhone, Android, etc.

- **Processes are in *separate* address spaces**
 - By default, no shared memory
- **Processes are unit of scheduling**
 - A process is *ready*, *waiting*, or *running*
- **Processes are unit of resource allocation**
 - Files, I/O, memory, privileges, ...
- **Processes are used for (almost) everything!**

A Note on Implementation

- Many OS implementations include (parts of) *kernel* in every address space
 - Protected
 - Easy to access
 - Allows *kernel* to see into client processes
 - Transferring data
 - Examining state
 - ...

Process Address Space (with kernel)



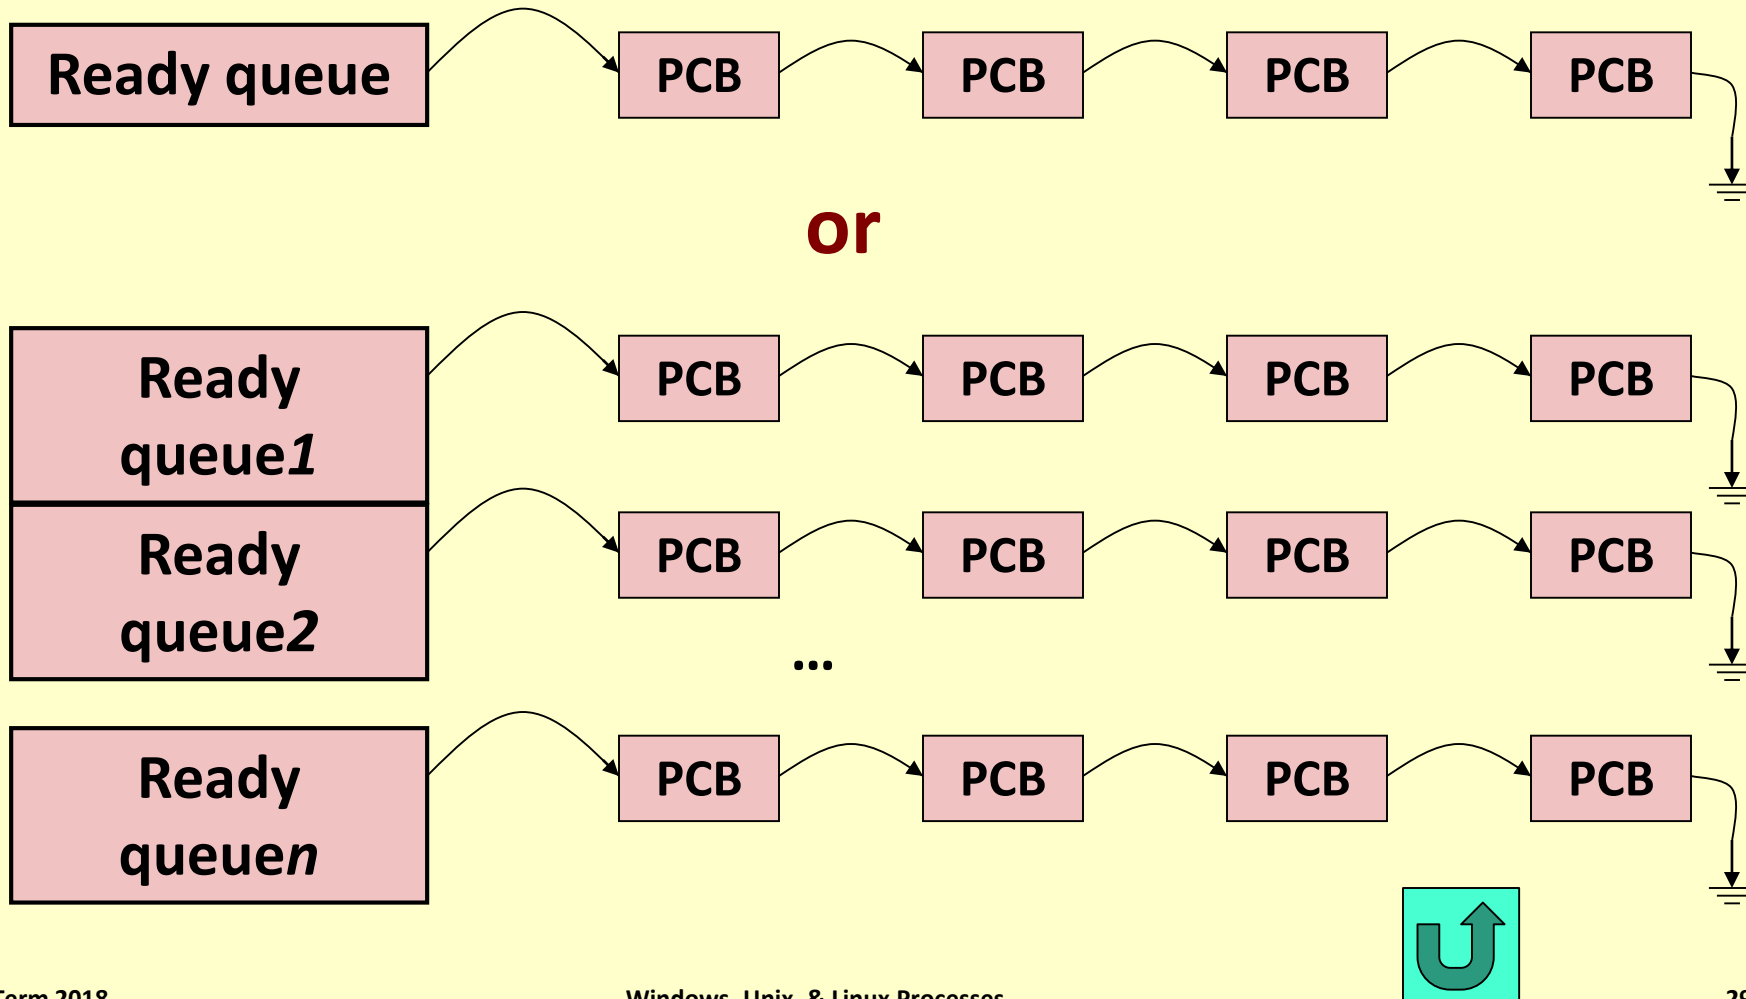
Linux Kernel Implementation

- Kernel may execute in either *Process context* vs. *Interrupt context*
- In *Process context*, kernel has access to
 - Virtual memory, files, other process resources
 - May sleep, take page faults, etc., on behalf of process
- In *Interrupt context*, no assumption about what process was executing (if any)
 - No access to virtual memory, files, resources
 - May not sleep, take page faults, etc.

Processes in Other Operating Systems

- Implementations will differ
- Sometimes a subset of *Unix/Linux/Windows*
Sometimes quite different
- May have more restricted set of resources
- Often, specialize in real-time constraints

Implementation of Processes



Reading Assignment

- **OSTEP, Chapters 4, 5, and 6**

Questions?