

Threads

Professor Hugh C. Lauer
CS-3013, Operating Systems

(Slides include copyright materials from *Operating Systems: Three Easy Steps*, by Remzi and Andrea Arpaci-Dusseau, from *Modern Operating Systems*, by Andrew S. Tanenbaum, 3rd edition, and from other sources)

Review

■ Process (the generic concept)

- A particular execution of a program
- Separate from all other executions of that same program
- Separate from all executions of other programs

Review — Processes in Unix, Linux, and Windows

- a particular execution of a program PLUS
- an *address space* – usually protected and virtual – mapped into memory
- the *code* for the running program
- the *data* for the running program
- an *execution stack* and *stack pointer* (SP); also *heap*
- the *program counter* (PC)
- a set of processor *registers* – general purpose and status
- a set of system *resources*
 - files, network connections, pipes, ...
 - privileges, (human) user association, ...
- ...

Processes in Unix, Linux, and Windows (continued)

- Evolutionary direction:– *Isolation*
- I.e., main goal was to keep processes separate from each other
 - Multi-user computer systems (time-sharing)
 - Independent applications
- Communication and cooperation among processes was secondary & limited

Processes in Unix, Linux, and Windows (continued)

■ Processes in Unix, Linux, and Windows are “heavyweight”

- Lots of resources
- Expensive context switches
- Inflexible
- ...

OS-Centric View

■ Lots of data in PCB & other data structures

- Even more when we study memory management
- More than that when we study file systems, etc.

■ Processor caches a lot of information

- Memory Management information
- Caches of active pages

■ Costly context switches and traps

- Many 10s of microseconds (even 100s of microseconds)

Application-Centric View

■ Separate processes have separate address spaces

- Shared memory is limited or nonexistent
- Applications with internal concurrency are difficult

■ Isolation between independent processes vs. cooperating activities

- Fundamentally different goals

Example

■ *Web Server*

- Need to support multiple, concurrent, independent requests pertaining to common data

■ **One solution:—**

- create several processes that execute in parallel
- Use shared memory — **shmget ()** — to map to the *same* address space into multiple processes
- have the OS schedule them in parallel

■ **Clumsy and inefficient**

- *Space and time:* PCB, page tables, cloning entire process, etc.
- *programming:* **shmget ()** is really hard to program!

Example 2

- **Transaction processing systems**
 - E.g, airline reservations or bank ATM transactions
- **1000's of transactions *per second***
 - Very small computation per transaction
 - Long wait times for data base access
- **Separate processes per transaction are too costly**
- **Other techniques (e.g., message passing) are much more complex**

Example 3

- **Games have multiple active characters**
 - Independent behaviors
 - Common context or environment
- **Need “real-time” response to user**
 - For interactive gaming experience
- **Programming all characters in separate processes is really, really hard!**
- **Programming them in a single process is much harder without concurrency support.**

This problem ...

■ ... is partly an artifact of

- Unix, Linux, and Windows

and of

- Big, powerful processors (e.g., Core i7, Xeon, Ryzen)

■ ... partly a consequence of history

- Shared computers rather than personal computers

■ ... tends to occur in most large systems

■ ~~... is infrequent in small scale systems~~

- ~~PDA's, cell phones~~
- ~~Closed systems (i.e., controlled applications)~~

Solution:– *Threads*

- **A *thread* is a particular execution of a program, function, or procedure *within the context* of a Linux or Windows process**

- I.e., a specialization of the concept of *process*

- **A thread has its own**

- Program counter, registers, PSW
 - Stack

- **A thread shares**

- Address space, heap, global and static data, program code
 - Files, privileges, all other resources

with all other threads of the same process

Reading Assignment

- **OSTEP**
 - §25 – §27

- **Robert Love, *Linux Kernel Development***
 - Chapter 3 – “Process Management”

Definition:– *Thread*

- **A *thread* is a particular execution of a program or procedure *within the context* of a Unix, Linux, or Windows process**

- I.e., a specialization of the concept of *process*

- **A thread has its own**

- Program counter, registers, PSW
 - Stack

- **A thread shares**

- Address space, heap, global and static data, program code
 - Files, privileges, all other resources

with all other threads of the same process

Illustration

```
void main(int argc, char** argv);  
char *f(char *c, int i);
```

```
main(...) {
```

```
    int j = ...;  
    char *msg = malloc(...);  
    thread_t T;  
    T = new_thread f(msg, j);  
  
    ...  
    while (...) {  
        /* loop doing something  
           else */  
    }  
  
    ...  
  
    char *result = join T;  
}
```

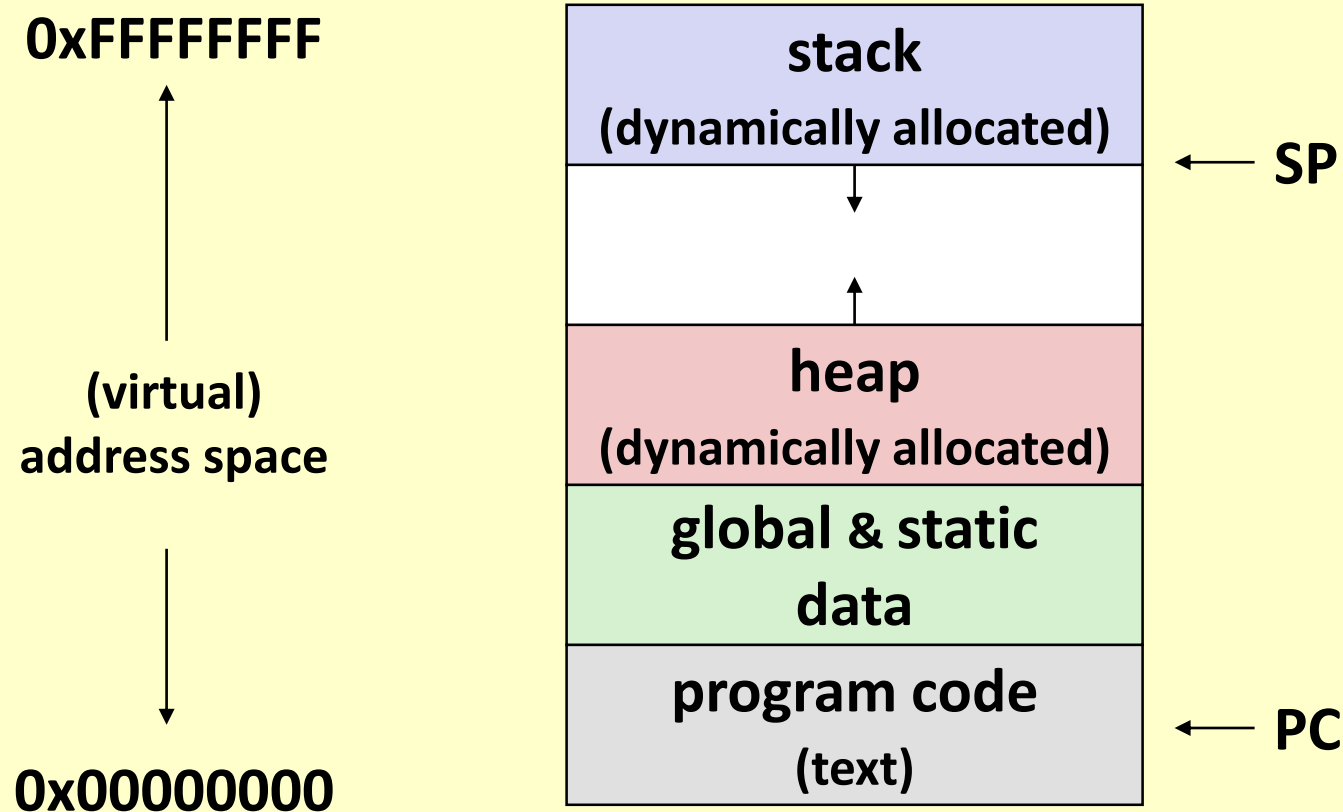
```
f(...) {
```

```
    char *d = malloc(...);  
    for (i=0; i<j; i++) {  
        ...  
        d[i] = c[i];  
    }  
    return d;  
}
```

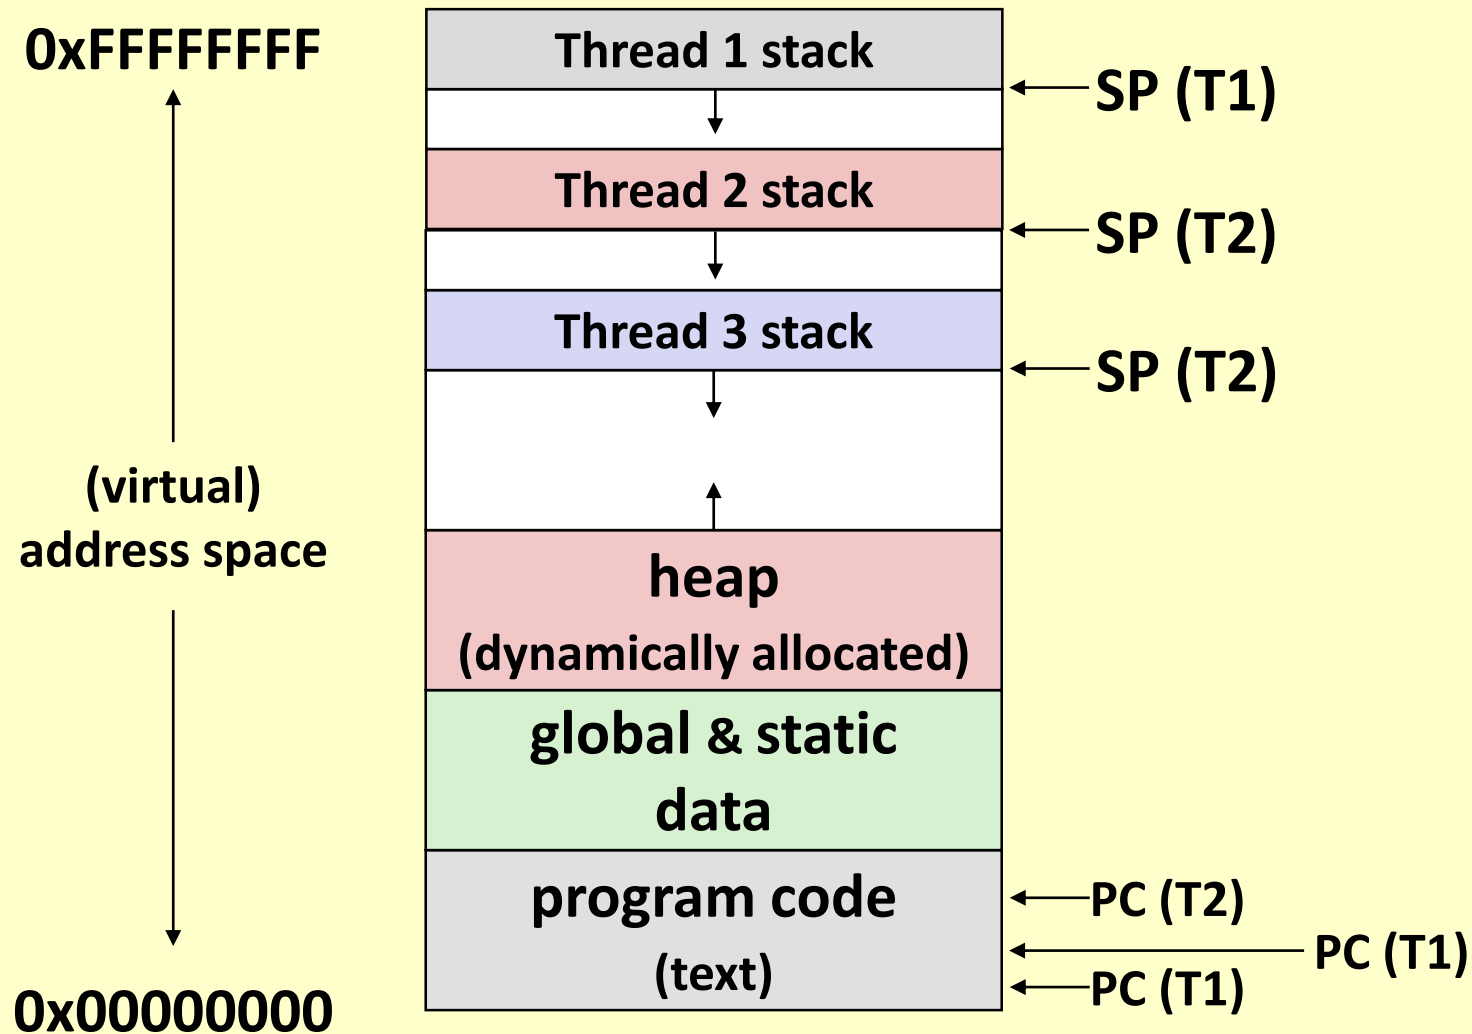
Same heap

From previous lesson

Process Address Space (traditional Linux)



Address Space for Multiple Threads



Basic Thread Functions

■ **Create a thread**

- Make a function call that executes *concurrently* with caller

■ **Exit a thread**

- *I.e.*, thread terminates or returns from its function

■ **Join a thread**

- Wait until the designated thread terminates, capture its return value, delete its stack

■ ...

Basic Thread Functions (continued)

■ ...

■ ***Detach* a thread**

- Separate it from its creators
- Becomes a standalone, independent execution
- Not “*joinable*”; no value returned
- Stack goes away upon termination

■ ***Get own thread ID***

- *I.e.*, get the number of my own thread

■ ...

Another Example

```
// web server

while (true) {
    listen for web request
    allocate private socket
    create & detach new
    thread to serve
    request
    ...
}
```

```
// individual request
// handler

interpret request
while (!done) {
    interact with client
    over private socket
    ...
}
exit thread
```

Benefits

- **Responsiveness**
- **Resource Sharing**
- **Economy**
- **Better utilization of multi-processor architectures than is achievable with just processes**

Using Threads

- **Everyone uses threads nowadays**
 - Many purposes
 - Many languages
- **Programming applications with concurrent activity within them has become an essential skill**
 - I.e., some form of threads

Questions?

Tools for using Threads

- **Three primary thread libraries: –**
 - POSIX *pthread*s
 - Win32 threads
 - Java threads

- Python
- Others?

See the man pages for these functions

Thread Interface

■ E.g., POSIX pthreads API:

- `int pthread_create(pthread_t *thread, const pthread_attr_t *attr, void*(*start_routine)(void), void *arg)`
 - creates a new thread of control
 - new thread begins executing the function `start_routine(arg)`
- `pthread_exit(void *value_ptr)`
 - terminates the calling thread
- `pthread_join(pthread_t thread, void **value_ptr)`
 - blocks the calling thread until the specified thread terminates
- `pthread_t pthread_self()`
 - Returns the calling thread's identifier

Project 3 will use pthreads

Implementing Threads

■ In *User space*

- User-space function library
- Runtime system – similar to process management except in user space
- Windows NT – *fibers*: a user-level thread mechanism

■ In *Kernel*

- Primitive objects known to and scheduled by kernel
- Linux: *lightweight process* (LWP)
- Windows NT, XP, Vista, 7, 10:– *threads*

Implementing Threads

■ In *User space*

Obsolete

- User-space function library
- Runtime system – similar to process management except in user space
- Windows NT – *fibers*: a user-level thread mechanism

■ In *Kernel*

- Primitive objects known to and scheduled by kernel
- Linux: *lightweight process* (LWP)
- Windows NT, XP, Vista, 7:– *threads*

Threads Implemented in User Space

- Thread management done by non-privileged threads library
- Runs in process address space

User-Space Threads (continued)

- **Can be implemented without kernel support**
 - ... or knowledge!
- **Program links with a runtime system that does thread management**
 - Operations are very efficient (function calls)
 - Space efficient and all in user space (TCB)
 - Task switching is very fast
- **Since kernel not aware of threads, there can be scheduling inefficiencies or issues**
 - E.g., blocking I/O calls
 - Non-concurrency of threads on multiple processors

User-Space Threads

- **Obsolete because all modern kernels support threads at kernel level**
- **Still (somewhat of) a performance issue**
- **Research focused on how to avoid system calls for thread synchronization, creation, deletion, etc.**

Questions?

Threads Implemented in Kernel

■ Supported by the Kernel

- OS maintains data structures for thread state and does all of the work of thread implementation.

■ Examples

- Windows XP/2000/Vista/7/10
- Solaris
- Linux version 2.6, etc.
- Tru64 UNIX
- Mac OS X
- ...

Kernel-level Threads (continued)

- OS schedules *threads* instead of *processes*

- **Benefits**

- Overlap I/O and computing in a process
- Creation is cheaper than processes
- Context switch *can* be faster than processes

- **Negatives**

- System calls needed for thread operations
 - High overhead
- Additional OS data space for each thread

Threads Supported by Processor

- **E.g., Intel Hyperthreading™**
 - www.intel.com/products/ht/hyperthreading_more.htm
- **Multiple processor cores on a single chip**
 - True concurrent execution within a single process
- **Requires kernel support**
- **Exposes many issues**
 - Critical section management of synchronization primitives at kernel level
 - Multiple threads scheduling from same ready queue
 - Multiple interrupts in progress at one time

Unix Processes vs. Threads

- On a 700 Mhz Pentium running Linux
 - Processes:
 - `fork()` / `exit()`: 250 microsec
 - Kernel threads:
 - `pthread_create()` / `pthread_join()`: 90 microsec
 - User-level threads:
 - `pthread_create()` / `pthread_join()`: 5 microsec

Some Issues Pertaining to Threads

- **Process global variables**

- E.g., `ERRNO` in Unix — a global variable set by system calls

- **Semantics of `fork()` and `exec()` system calls for processes**

- **Thread cancellation**

- **Signal handling**

- **Thread pools**

- **Thread specific data**

- **Scheduler activations**

Semantics of `fork()` and `exec()`

- Does `fork()` duplicate only the calling thread or all threads?
 - Easy if user-level threads
 - All threads are duplicated, unbeknownst to kernel
 - Not so easy with kernel-level threads
 - Linux has special `clone()` operation
 - Creates new thread in current address space
 - Specifies function to run
 - See **man** page
 - Windows XP/Vista/7/10 has something similar

Thread Cancellation

- **Terminating a thread before it has finished**
- **Two general approaches:**
 - *Asynchronous cancellation* terminates the target thread immediately
 - *Deferred cancellation* allows the target thread to periodically check if it should be cancelled

Modern Linux Threads

- Implemented in kernel
- “A thread is just a special kind of process.”
 - Robert Love, *Linux Kernel Development*, p.23
- ***The* primary unit of scheduling and computation implemented by Linux 2.6 (and later) kernel**
- Every thread has its own `task_struct` in kernel
- ...

Modern Linux Threads (continued)

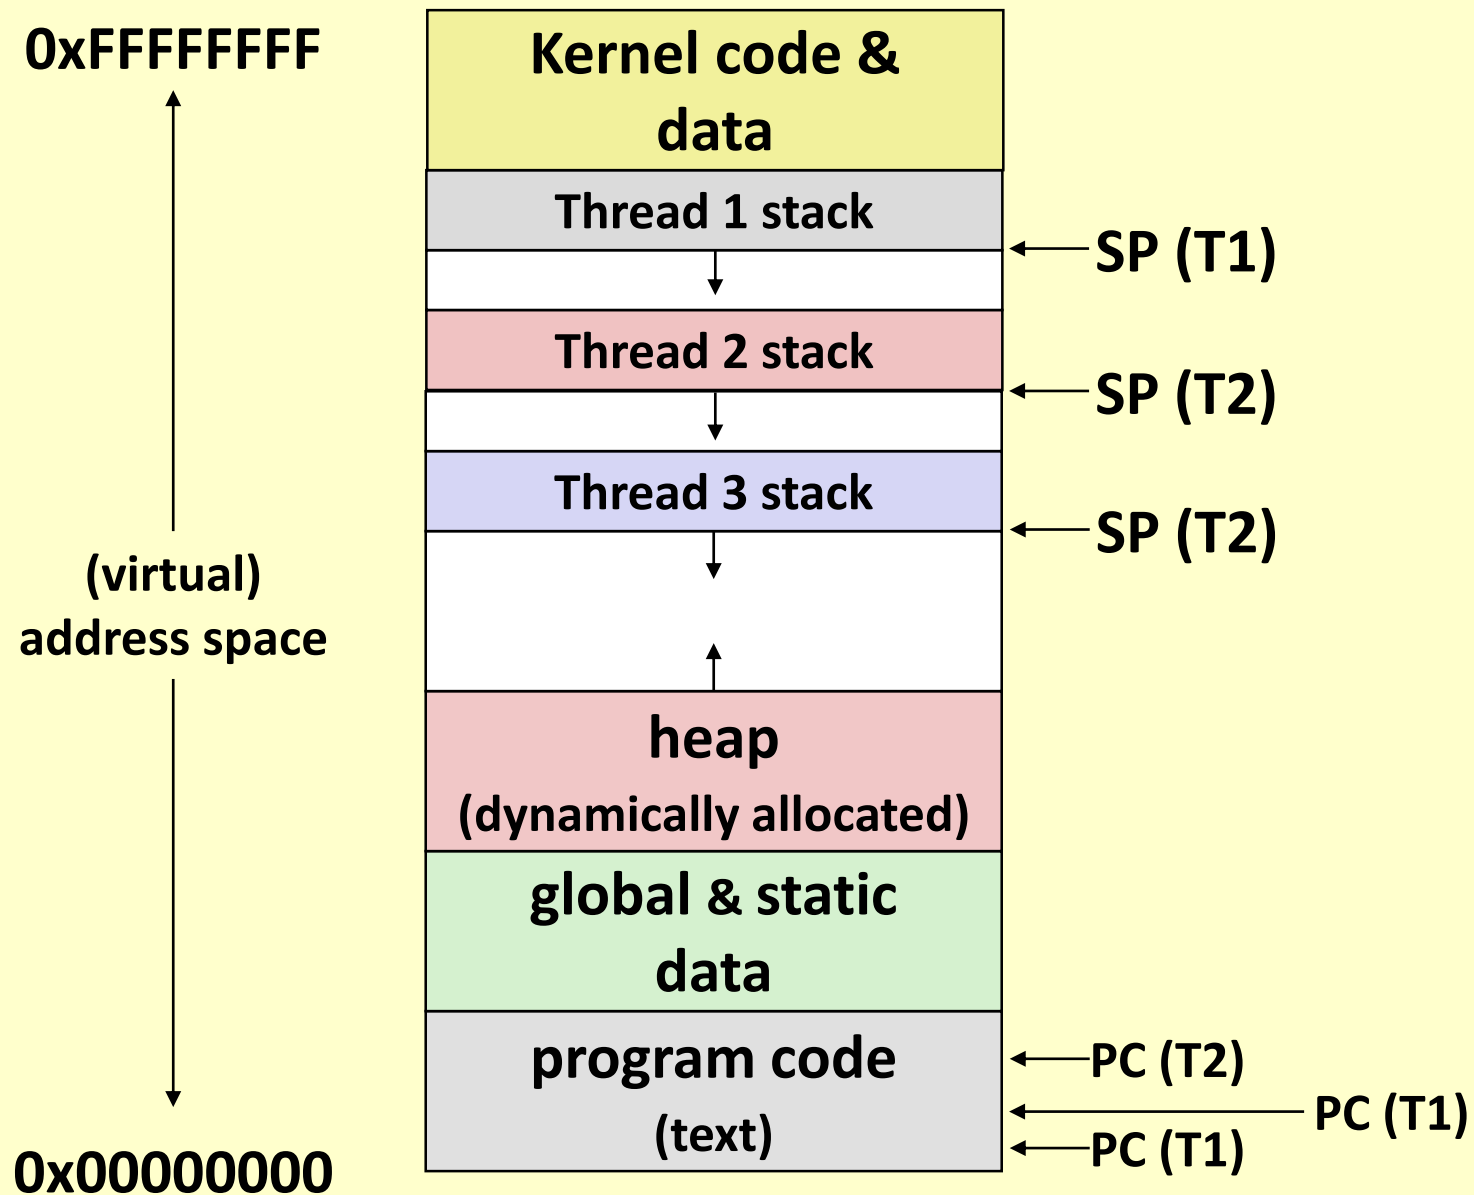
- *Process* `task_struct` has pointers to own memory & resources
- *Thread* `task_struct` has pointer to process's memory & resources
- `fork()` and `thread_create()` are library functions implemented by `clone()` kernel call
- ...

Modern Linux Threads (continued)

- **Threads are scheduled independently of each other**
- **Threads can block independently of each other**
 - Even threads of same process
- **Threads can make their own kernel calls**
 - Kernel maintains a small *kernel stack* per thread
 - During kernel call, kernel is in *process context* (next slides)

Digression – Process Address Space

- Linux includes (parts of) *kernel* in every address space
 - Protected
 - Easy to access
 - Allows *kernel* to see into client processes
 - Transferring data
 - Examining state
 - ...
- Also many other operating systems



Linux Kernel Implementation

- Kernel may execute in either *Process context* vs. *Interrupt context*
- In *Process context*, kernel has access to
 - Virtual memory, files, other process resources (of one process)
 - May sleep, take page faults, etc., on behalf of process
- In *Interrupt context*, no assumption about what process was executing (if any)
 - No access to virtual memory, files, resources
 - May not sleep, take page faults, etc.

Modern Linux Threads (continued)

- Multiple threads can be executing *in kernel* at same time
- When in *process context*, kernel can
 - sleep on behalf of its thread
 - take pages faults on behalf of its thread
 - move data between kernel and process or thread
 - ...

Implications of Threads in Kernel

- Kernel maintains a small stack for each thread in kernel space
- Large enough for *one* system call
 - 4 kBytes (32 bit architectures)
 - 8 kBytes (64-bit architectures)
- Limit on total number of threads supportable by kernel

Bound on number
of function calls
within system call

Threads in Linux Kernel

■ Kernel has its own threads

- No associated *process context*

■ Supports concurrent activity within kernel

- Multiple devices operating at one time
- Multiple application activities at one time
- Multiple processors in kernel at one time

■ A useful tool

- Special kernel thread packages, synchronization primitives, etc.
- Useful for complex OS environments

Windows NT/XP/Vista Threads

■ Much like Linux 2.6 threads

- Primitive unit of scheduling defined by kernel
- Threads can block independently of each other
- Threads can make kernel calls
- ...

■ Process

- A higher level (non-kernel) abstraction
- A *container*

Threads – Summary

- Threads were invented to counteract the heavyweight nature of *Processes* in Unix, Windows, etc.
- Provide lightweight concurrency *within* a single address space
- Have evolved to become *the* primitive abstraction defined by kernel
 - Fundamental unit of scheduling in Linux, Windows, etc

Reading Assignment

- **OSTEP**
 - §25 – §27

- **Robert Love, *Linux Kernel Development***
 - Chapter 3 – “Process Management”

Questions?