

Introduction to Synchronization

Professor Hugh C. Lauer

CS-3013 — Operating Systems

(Slides include copyright materials from *Operating Systems: Three Easy Steps*, by Remzi and Andrea Arpaci-Dusseau, from *Modern Operating Systems*, by Andrew S. Tanenbaum, 3rd edition, and from other sources)

Challenge

- Now that concurrent execution in the same address space has been established, ...
- ... how do threads or concurrent actions synchronize themselves with each other?

Challenge (continued)

- How does one thread “know” that another has completed an action?
- How do separate threads “keep out of each others’ way” with respect to some shared resource?
- How do threads divide up and share the load of particularly long computations?

Context

- **Separate threads in same process**
- **Separate threads or processes making system calls to same kernel**
- **Separate processes sharing some common resource**
- **...**

Digression (thought experiment)

```
int y = 0;
```

```
int main(int argc, char **argv)
{
    extern int y;

    y = y + 1;

    return y;
}
```

Upon completion of `main`, `y == 1`

Thought experiment (continued)

```
int y = 0;
```

Thread 1

```
int main(int argc, char
    **argv)
{
    extern int y;

    y = y + 1;

    return y;
}
```

Thread 2

```
int main2(int argc, char
    **argv)
{
    extern int y;

    y = y - 1;

    return y;
}
```

Assuming threads run “at the same time,” what are *possible values* of *y* after both threads terminate?



Definition – *Atomic Operation*

- **An operation that either happens entirely or not at all**
 - No partial result is visible or apparent
 - Appears to be non-interruptible
- **If two atomic operations happen “at the same time”**
 - Effect is as if one is first and the other is second
 - (Usually) don't know which is which

Hardware Atomic Operations

- On (essentially) all computers, reading and writing of machine words can be considered as *atomic*
 - Non-interruptible
 - It either happens or it doesn't
 - Not in conflict with any other operation

- When two attempts to read or write the *same* data, one is first and the other is second
 - Don't know which is which!

- No other guarantees
 - (unless we take extraordinary measures)

Atomic Operations (continued)

- Most modern processors provide some additional atomic operations
- Read-modify-write of memory values — e.g.,
 - Test-and-Set
 - Fetch-and-add
 - Swap Register with Memory
 - Memory-memory Exchange
 - ...

Generally confined to one memory location

Difficult to implement with multiple processors and multiple caches

Definitions

■ Definition: *race condition*

- When two or more concurrent activities try to do something with the same set of variables resulting in different values
- Random outcome

■ *Critical Region (aka critical section)*

- One or more fragments of code that operate on the same data, such that at most one activity at a time may be permitted to execute anywhere in that set of fragments.

Synchronization – Critical Regions

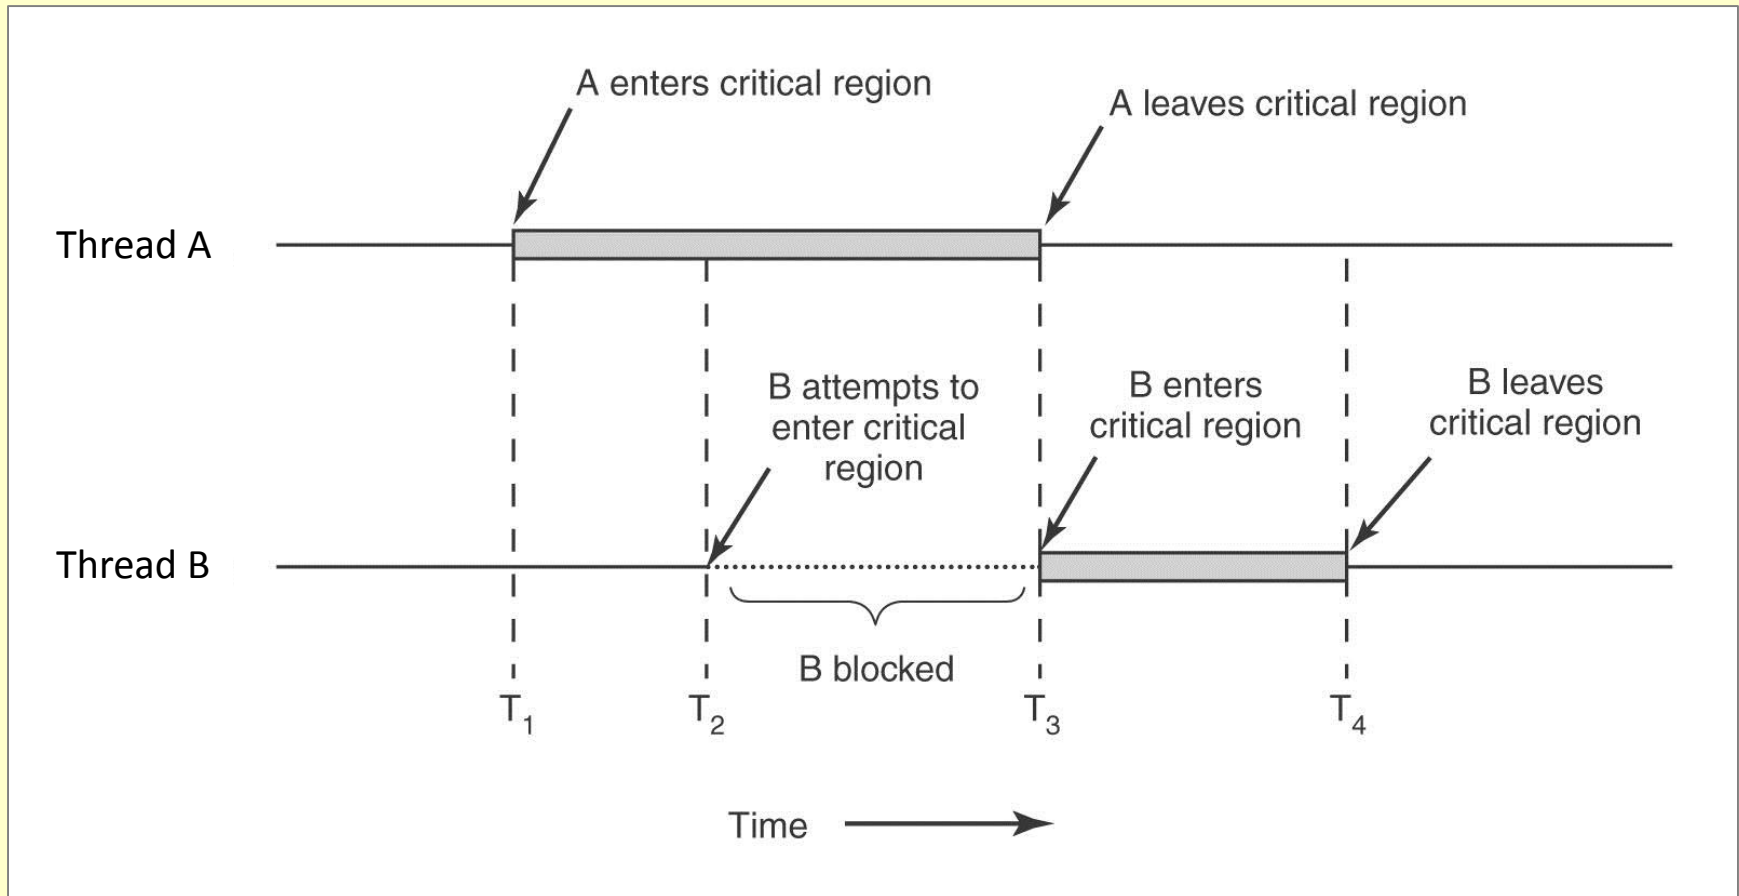


Fig 2-22, Tanenbaum, See also OSTEP Fig. 26.7

Class Discussion

- **How do we keep multiple computations from being in a critical region at the same time?**
 - Especially when number of computations is > 2
 - Remembering that read and write operations are atomic

example

Possible ways to protect critical regions

■ Without OS assistance

- Locking variables & busy waiting
- Wait-free solutions

■ With OS assistance — abstract data synchronization operations

- Single processor
- Multiple processors

Controlling access to a critical region

■ Classical requirements (Dijkstra):—

- Only one computation in critical section at a time
- Symmetrical among n computations
- No assumption about relative speeds
- A stoppage outside critical section does not lead to potential blocking of access to region by others
- No *starvation* — i.e. no combination of timings that could cause a computation to wait forever to enter its critical section

■ Practical requirement:—

- Completion in bounded time, regardless of behavior of other processes — i.e., *wait-free*



Non-solution

```
int turn = 0;
```

Thread 1

```
while (TRUE) {  
    while (turn !=0)  
        /*loop*/;  
    critical_region();  
    turn = 1;  
    noncritical_region1();  
};
```

Thread 2

```
while (TRUE) {  
    while (turn !=1)  
        /*loop*/;  
    critical_region();  
    turn = 0;  
    noncritical_region2();  
};
```



What is wrong with this approach?

Peterson's solution (2 processes or threads)

```
int turn = 0;
int interested[2];

void enter_region(int process) {
    int other = 1 - process;

    interested[process] = TRUE;
    turn = process;
    while (turn == process &&
           interested[other] == TRUE)
        /*loop*/;
};

void leave_region(int process) {
    interested[process] = FALSE;
};
```

This is a simplification of Dijkstra's 1965 solution for n processes. See [.pdf](#)

Busy waiting!



Another approach: Test & Set

(atomic read-modify-write instruction built into CPU hardware)

```
int lock = 0;
extern int TestAndSet(int *i);
    /* atomically sets the value of
    i to 1 and returns the previous
    value of i.  */

void enter_region(int *lock) {
    while (TestAndSet(lock) == 1)
        /* loop */ ;
};

void leave_region(int *lock) {
    *lock = 0;
};
```

What about this solution?



Busy waiting!



Variations of Atomic Operations

■ Exchange (a, b)

- temp = b
- b = a
- a = temp

■ Compare and Swap (var, old, new)

- previous = var;
- If (previous == old)
var = new;
- return previous;



Net effect

- We can simulate the *atomicity* of critical sections using instructions available in computer processor
- Is this the best approach?
- Sometimes yes, sometimes no!

There are entire courses and textbooks devoted to this subject!

The Art of Multiprocessor Programming, by Maurice Herlihy
& Nir Shavit, Morgan Kaufman, 2012

Possible Ways to Protect Critical Sections

■ Without OS assistance

- Locking variables & busy waiting
- Wait-free solutions

■ With OS assistance — abstract data synchronization operations

- Single processor
- Multiple processors

Protecting Critical Section with OS Assistance

Implement an *abstraction*:-

- A data type called *semaphore*
 - Non-negative integer values plus a queue
- An operation *wait_s(semaphore *s)* such that
 - Atomically test $s > 0$ and, if so, decrement s and proceed.
 - if $s = 0$, block the process or thread until some other process or thread executes *post_s(s)*
- An operation *post_s(semaphore *s):-*
 - If one or more processes or threads are blocked on s , allow precisely one of them to unblock and proceed
 - Otherwise, atomically increment s and continue

Critical Section control with Semaphore

```
semaphore mutex = 1;
```

Thread 1

```
while (TRUE) {  
    wait_s(mutex);  
  
    critical_region();  
  
    post_s(mutex);  
  
    noncritical_region1();  
};
```

Thread 2

```
while (TRUE) {  
    wait_s(mutex);  
  
    critical_region();  
  
    post_s(mutex);  
  
    noncritical_region2();  
};
```

Does this meet the requirements for controlling access to critical sections?



Semaphores – History

- Introduced by E. Dijkstra in 1965
- *wait_s()* was called *P()*
 - Initial letter of a Dutch word meaning “test”
- *post_s()* was called *V()*
 - Initial letter of a Dutch word meaning “increase”
- In Linux kernel (and some other modern systems)
 - *wait_s()* is called *down()*
 - *post_s()* is called *up()*

Abstractions

- The *semaphore* is an example of a powerful *abstraction* defined by OS
 - I.e., a data type and some operations that add a capability that was not in the underlying hardware or system.
- Any program can use this abstraction to control critical sections and to create more powerful forms of synchronization among computations.
- OSTEP, Ch 28




Data Structures for Implementing Semaphores

```
class State {  
    long int PSW;  
    long int regs[R];  
    /*other stuff*/  
}  
class PCB {  
    PCB *next, *prev, *queue;  
    State s;  
  
    PCB (...); /*constructor*/  
    ~PCB(); /*destructor*/  
}
```


```
class Semaphore {  
    int count;  
    PCB *queue;  
  
    friend wait_s(Semaphore *s);  
    friend post_s(Semaphore *s);  
  
    Semaphore(int initial);  
        /*constructor*/  
    ~Semaphore();  
        /*destructor*/  
}
```



Semaphore Data Structures (continued)



From earlier lesson




■ Action – dispatch a process or thread to CPU

- Remove first PCB from *ReadyQueue*
- Load registers and PSW
- Return from interrupt or trap

■ Action – interrupt a process or thread

- Save PSW and registers in PCB
- If not blocked, insert PCB into *ReadyQueue* (in some order)
- Take appropriate action
- Dispatch same or another process from *ReadyQueue*

Implementation – Semaphore actions



■ Action – `wait_s (Semaphore *s)`

- Implement as a Trap (with interrupts *disabled*)

`if (s->count == 0)`


- Save registers and PSW in PCB
- Queue PCB on `s->queue`
- Dispatch next process on ReadyQueue

`else`

- `s->count = s->count - 1;`
- Re-dispatch current process

Event wait

Implementation – Semaphore actions



■ Action – `post_s (Semaphore *s)`

- Implement as a Trap (with interrupts disabled)

```
if (s->queue != null)
```

- Save current process in ReadyQueue
- Move first process on `s->queue` to ReadyQueue
- Dispatch some process on ReadyQueue

```
else
```

- `s->count = s->count + 1;`
- Re-dispatch current process

Event completion

Interrupt Handling

- (Quickly) analyze reason for interrupt
- Execute equivalent of *post_s* to appropriate semaphore as necessary
 - Implemented in device-specific routines
 - Real work of interrupt handler is done in a separate task-like entity in the kernel
- More about interrupt handling later in the course

Semaphores — Summary

- Interrupts transparent to processes
- Can be used to simulate *atomic* actions
 - On single processor systems
- `wait_s()` and `post_s()` behave as if they are atomic
- Useful for synchronizing among processes and threads

Semaphores – Epilogue

- A way for generic processes to synchronize with each other
- Not the *only* way
- Not even the *best* way in most cases
- More in next topic

Questions?

Synchronization and Atomic Operations within Linux Kernel

■ Reading assignment

- Robert Love, *Linux Kernel Development*, 3rd edition, Chapter 10

■ Lots of tools

- Atomic operations on integer (32- and 64-bit) and bits
- Spin Locks
- Kernel Semaphores
- Kernel Mutexes
- ...

Complications for Multiple Processors

- **Disabling interrupts is not enough for implementing “atomic” updates to kernel data**
 - Semaphore operations must themselves be implemented in critical sections!
 - Queuing and dequeuing PCB's must also be implemented atomically
 - Interrupt handlers need to update status atomically
 - Other control operations need protection
- **These problems need deeper thought!**
 - *The* central issue in transition from single processor kernel to multi-threaded kernel (Linux 2.4.x to 2.6.x)

Hierarchy of Solutions

- **Wait-free operations**
 - Supported by hardware
- **Disciplined use of spin locks**
 - Avoid unbounded locking
- **Higher-level synchronization primitives within kernel**

Wait-Free Synchronization

- **A whole mathematical theory about efficacy of atomic hardware operations**
 - Atomic Read-Modify-Write is the weakest
 - Exchange is stronger
 - Compare-and-Swap is the strongest
- **All require extraordinary circuitry in processor memory, and bus to implement atomically**
- **Herlihy, Maurice, “Wait-Free Synchronization,” *ACM Transactions on Programming Languages and Systems*, vol 11, #1, January 1992, pp. 124-149 ([.pdf](#))**

Compare-and-Swap

```
int compare_and_swap (int* reg,
                      int oldval, int newval)
{
    int old_reg_val = *reg;
    if (old_reg_val == oldval)
        *reg = newval;
    return old_reg_val;
}
```

- As implemented in IBM 370 and successors
 - Test `old_reg_val` against `oldval` to determine if success
- Similar to `CMPXCH` on Intel architectures

What can we do with Compare-and-Swap?

- Atomic arithmetic or bitwise operations
- Add an object to or remove an object from head of a linked list
- ...
- All with no lock variable!

Used to implement
Linux kernel atomic
operations

Multi-word Compare-and-Swap Operations

- Harris, Timothy L., Fraser, Keir, and Pratt, Ian A., “A Practical Multi-Word Compare-and-Swap Operation,” University of Cambridge Computer Laboratory, Cambridge, UK ([.pdf](#))
- For updating queues, linked lists, etc., without lock variables

Load Linked / Store Conditional

■ **Alternative to Compare-and-Swap**

- DEC Alpha, PowerPC, MIPS, ARM, etc.

■ ***Load Linked (LL)***

- Like a regular load instruction, but keep tabs on that memory address
- Keep a hardware flag on memory location (in cache)

■ ***Store Conditional (SC)***

- Like a regular store instruction, but *fails* if
 - Flagged location has been updated since LL instruction
 - Context switch occurs since LL instruction
 - Anything else that might have changed target location

Load Linked / Store Conditional (continued)

- Works well with RISC processors
- Works well with multi-level caches
 - Provided cached implement cache-consistency
- Easy to simulate Compare-and-Swap
- May *not* be nested

Load Linked / Store Conditional (example)

Atomic increment

```
try:    LL        R2, 0(R1)
        ADD       R3, R2, #1
        SC        R3, 0(R1)
        BEQZ      R3, try
```

Questions?

Linux Kernel — Hierarchy of Solutions

- **Wait-free operations**
 - Supported by hardware
- **Disciplined use of spin locks**
 - Avoid unbounded locking
- **Higher-level synchronization primitives within kernel**

Disciplined Use of Spin Locks

- **Spin Lock:– A lock variable that is waited-for by busy waiting**
 - E.g., *Test-and-Set*

- **Widely used in Linux kernel**
 - Protect shared data
 - Critical sections must be *very short*
 - Robust in multi-processor environment
 - May be statically or dynamically allocated



Disciplined Use of Spin Locks

■ Rules

- Critical sections must *very* short – a few nanoseconds!
- Process holding a spinlock may not be pre-empted or rescheduled by any other process or kernel routine
 - I.e., disable interrupts!
- Process may not sleep, take page fault, or wait for *any reason* while holding a spin lock

■ Interrupt handler must

- Disable interrupts on current processor before locking

Reader-Writer Spin Locks

- Multiple tasks may hold a lock as “reader”
- Only one task may hold lock as “writer”
 - Must wait till all readers clear
- Reader may upgrade status to “writer”

Questions?

Hierarchy of Solutions

- **Wait-free operations**
 - Supported by hardware

- **Disciplined use of spin locks**
 - Avoid unbounded locking

- **Higher-level synchronization primitives within kernel**

Additional Synchronization Primitives in Linux Kernel

■ Semaphores

- Counting and binary

■ Mutexes

- Specialized variation of binary semaphore

■ Completion Variable

■ Sequential Lock

These all enable processes to sleep, be pre-empted, wait for events, etc.

Questions?