

Threads in the Operating System Kernel

Professor Hugh C. Lauer

CS-3013 — Operating Systems

(Slides include copyright materials from *Operating Systems: Three Easy Steps*, by Remzi and Andrea Arpaci-Dusseau, from *Modern Operating Systems*, by Andrew S. Tanenbaum, 3rd edition, and from other sources)

In the old days ...

- Operating system kernels did *one* thing at a time
- ... with interrupts disabled
- ... and all processes and threads suspended!
- Challenging enough to keep track of everything

No longer!

■ Desktop PC

- ~100 processes
- > 1,000 threads
- 8 processors

■ Shared system

- 1000s of processes
- Many 1000s of threads
- 100s of processors

■ Single-threaded kernel becomes serious bottleneck

Alternatives

■ Microkernels — e.g., MACH

- Different subsystems operate in separate address spaces
- Communication via message passing
- Performance issues

■ Cluster systems

- Partition applications across computers
- Shared files, but ...
- ... not much else

Need for multi-threaded kernel!

Multi-threaded Kernel

- **Linux kernel became multi-threaded in mid-2000s**
 - Between Linux 2.4.x and 2.6.x
- **Windows, other forms of Unix at about same time**
- **CS-3013 was using a multithreaded kernel in virtual machines by C-term 2008**

Linux approach

- ***Thread* is unit of scheduling**
- **Kernel maintains**
 - Interrupt stack for each processor (or core)
 - 4-8 kilobytes
 - Kernel stack for each thread
 - 4-8 kilobytes
 - Fixed location within address space

Interrupt handler

- Entered with interrupts disabled
- Do minimal processing to handle interrupt
 - Using interrupt stack of interrupted processor!
- Hand off to some thread for real work

More later in the course!

Definition – System Call

- A *structured* function call across a protection boundary between less privileged applications and more privileged operating system functions
- Also, across privilege layers of the operating system itself

Protection Boundary

- **Application programs are *not allowed* to**
 - Read or write data structures in the kernel
 - Call functions in the kernel directly
 - Change settings of the machine
 - Control arbitrary devices directly
 - Interfere with the operation of the kernel in any way

- **Enforced by hardware**
 - I.e., one or more bits in PSW

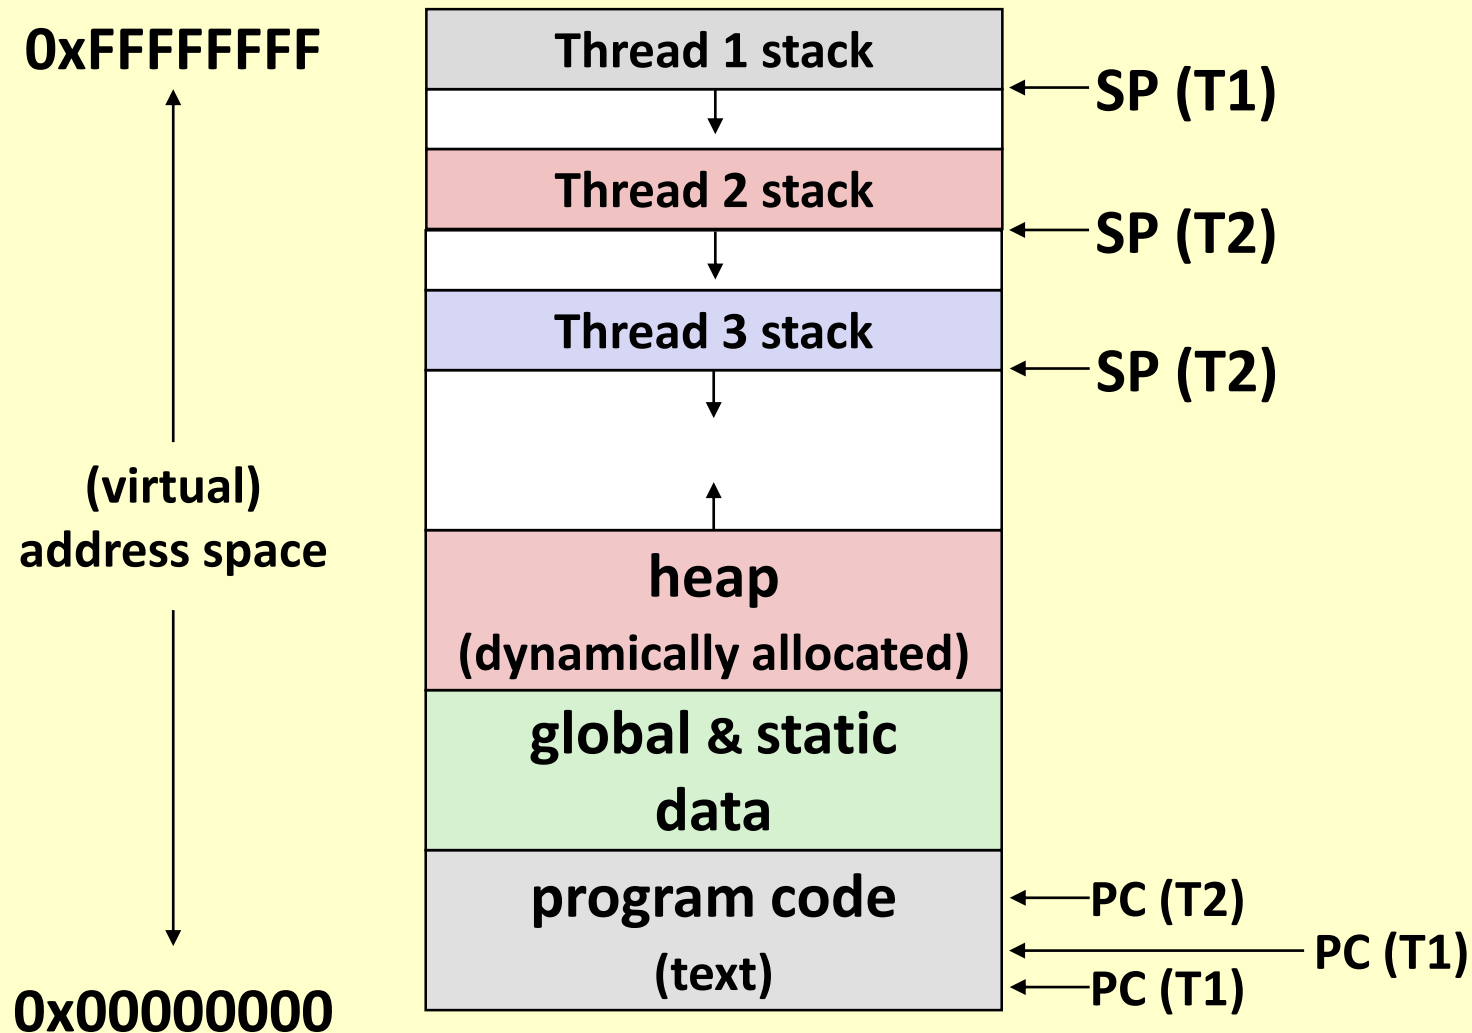
System Call

- A *trap* caused by executing a special machine language instruction
- Causes a *synchronous interrupt* to a specific interrupt/trap handler in the OS
- Allows the OS to control access, check arguments, manage behavior, etc.
- Causes machine to switch modes from “user” to “system” or “privileged”
 - As indicated by bits in the PSW

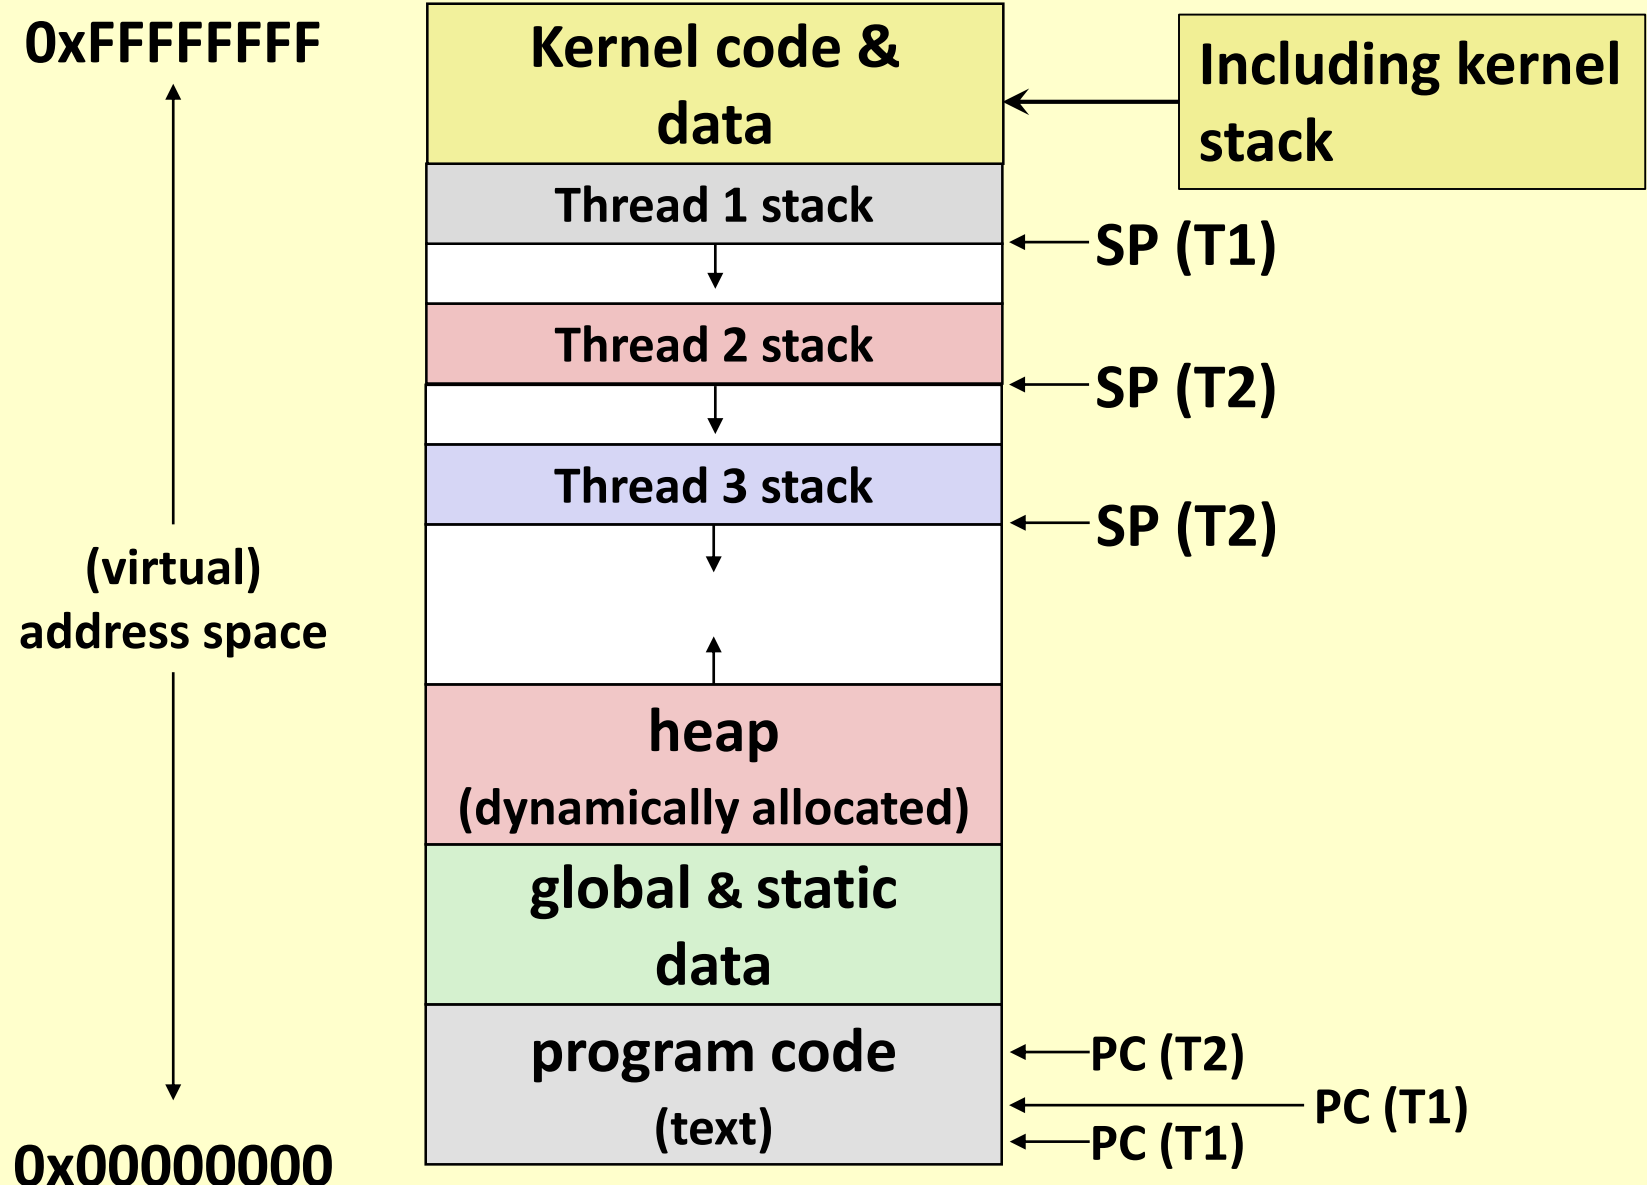
Trap handling

- Find kernel stack for *this thread*
- Enable interrupts
- Handle the trap or system call
 - As a kernel function
 - Like the syscall stubs that we implemented in Project 0

From previous topic



In reality:–



Digression – Process Address Space

■ Linux includes (parts of) *kernel* in every address space

- Protected ← But not touchable in non-privileged mode
- Easy to access ← In privileged mode
- Allows *kernel functions* to see into client processes
 - Transferring data
 - Examining state
 - ...

■ Also many other operating systems

Linux Kernel Implementation

- Kernel may execute in either *Interrupt context* or *Process context*
- In *Interrupt context*, no assumption about what process was executing (if any)
 - No access to virtual memory, files, resources
 - May not sleep, take page faults, wait for input, etc.
- In *Process context*, kernel has access to
 - Virtual memory, files, other process resources
 - May sleep, take page faults, etc., on behalf of process
 - May access shared resources & wait till available, etc.

Modern Linux Threads (continued)

- **Multiple threads can be executing *in kernel* at same time**
 - In various states of activity
- **Multiple processors can be executing *in kernel* at the same time**
 - Handling interrupts
 - In process context on behalf of some thread
- **Made possible by**
 - One kernel stack per thread
 - One interrupt stack per processor

Threads in Linux Kernel

■ Kernel has its own threads

- No associated *process context*

■ Supports concurrent activity within kernel

- Multiple devices operating at one time
- Multiple application activities at one time
- Multiple processors in kernel at one time

■ A useful tool

- Special kernel thread packages, synchronization primitives, etc.
- Useful for complex OS environments

Questions?