

CS-3013 — Operating Systems

Professor Hugh C. Lauer

CS-3013 — Operating Systems

(Slides include copyright materials from *Operating Systems: Three Easy Steps*, by Remzi and Andrea Arpaci-Dusseau, from *Modern Operating Systems*, by Andrew S. Tanenbaum, 3rd edition, and from other sources)

In the beginning (prehistory)...

- **Single usage (or batch processing) systems**
 - One program loaded in physical memory at a time
 - Runs to completion
- **If job larger than physical memory, use *overlays***
 - Identify sections of program that
 - Can run to a result
 - Can fit into the available memory
 - Add commands after result to load a new section
 - Example: passes of a compiler
 - Example: SAGE – *North American Air Defense System*

Still near the beginning (multi-tasking) ...

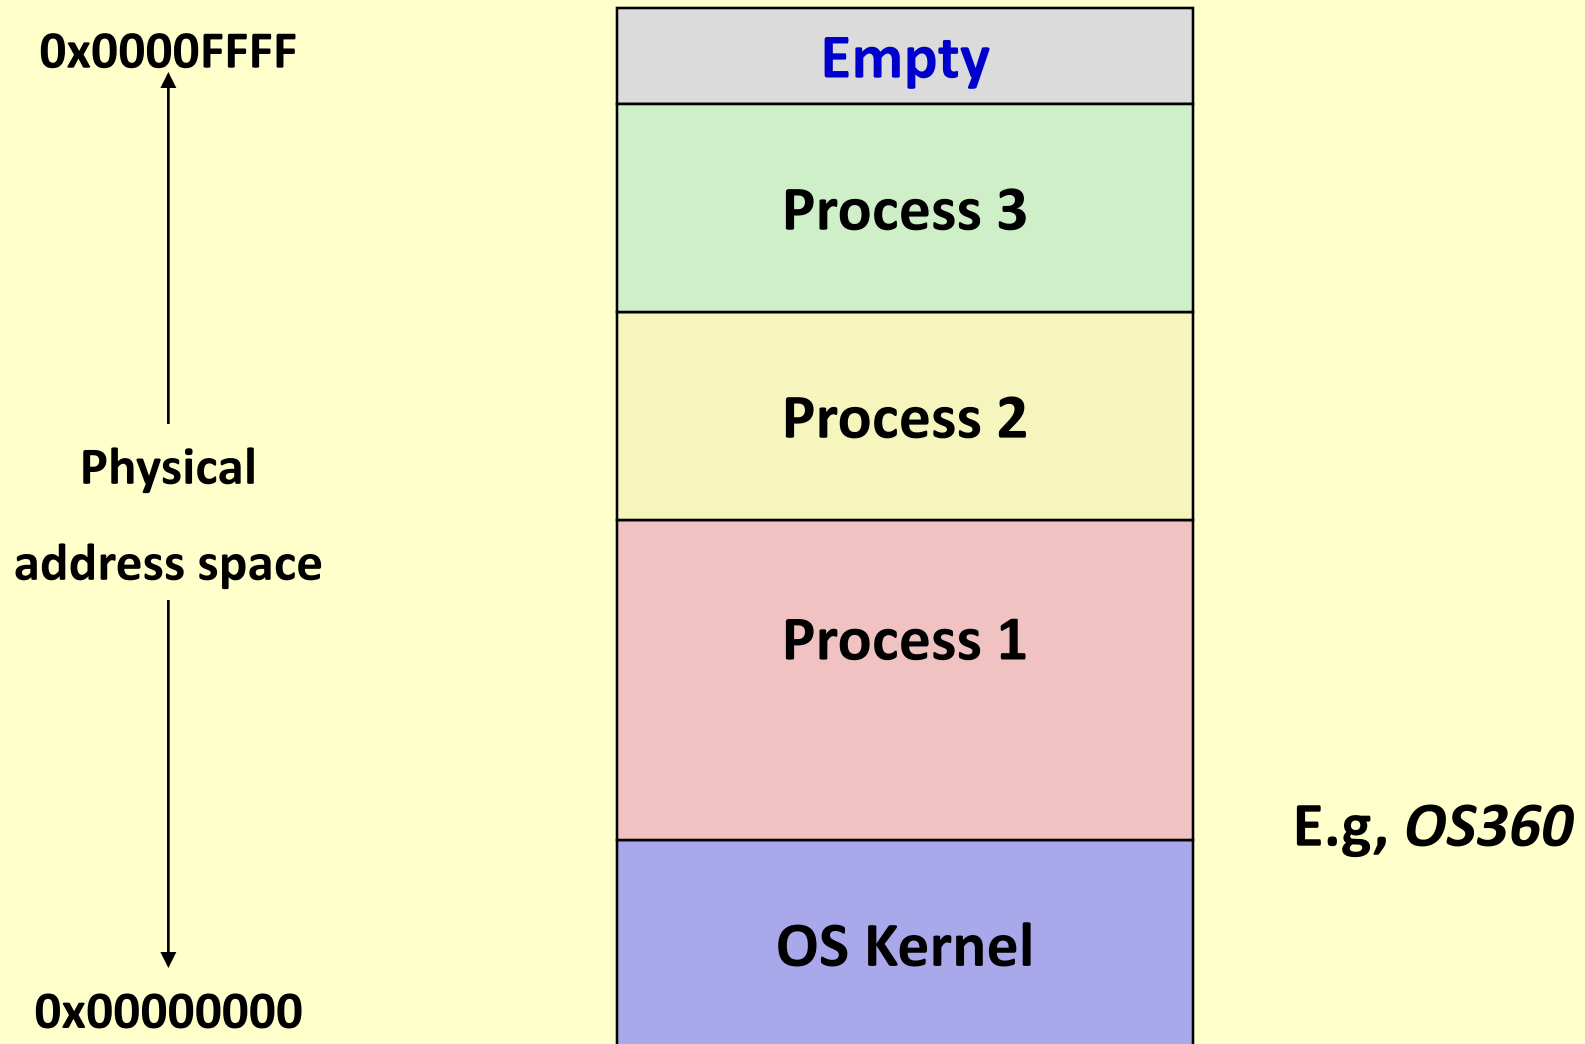
■ Multiple processes in physical memory at the same time

- allows fast switching to a ready process
- *Partition* physical memory into multiple pieces
 - One partition for each program
- Some modern operating systems
 - *Real-time* systems
 - Small, dedicated systems (mobile phone, automotive processors, etc.)

■ Partition requirements

- *Protection* – keep processes from smashing each other
- *Fast execution* – memory accesses can't be slowed by protection mechanisms
- *Fast context switch* – can't take forever to setup mapping of addresses

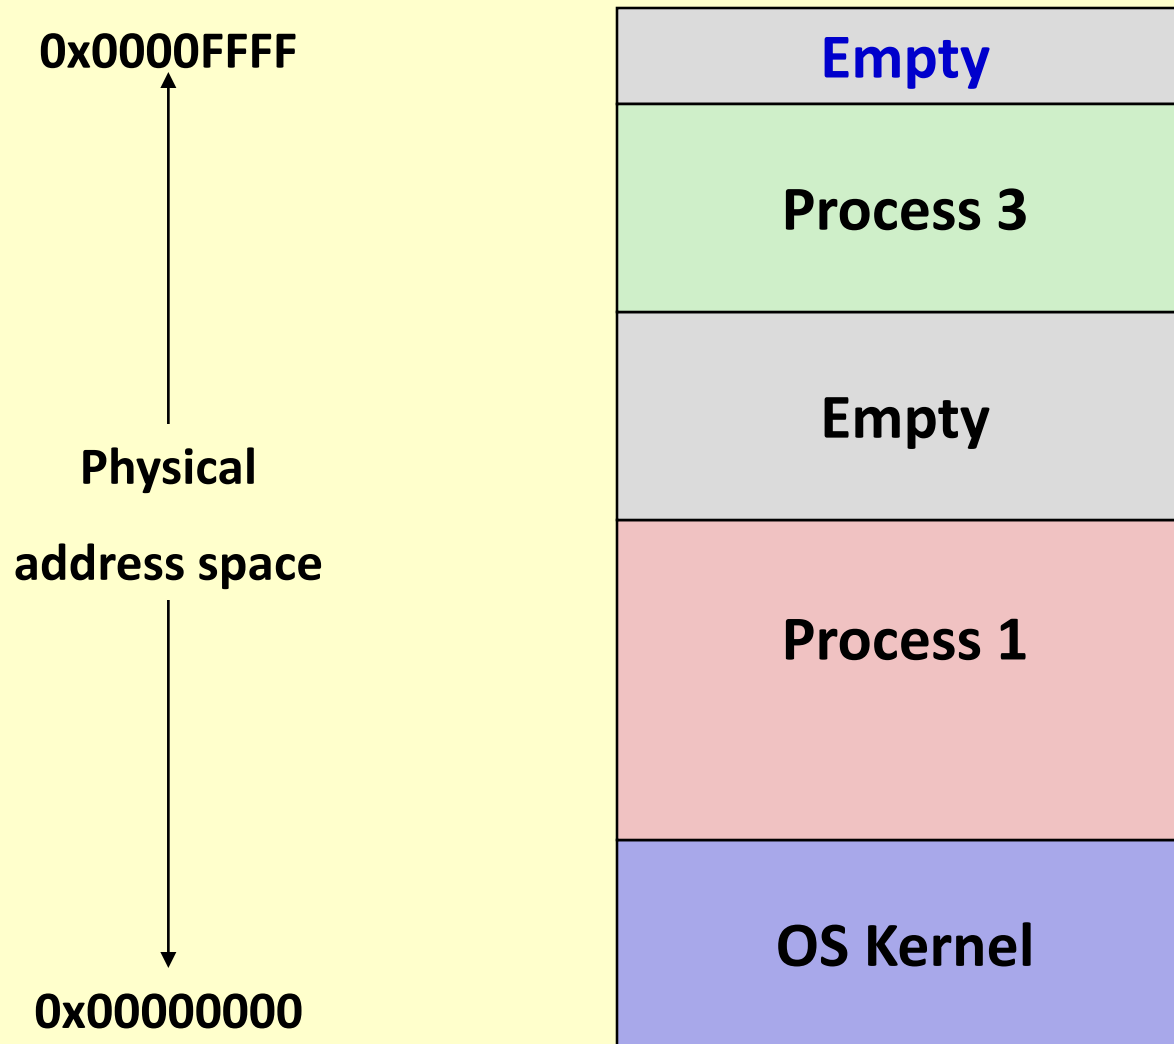
Physical memory



Loading a process

- Relocate all addresses relative to start of partition
 - See *Linking and Loading*
- Memory protection assigned by OS
 - Block-by-block to physical memory
 - Base and limit registers
- Once process starts
 - Partition cannot be moved in memory
 - *Why?*

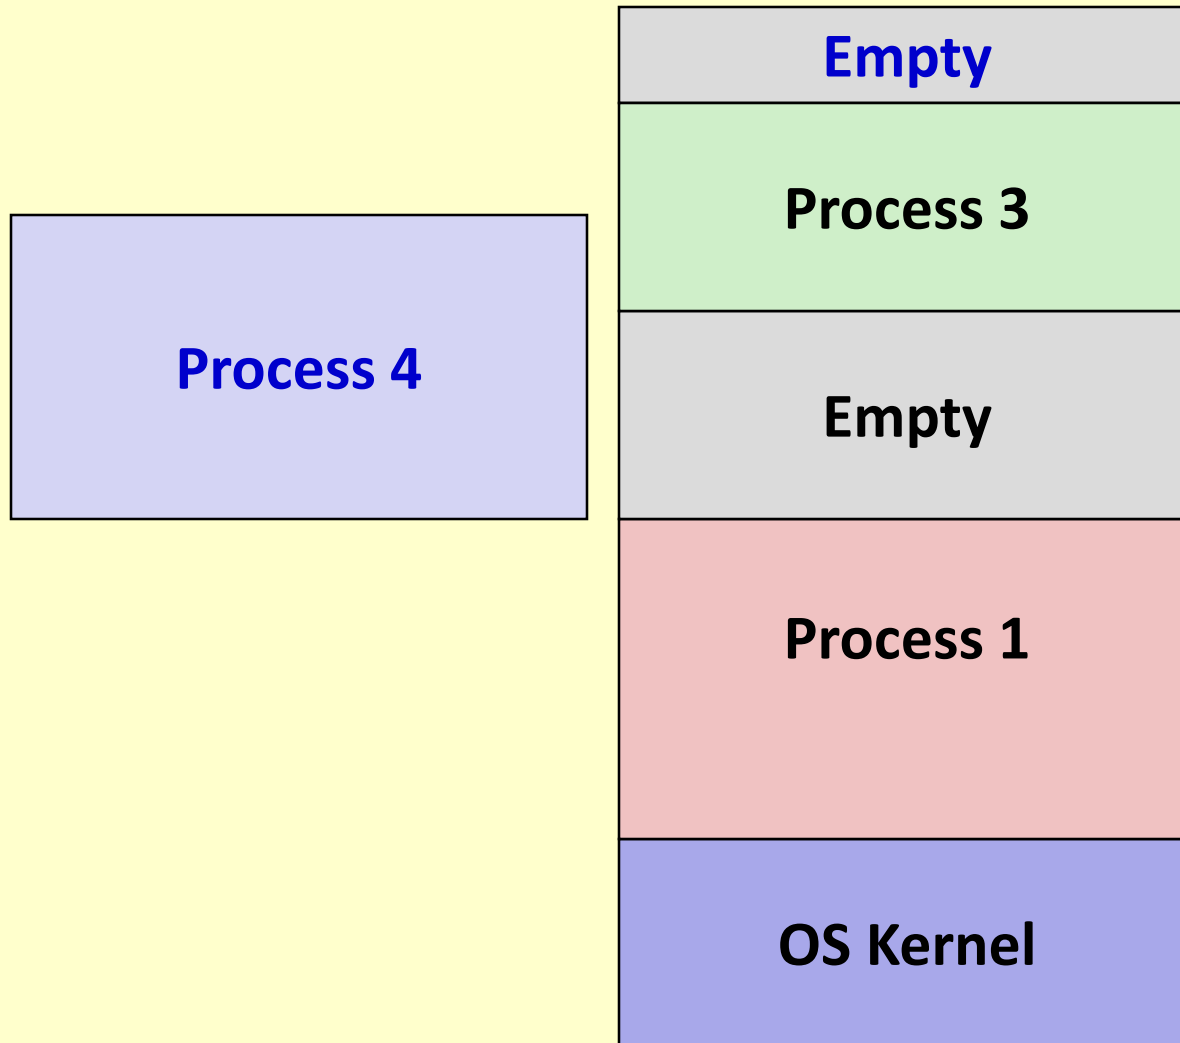
Physical memory – process 2 terminates



Problem

- **What happens when Process 4 comes along and requires space larger than the largest empty partition?**
 - Wait
 - Complex resource allocation problem for OS
 - Potential starvation

Physical memory



Solution

- ***Virtual Address***: an address used by the program that is translated by computer into a *physical address* **each time** it is used
 - Also called *Logical Address*
- When the program utters 0x00105C, ...
- ... the machine accesses 0x01605C

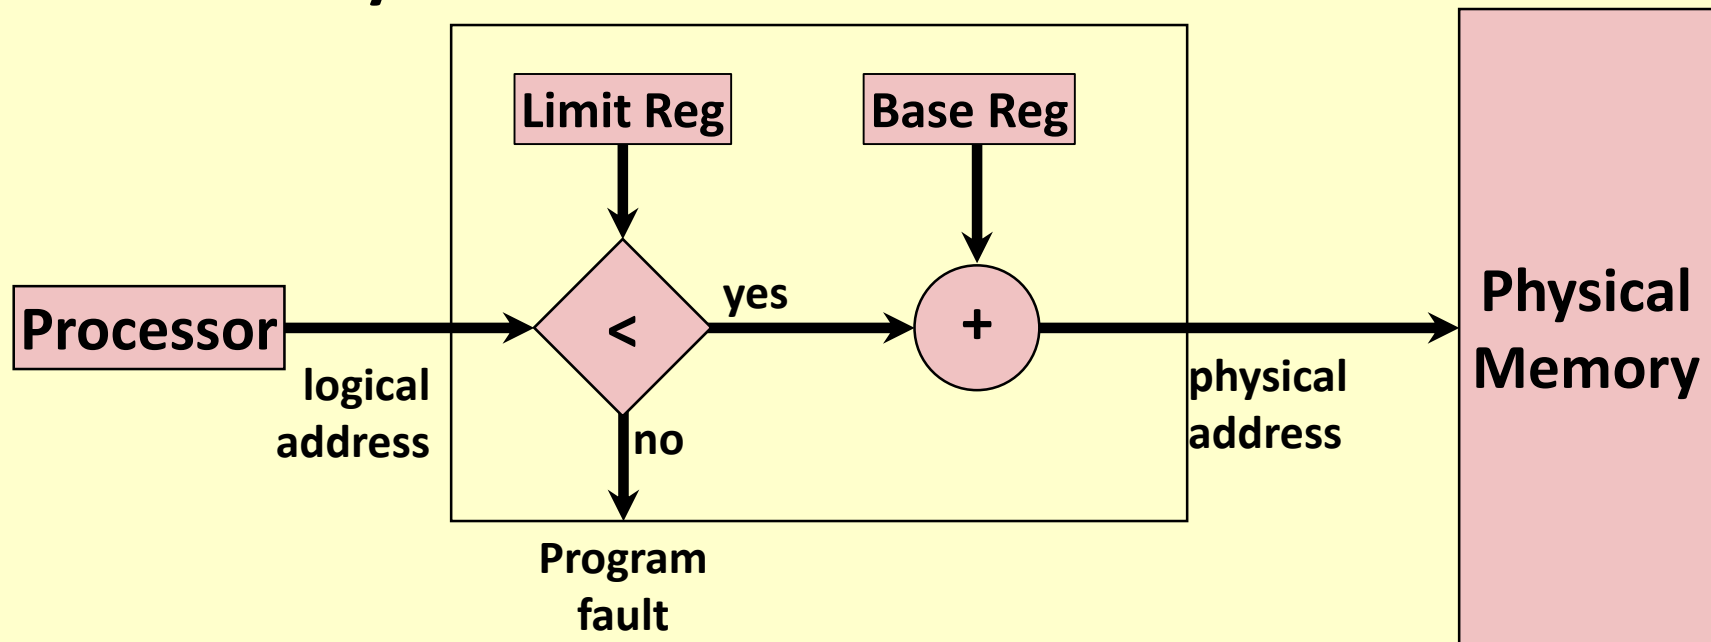
OSTEP §12-15

First implementation

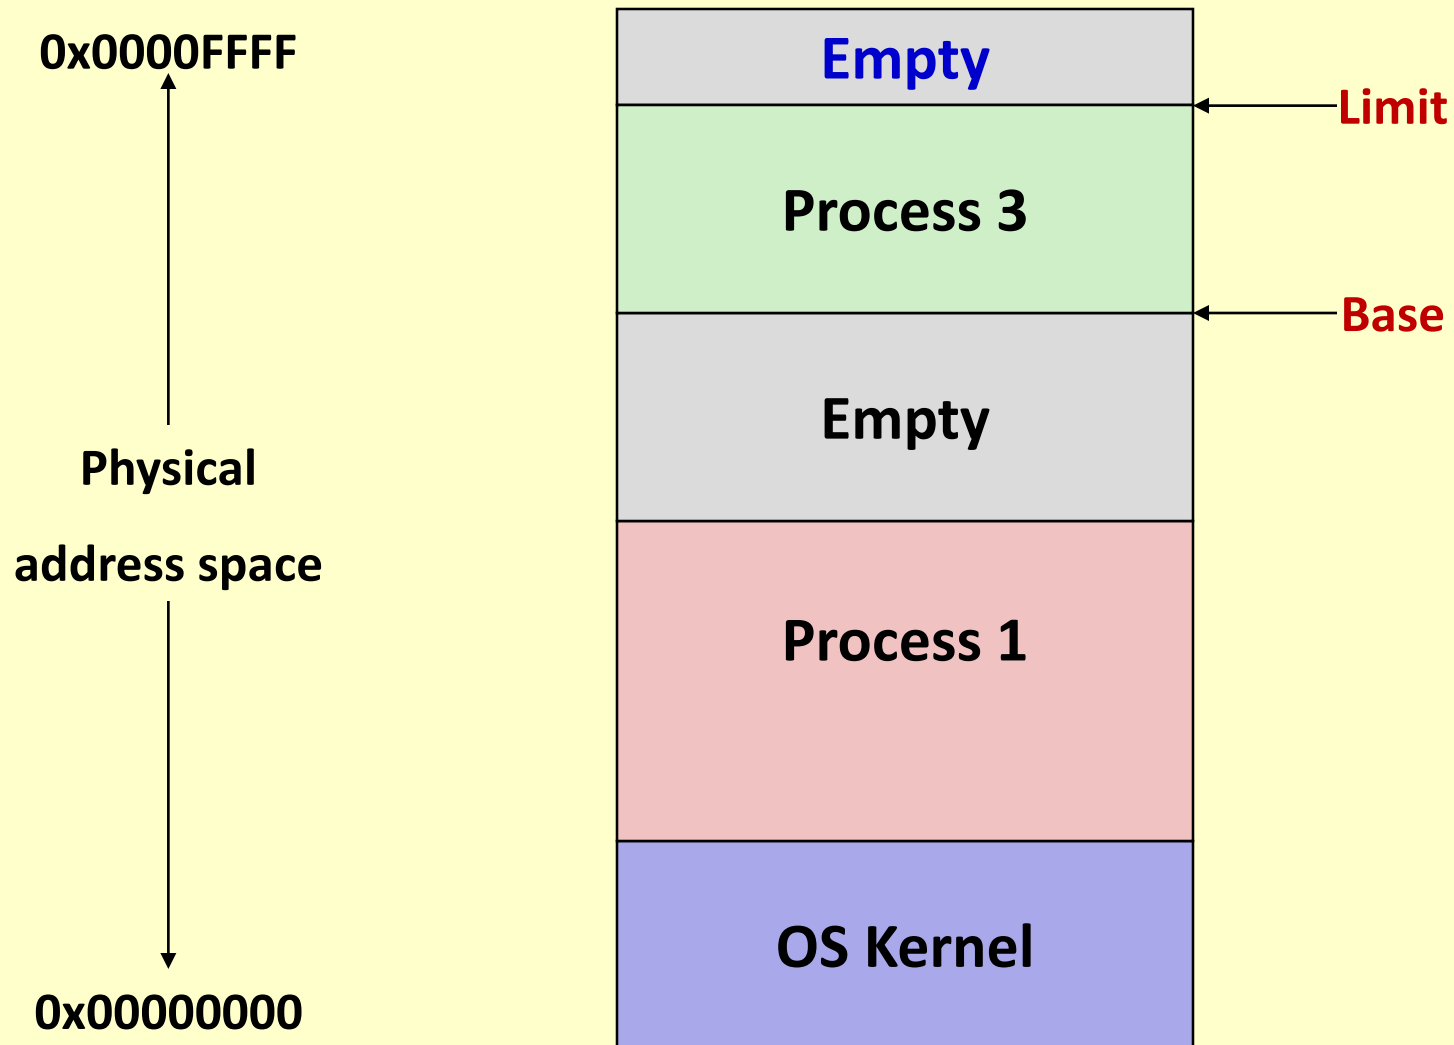
■ *Base* and *Limit* registers

- *Limit* is checked on all memory references
- *Base* is automatically added to all addresses
- Introduced in minicomputers of early 1970s

■ Loaded by OS at each context switch



Physical memory

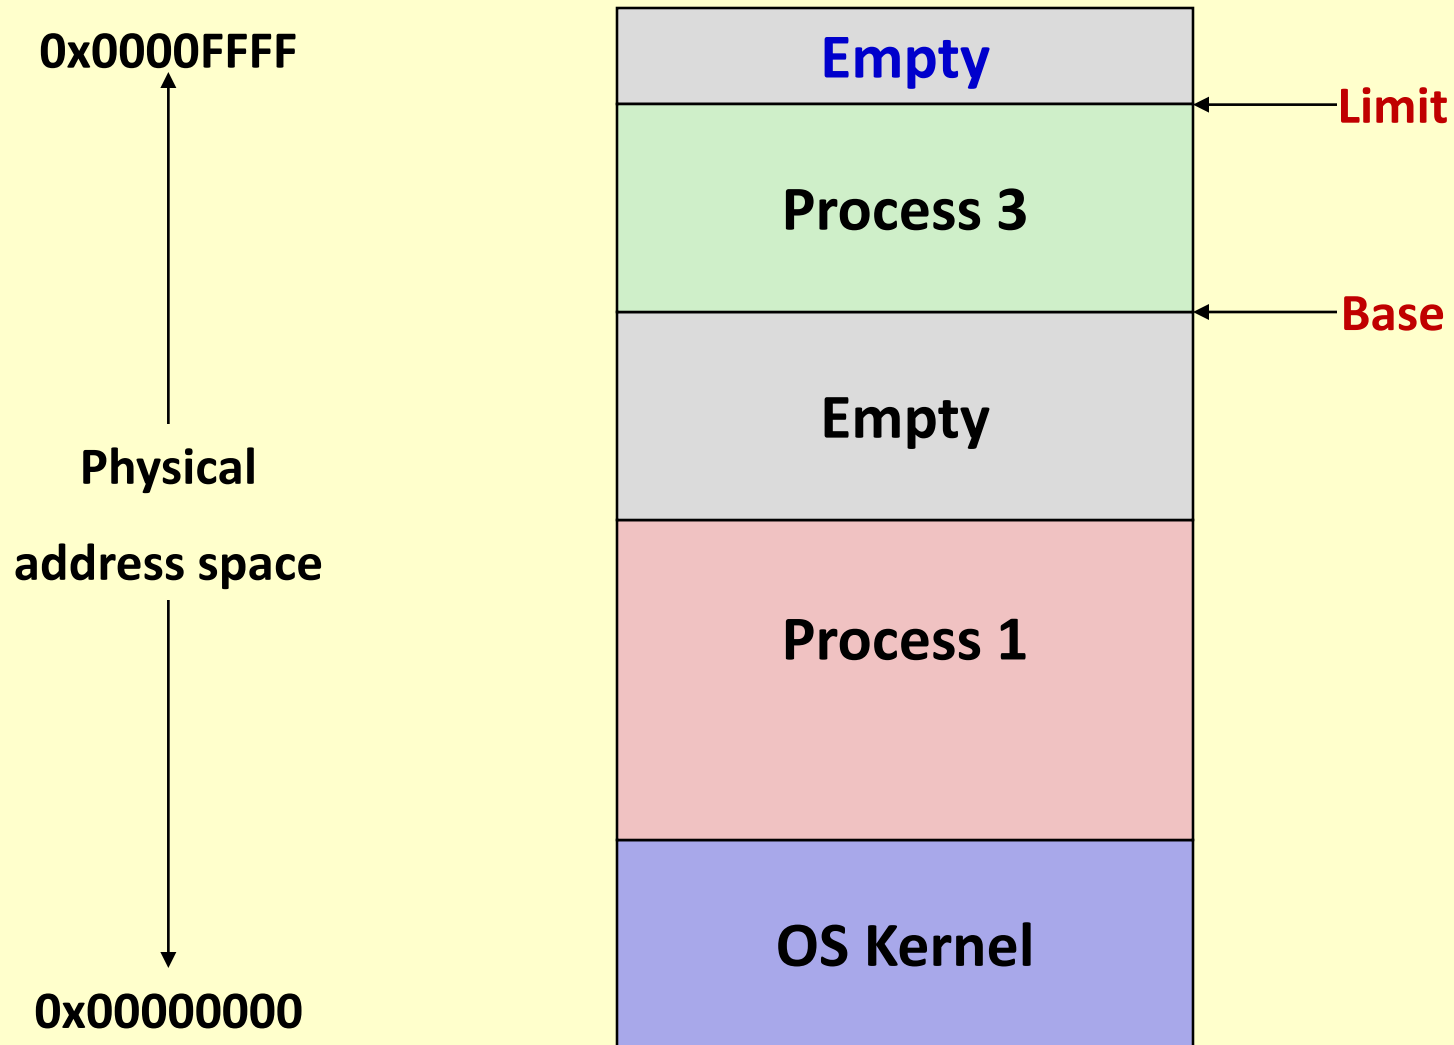


Advantages

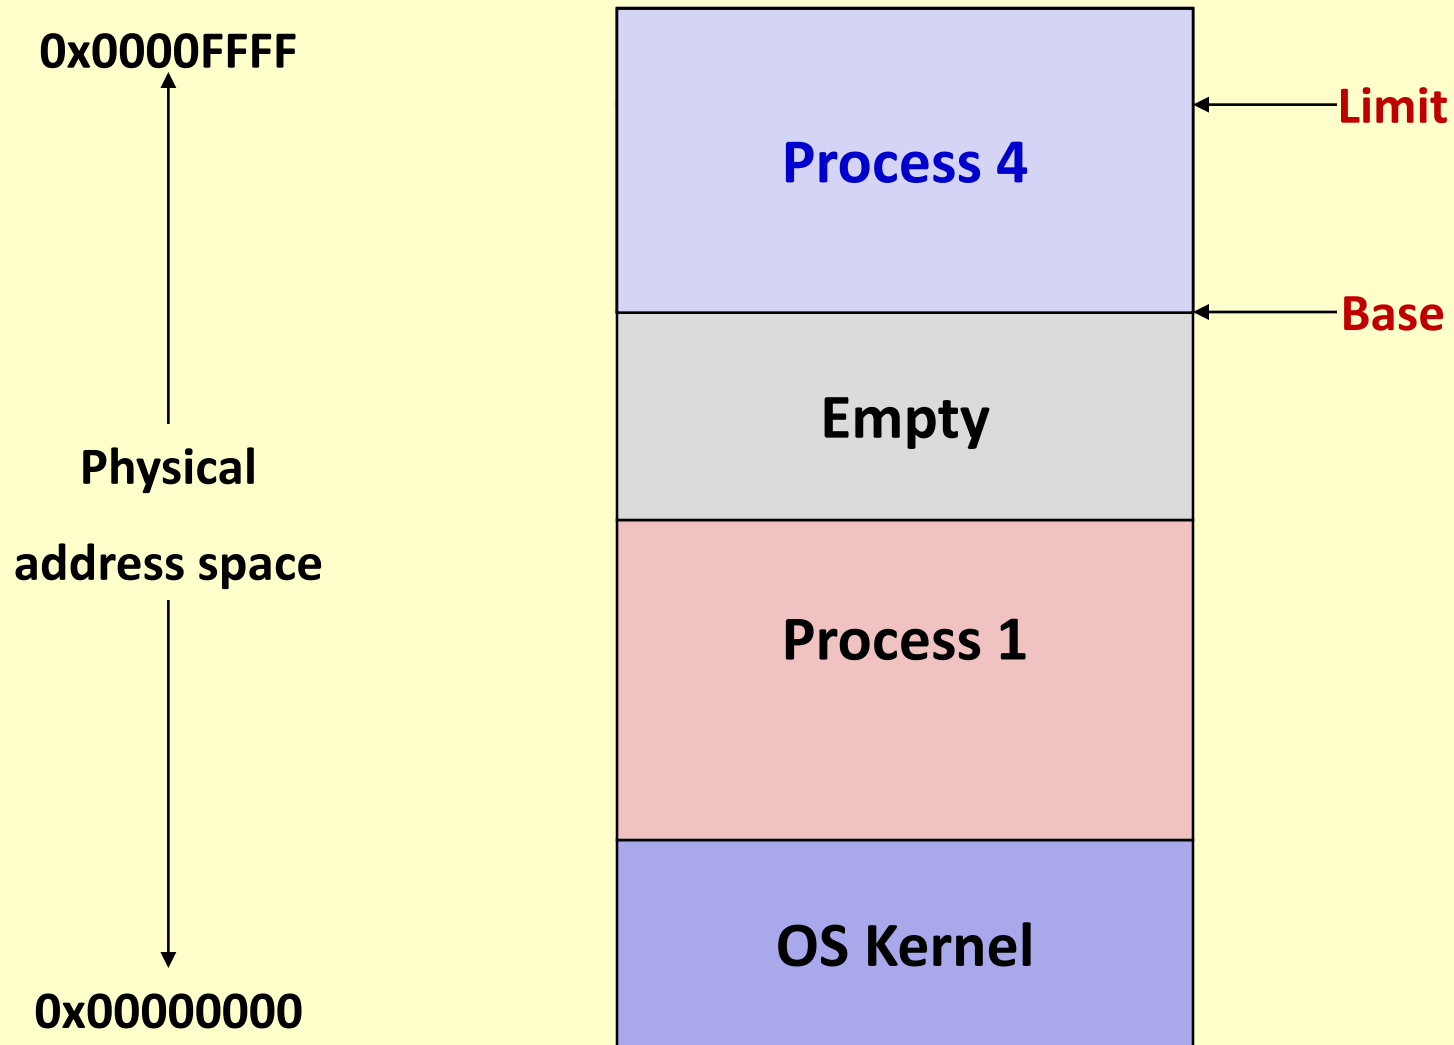
- **No relocation of program addresses at load time**
 - All addresses relative to zero!
- **Built-in protection provided by *Limit***
 - No physical protection per page or block
- **Fast execution**
 - Addition and limit check at hardware speeds within each instruction
- **Fast context switch**
 - Need only change base and limit registers
- **Partition can be suspended and moved at any time**
 - Process is unaware of change
 - Potentially expensive for large processes due to copy costs!



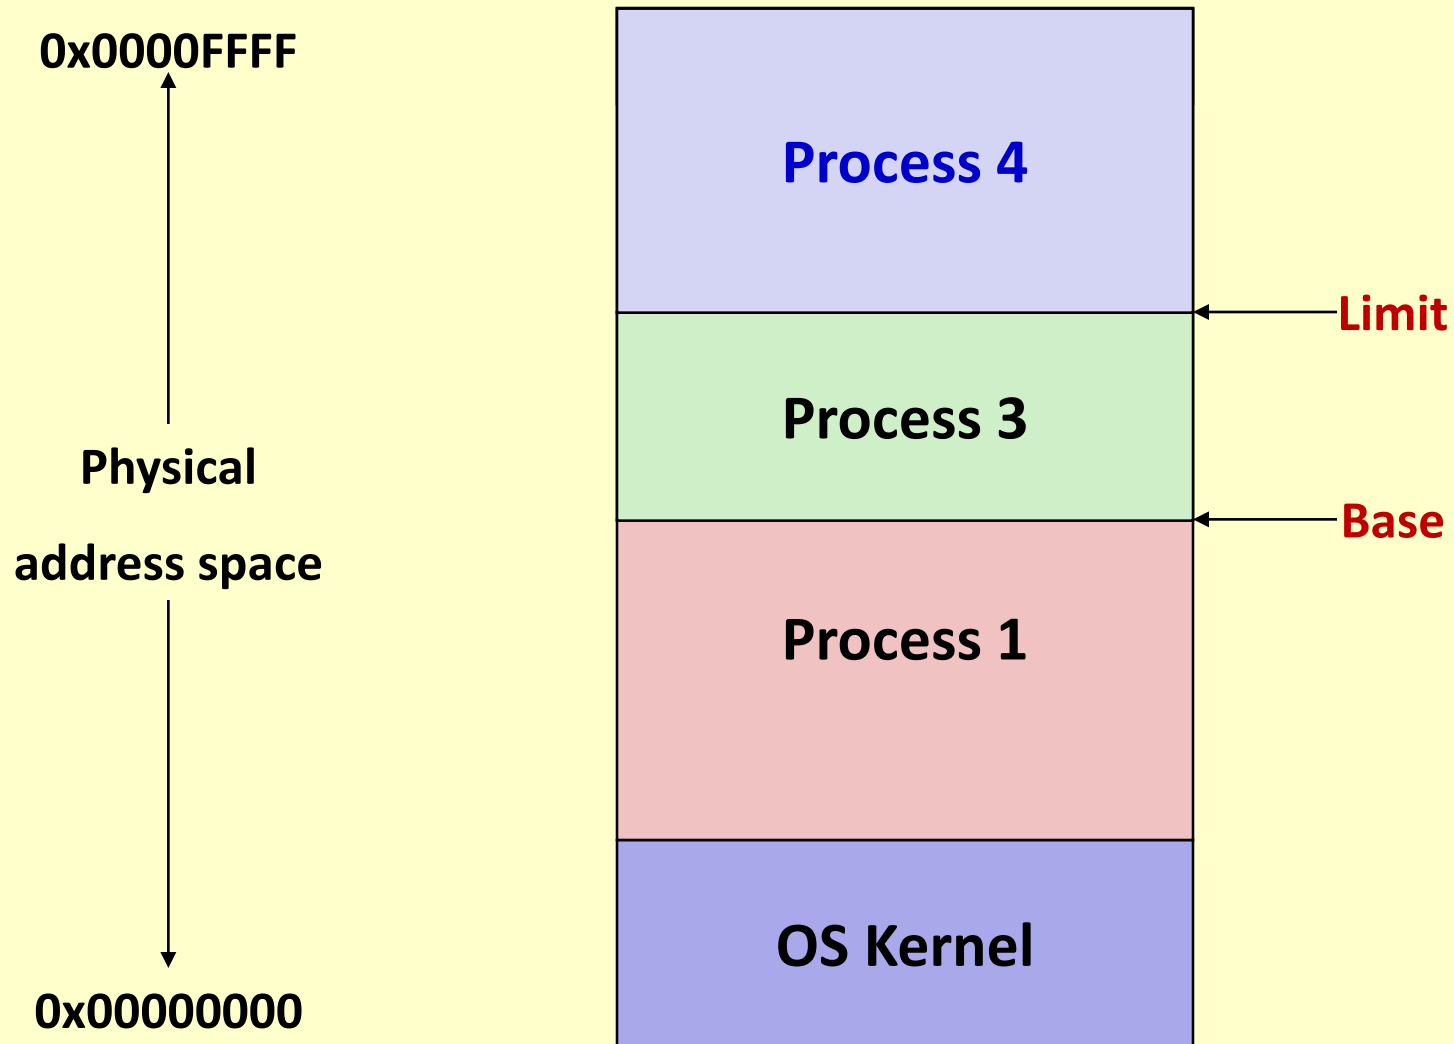
Physical memory



Physical memory



Physical memory

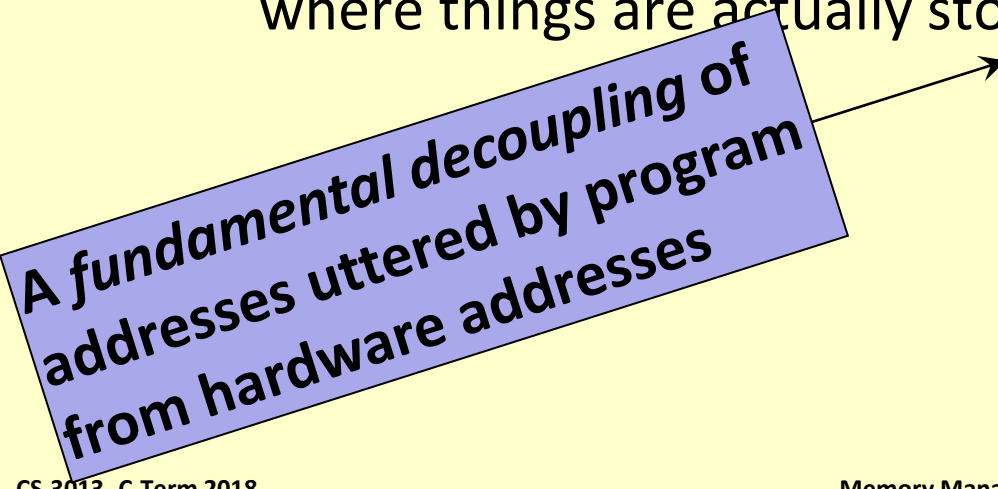


Definition

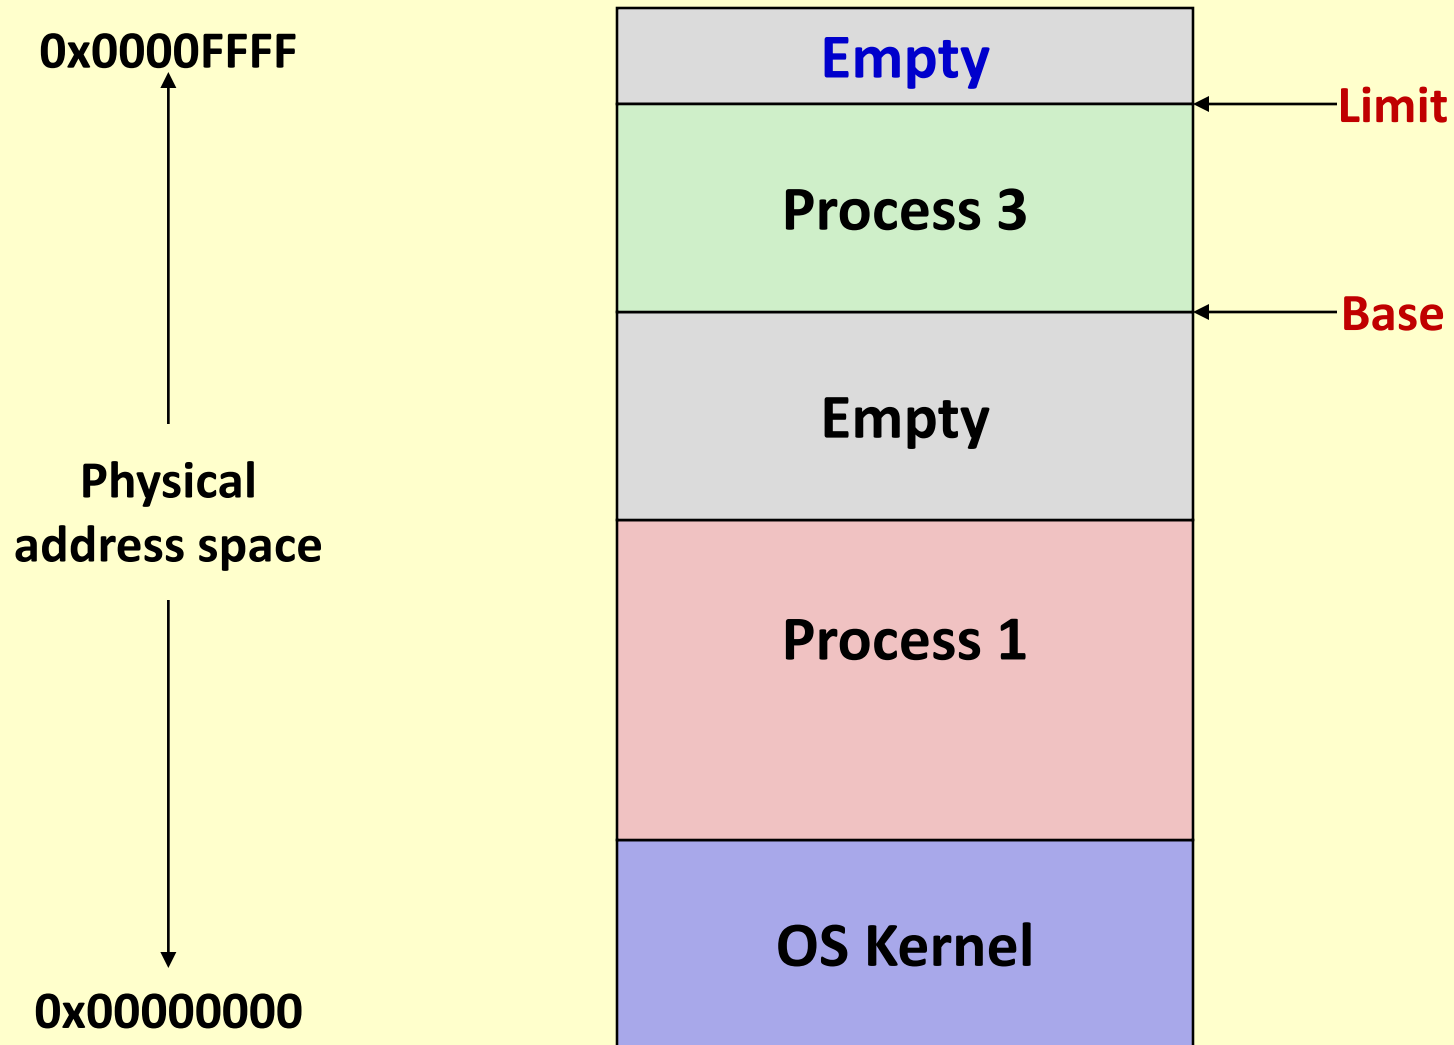
■ *Virtual Address Space:*

- The address space in which a process or thread “thinks”
- Address space with respect to which pointers, code & data addresses, etc., are interpreted
- Separate and independent of *physical address space* where things are actually stored

A fundamental decoupling of
addresses uttered by program
from hardware addresses



Physical memory



New problem:— how to manage memory

■ Fixed partitions

- Easy

■ Variable partitions

- Seems to make better use of space

Anything having to do with managing space — warehouse design, packaging, etc.

This is a general problem with broad applicability — e.g., to files systems, databases, etc.

Partitioning strategies – fixed

- **Fixed Partitions – divide memory into equal sized pieces (except for OS)**
 - Degree of multiprogramming = number of partitions
 - Simple policy to implement
 - All processes must fit into partition space
 - Find any free partition and load the process
- **Problem – what is the “right” partition size?**
 - Process size is limited
 - *Internal Fragmentation* – unused memory within a partition that is not available to other processes

Partitioning strategies – variable

- Idea: remove “wasted” memory that is not needed in each partition
 - Eliminating *internal fragmentation*
- Memory is dynamically divided into partitions based on process needs
- Definition:
 - *Hole*: a block of free or available memory
 - Holes are scattered throughout physical memory
- Memory is allocated to new process from hole large enough to fit it

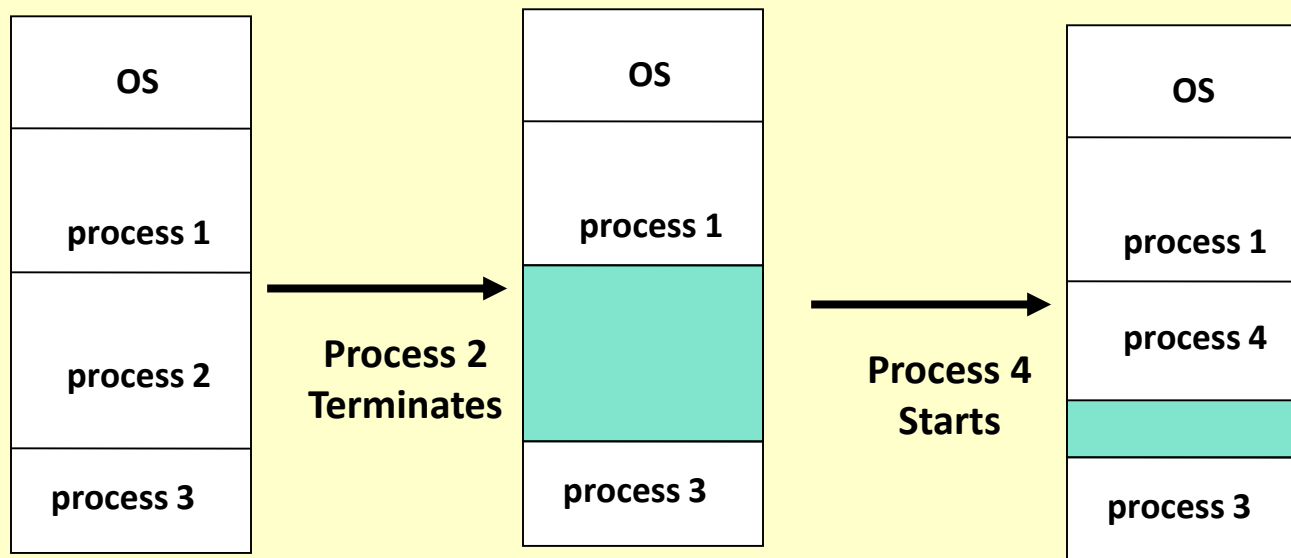
Variable partitions

■ More complex management problem

- Must track free and used memory
- Need data structures to do tracking
- What holes are used for a process?

■ *External fragmentation*

- memory that is outside any partition and is too small to be usable by any process



Definitions – *fragmentation*

- **Unused space that cannot be allocated to fill a need**
- ***Internal* fragmentation**
 - Unused or unneeded space *within* an allocated part of memory.
 - Cannot be allocated to another task/job/process
- ***External* fragmentation**
 - Unused space *between* allocations.
 - Too small to be used by other requests
- **Applies to all forms of *spatial* resource allocation**
 - RAM, Disk, Virtual memory within process
 - File systems
 - ...

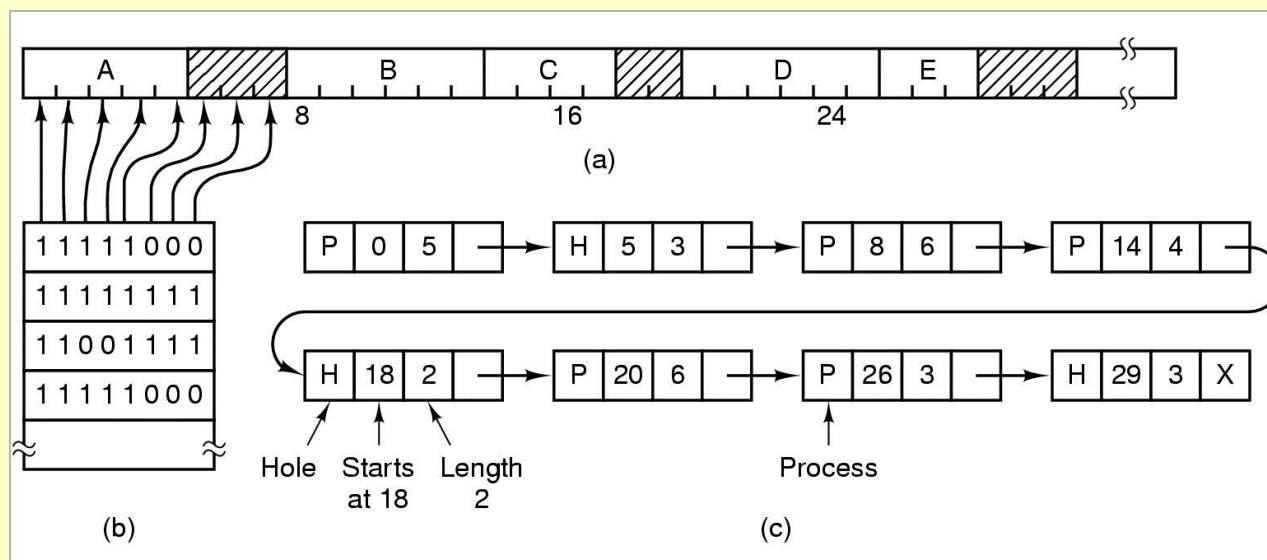
Memory allocation – mechanism

- **MM system maintains data about free and allocated memory alternatives**
 - *Bit maps* – 1 bit per “allocation unit”
 - *Linked Lists* – free list updated and coalesced when not allocated to a process
- **At swap-in or process create**
 - Find free memory that is large enough to hold the process
 - Allocate part (or all) of memory to process and mark remainder as free
- ***Compaction***
 - Moving things around so that *holes* can be consolidated
 - Expensive in OS time

See OSTEP, §17.1

Memory management – list vs. map

- Part of memory with 5 processes, 3 holes
 - tick marks show allocation units
 - shaded regions are free
- Corresponding bit map
- Same information as a list



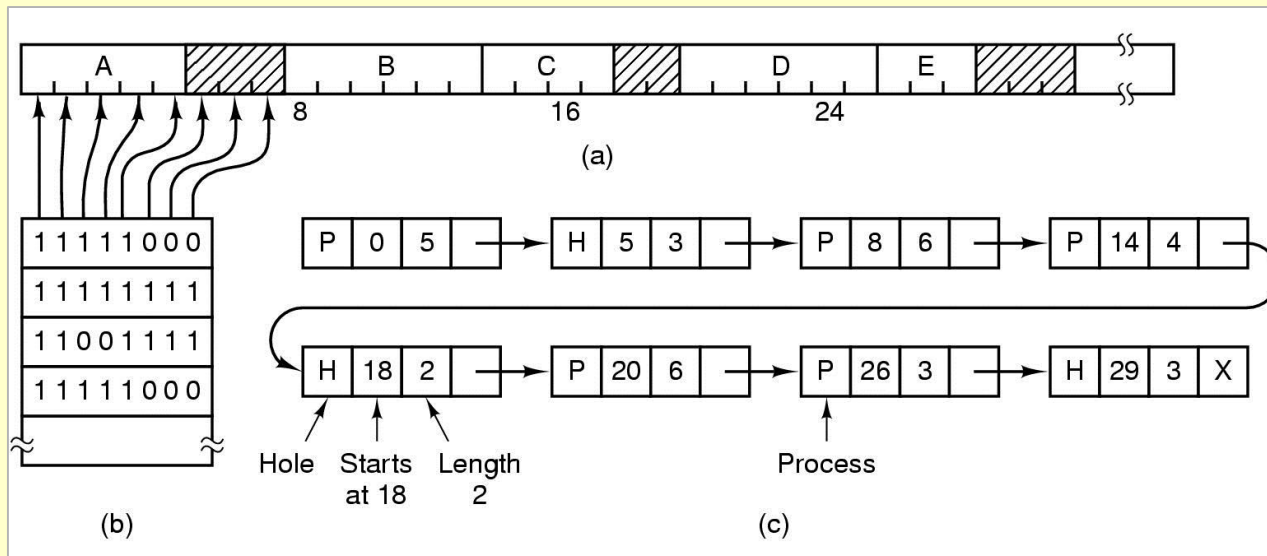
Memory management – bit map

■ Advantages:–

- Can see big picture
- Easy to search using bit instructions in processor
- Holes automatically coalesce

■ Disadvantage

- No association between blocks and processes that own them



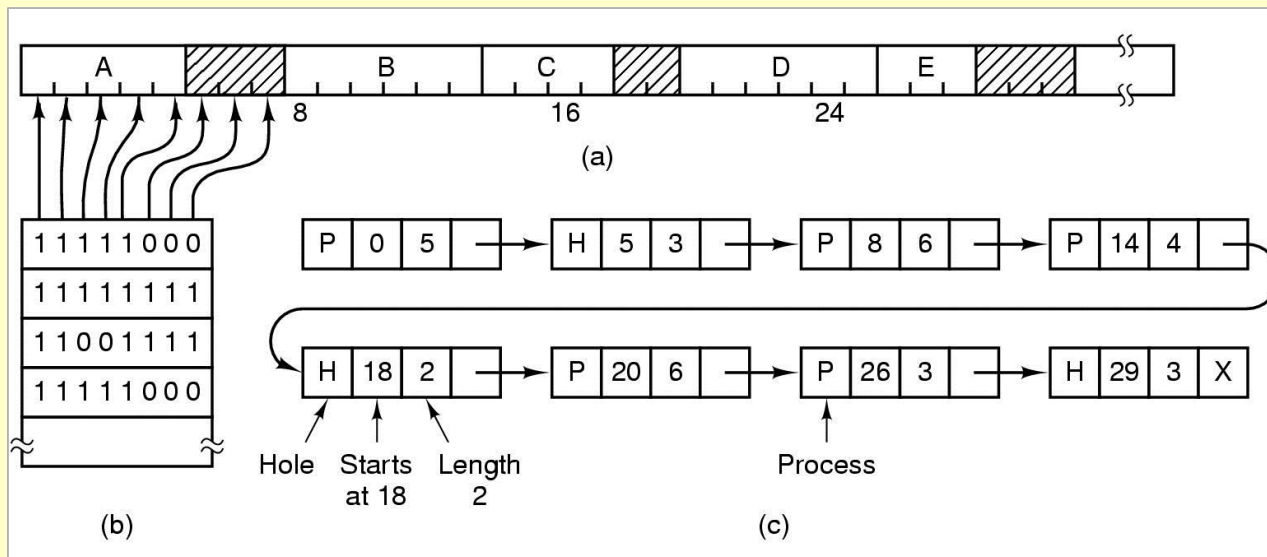
Memory management – list

■ Advantages:–

- Direct association between block and process owning it

■ Disadvantages:–

- Cannot see big picture
- Searching is expensive
- Coalescing adjacent blocks requires extra effort (sorted order)



Memory allocation – policies

■ Policy examples

- *First Fit*: scan free list from beginning and allocate first hole that is large enough – fast
- *Next Fit*: start search from end of last allocation
- *Best Fit*: find smallest hole that is adequate – slower and lots of fragmentation
- *Worst fit*: find largest hole

■ Simulation results show that *First Fit* usually works out to be the best

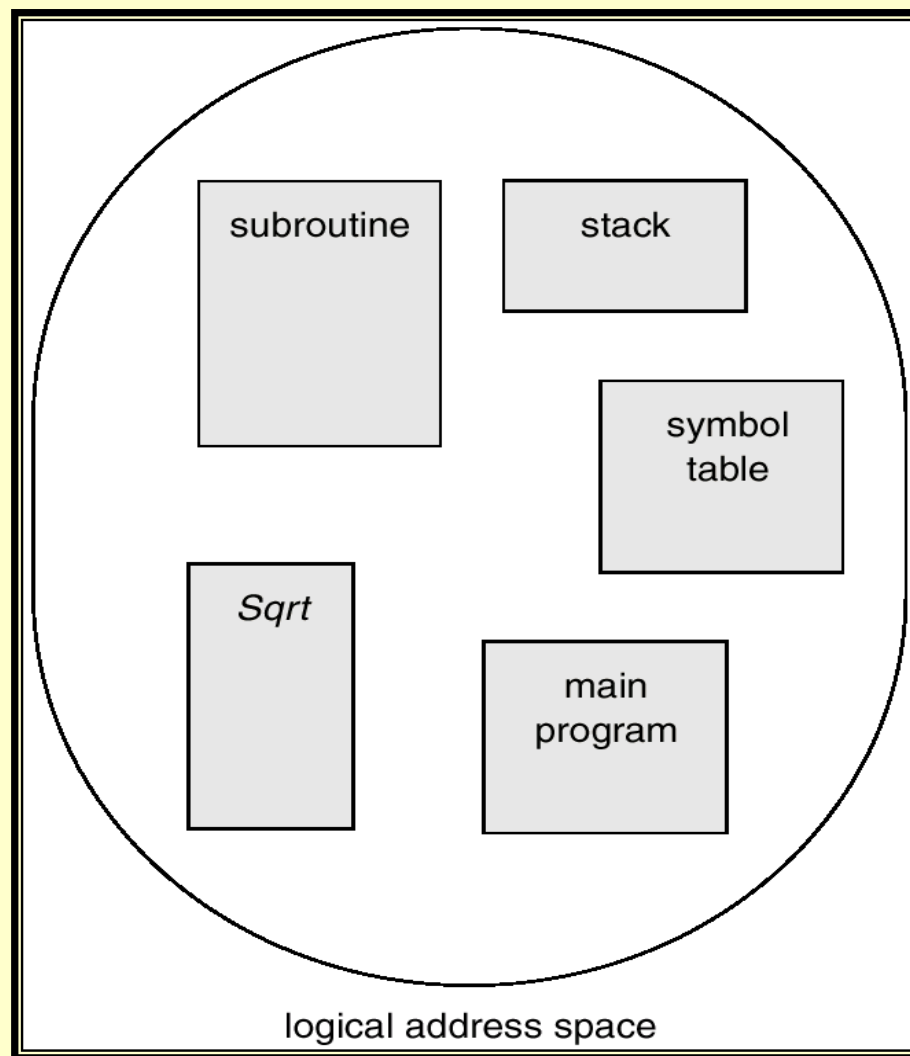
Swapping and scheduling

■ *Swapping*

- Move process from memory to disk (swap space)
 - Process is blocked or suspended
- Move process from swap space to big enough partition
 - Process is ready
 - Set up Base and Limit registers
- Memory Manager (MM) and Process scheduler work together
 - Scheduler keeps track of all processes
 - MM keeps track of memory
 - Scheduler marks processes as swap-able and notifies MM to swap in processes
 - Scheduler policy must account for swapping overhead
 - MM policy must account for need to have memory space for ready processes

Can we do better?

User's view of a program



Memory management — beyond partitions


■ Can we improve memory utilization & performance

- Processes have distinct parts
 - *Code* – program and maybe shared libraries
 - *Data* – pre-allocated and heap
 - *Stack*
- Solution – slightly more Memory Management hardware
 - Multiple sets of “base and limit” registers
 - Divide process into logical pieces called *segments*


■ Advantages of *segments*

- Code segments don't need to be swapped out and may be shared
- Stack and heap can be grown – may require segment swap
- With separate I and D spaces can have larger virtual address spaces
 - “I” = *Instruction* (i.e., code, always read-only)
 - “D” = *Data* (usually read-write)

Logical view of segmentation



user space



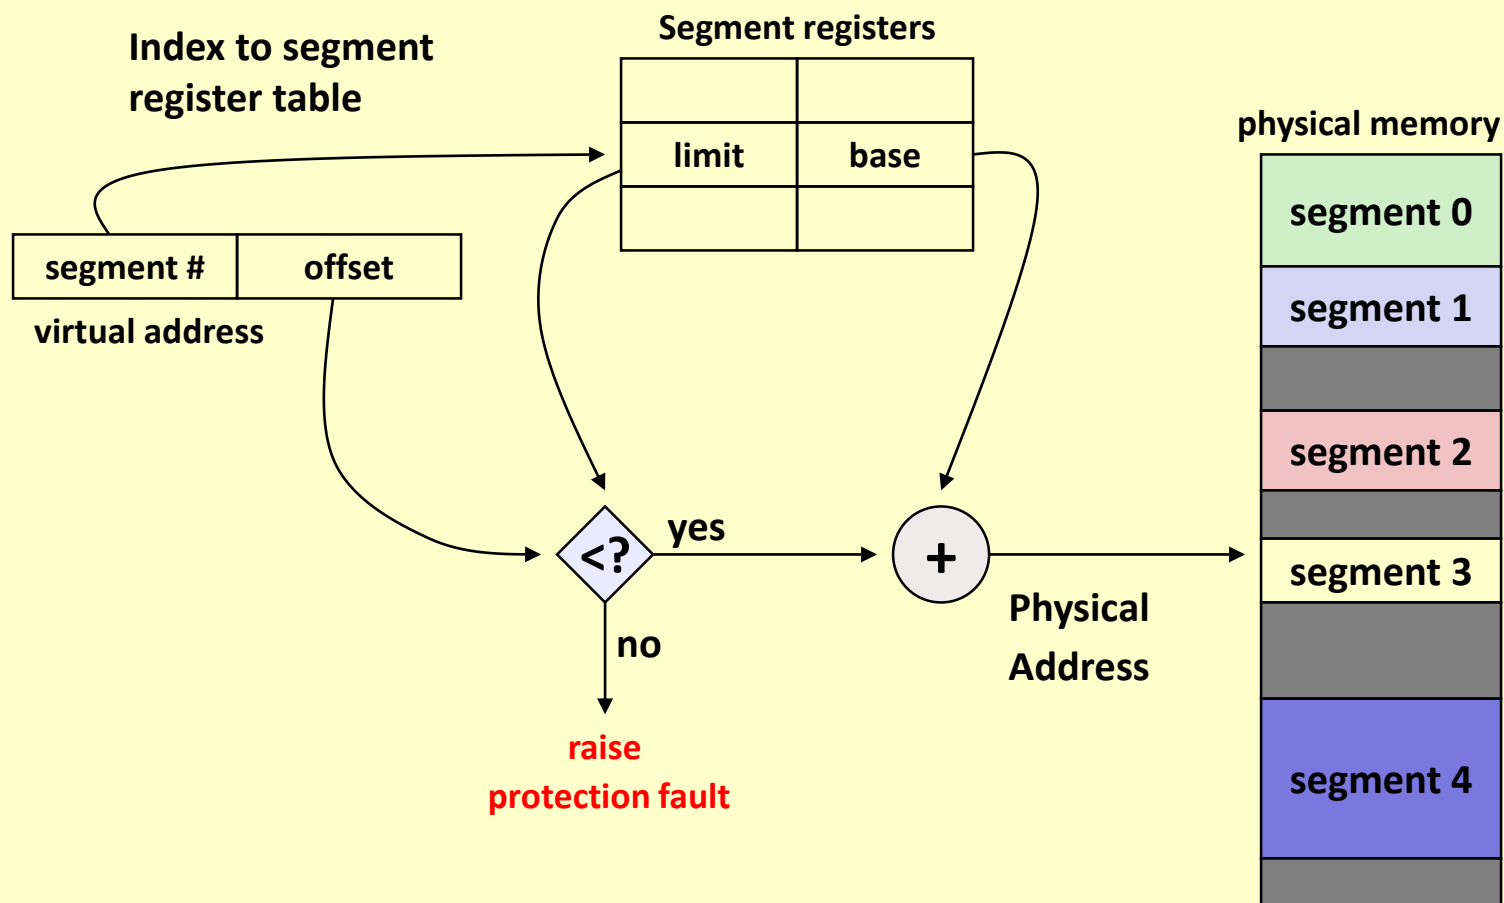
physical memory space

Segmentation

- Logical address consists of a pair:–
`<segment-number, offset>`
- Segment table – maps two-dimensional physical addresses; each table entry has:
 - *Base*: contains the starting physical address where the segments reside in memory.
 - *Limit*: specifies the length of the segment.

OSTEP §16-17

Segment lookup



Segmentation

- ***Protection.*** With each pair of segment registers, include:
 - *validation bit* = 0 \Rightarrow illegal segment
 - *read/write/execute* privileges
- ***Protection bits*** associated with segments; code sharing occurs at segment level.
- Since segments vary in length, memory allocation is a **dynamic storage-allocation problem**
 - With all the problems of fragmentation!

Segmentation

- **Common in early minicomputers**
 - Small amount of additional hardware – 4 or 8 segments
 - Used effectively in classical Unix
- **Good idea that has persisted and supported in current hardware and OSs**
 - Pentium, x86 supports segments
 - Linux supports segments (sort of)
- **Still have *external fragmentation* of memory**
- **What is the next level of Memory Management improvement?**
 - Next topic

Questions?